

Depth and Transparency

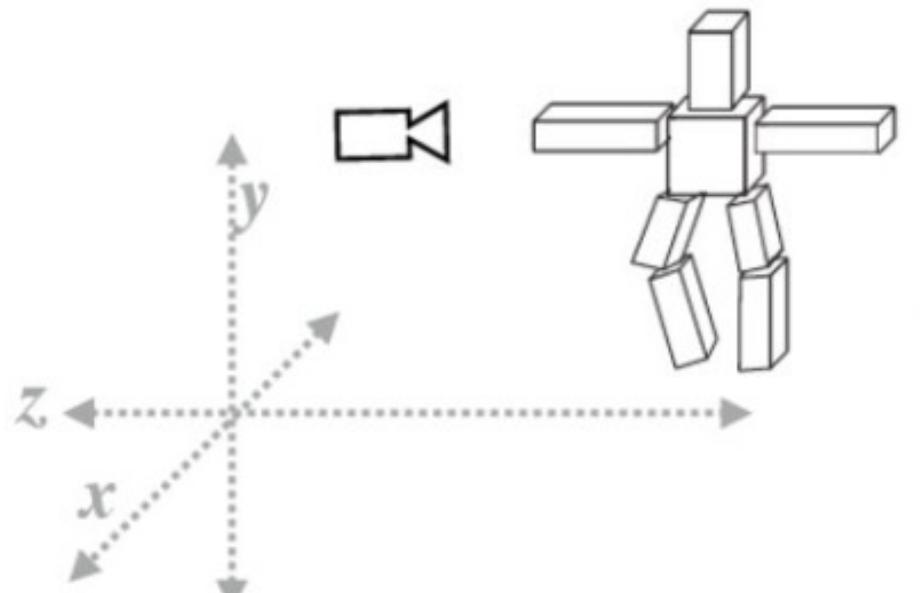
Today: Wrap up the rasterization pipeline!

Remember our goal:

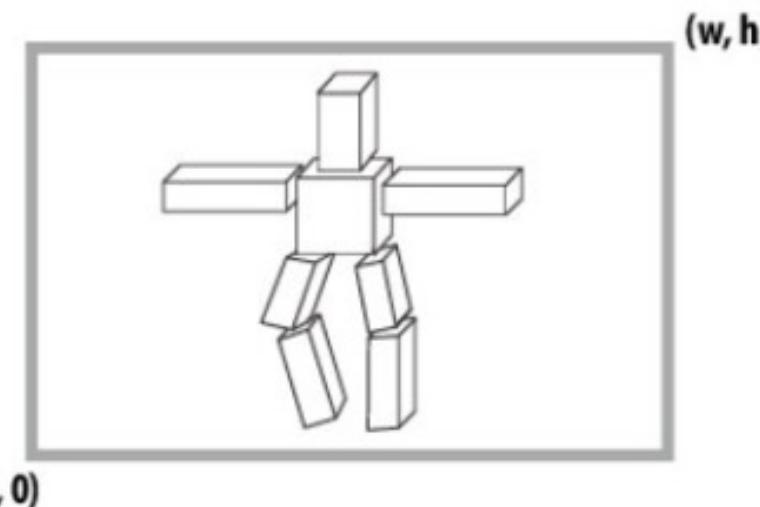
- Start with INPUTS (triangles)
 - possibly w/ other data (e.g., colors or texture coordinates)
- Apply a series of transformations: STAGES of pipeline
- Produce OUTPUT (final image)



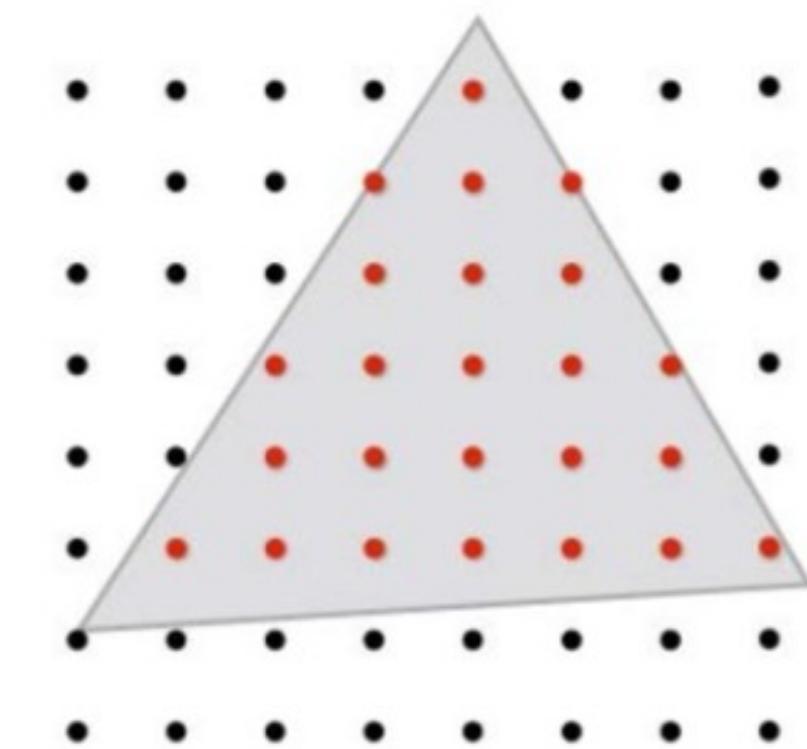
What we know how to do so far...



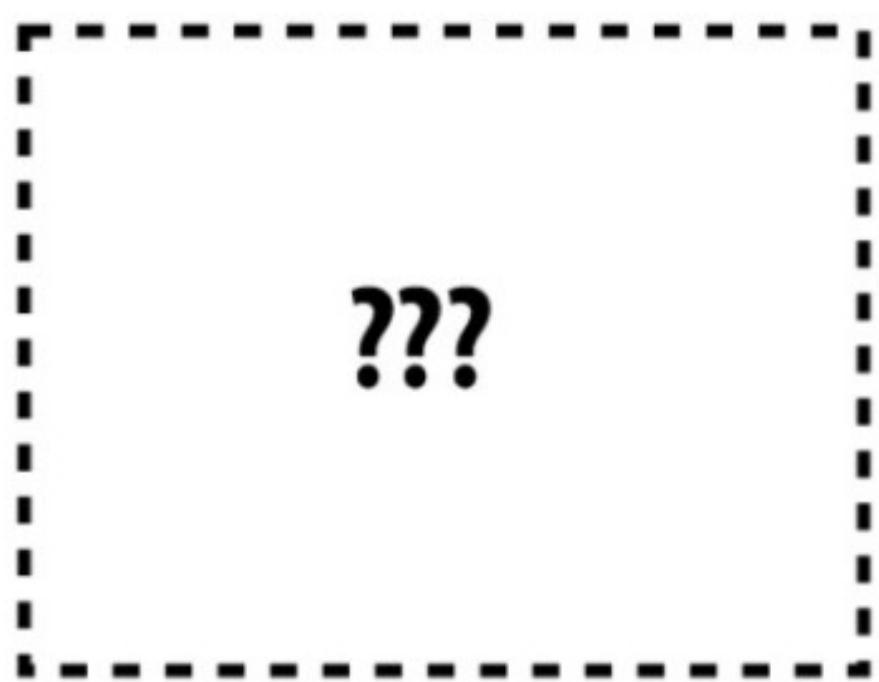
**position objects in the world
(3D transformations)**



**project objects onto the screen
(perspective projection)**



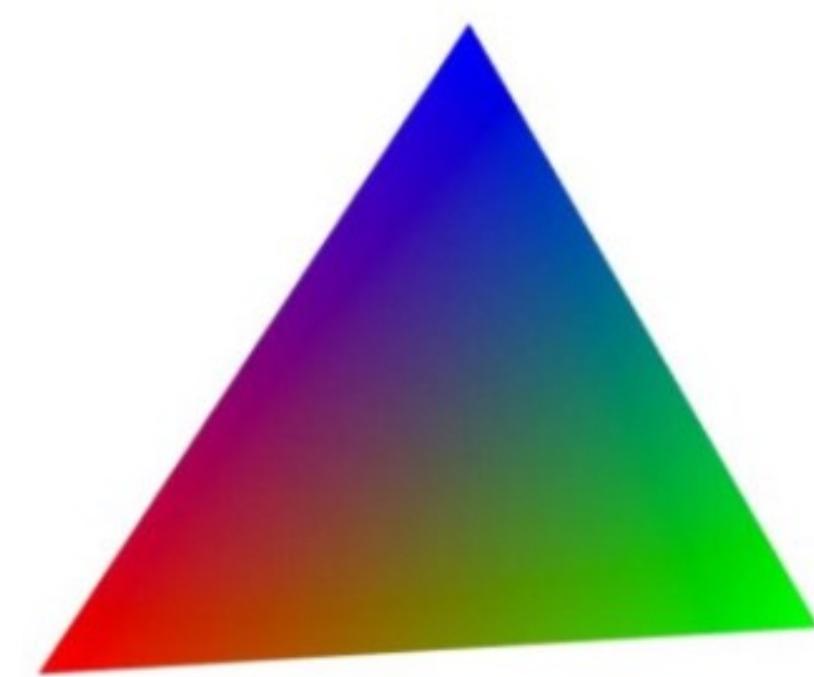
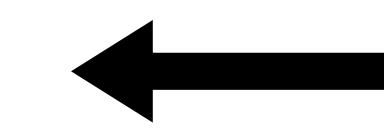
**sample triangle coverage
(rasterization)**



**put samples into frame buffer
(depth & alpha)**



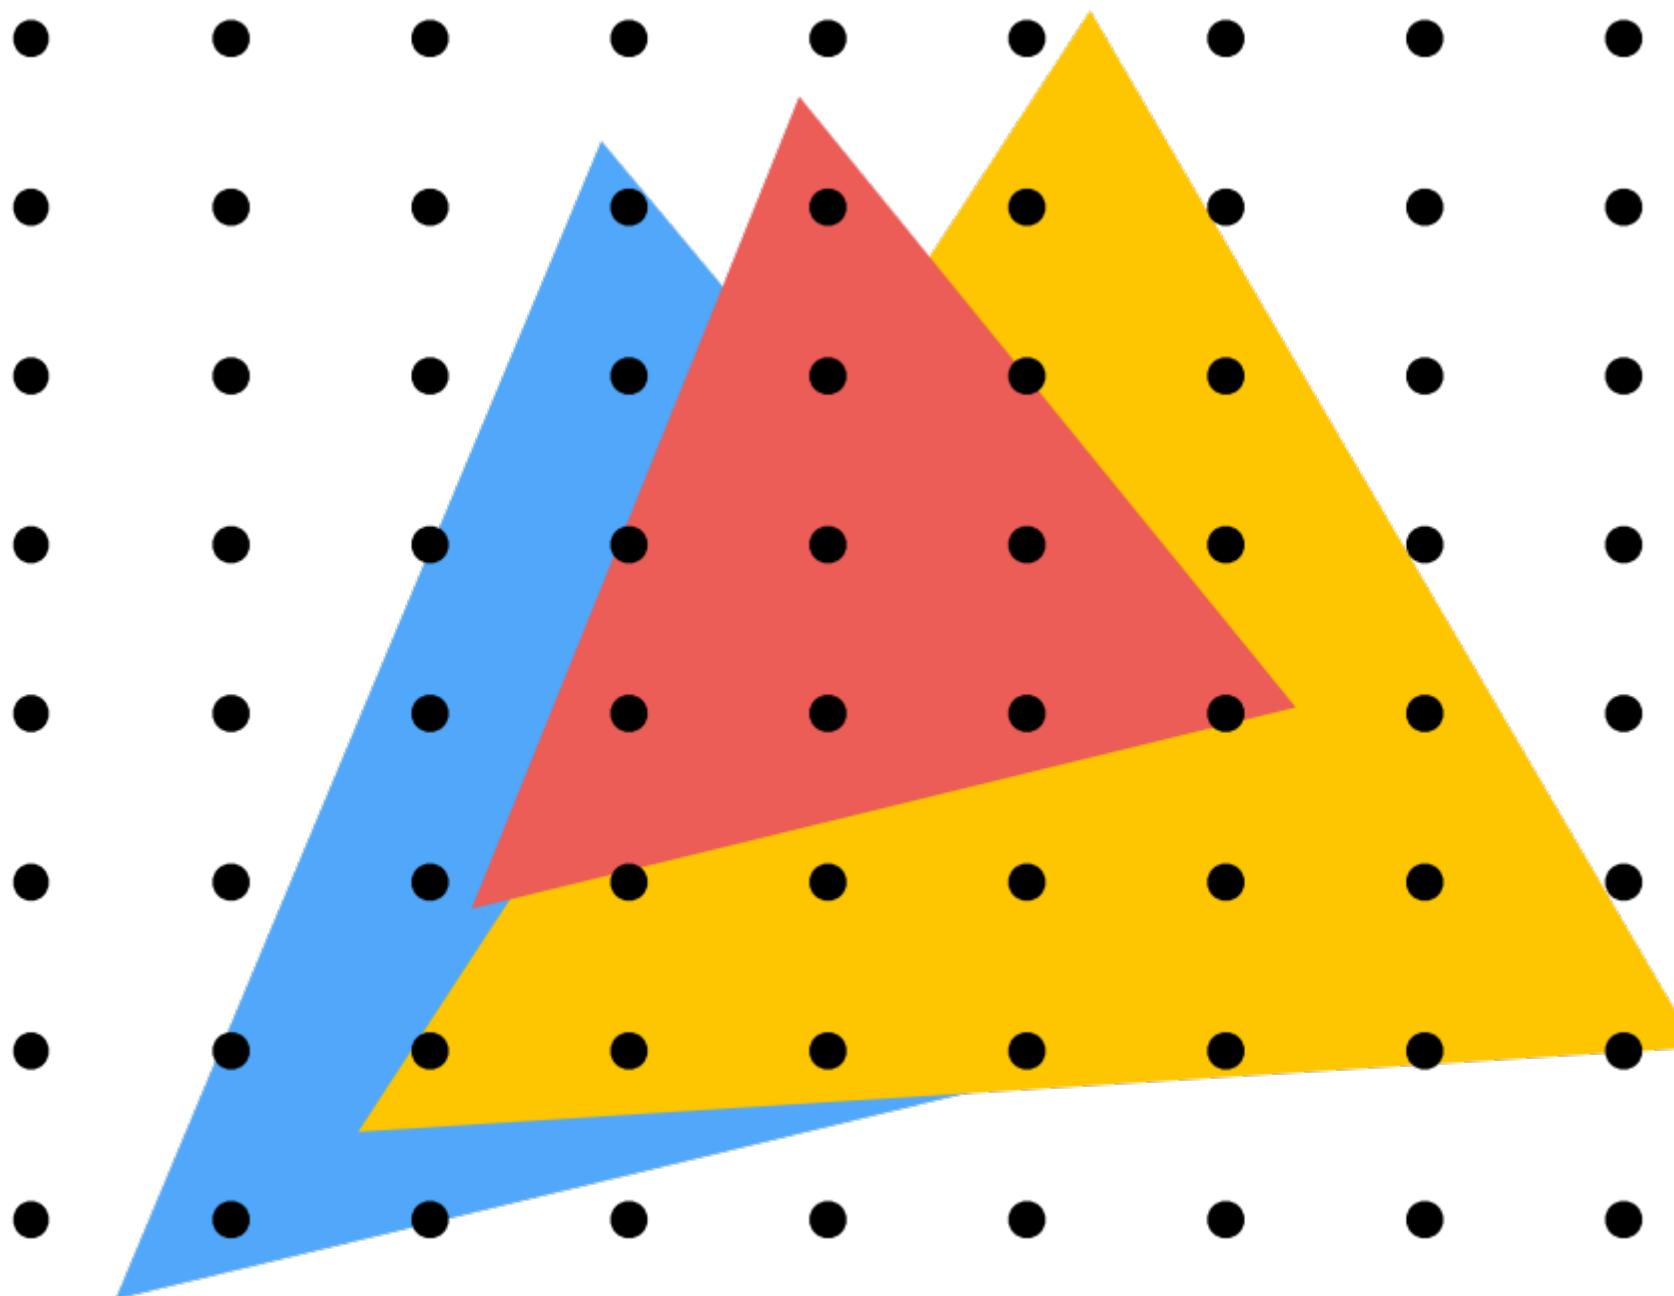
**sample texture maps
(filtering, mipmapping)**



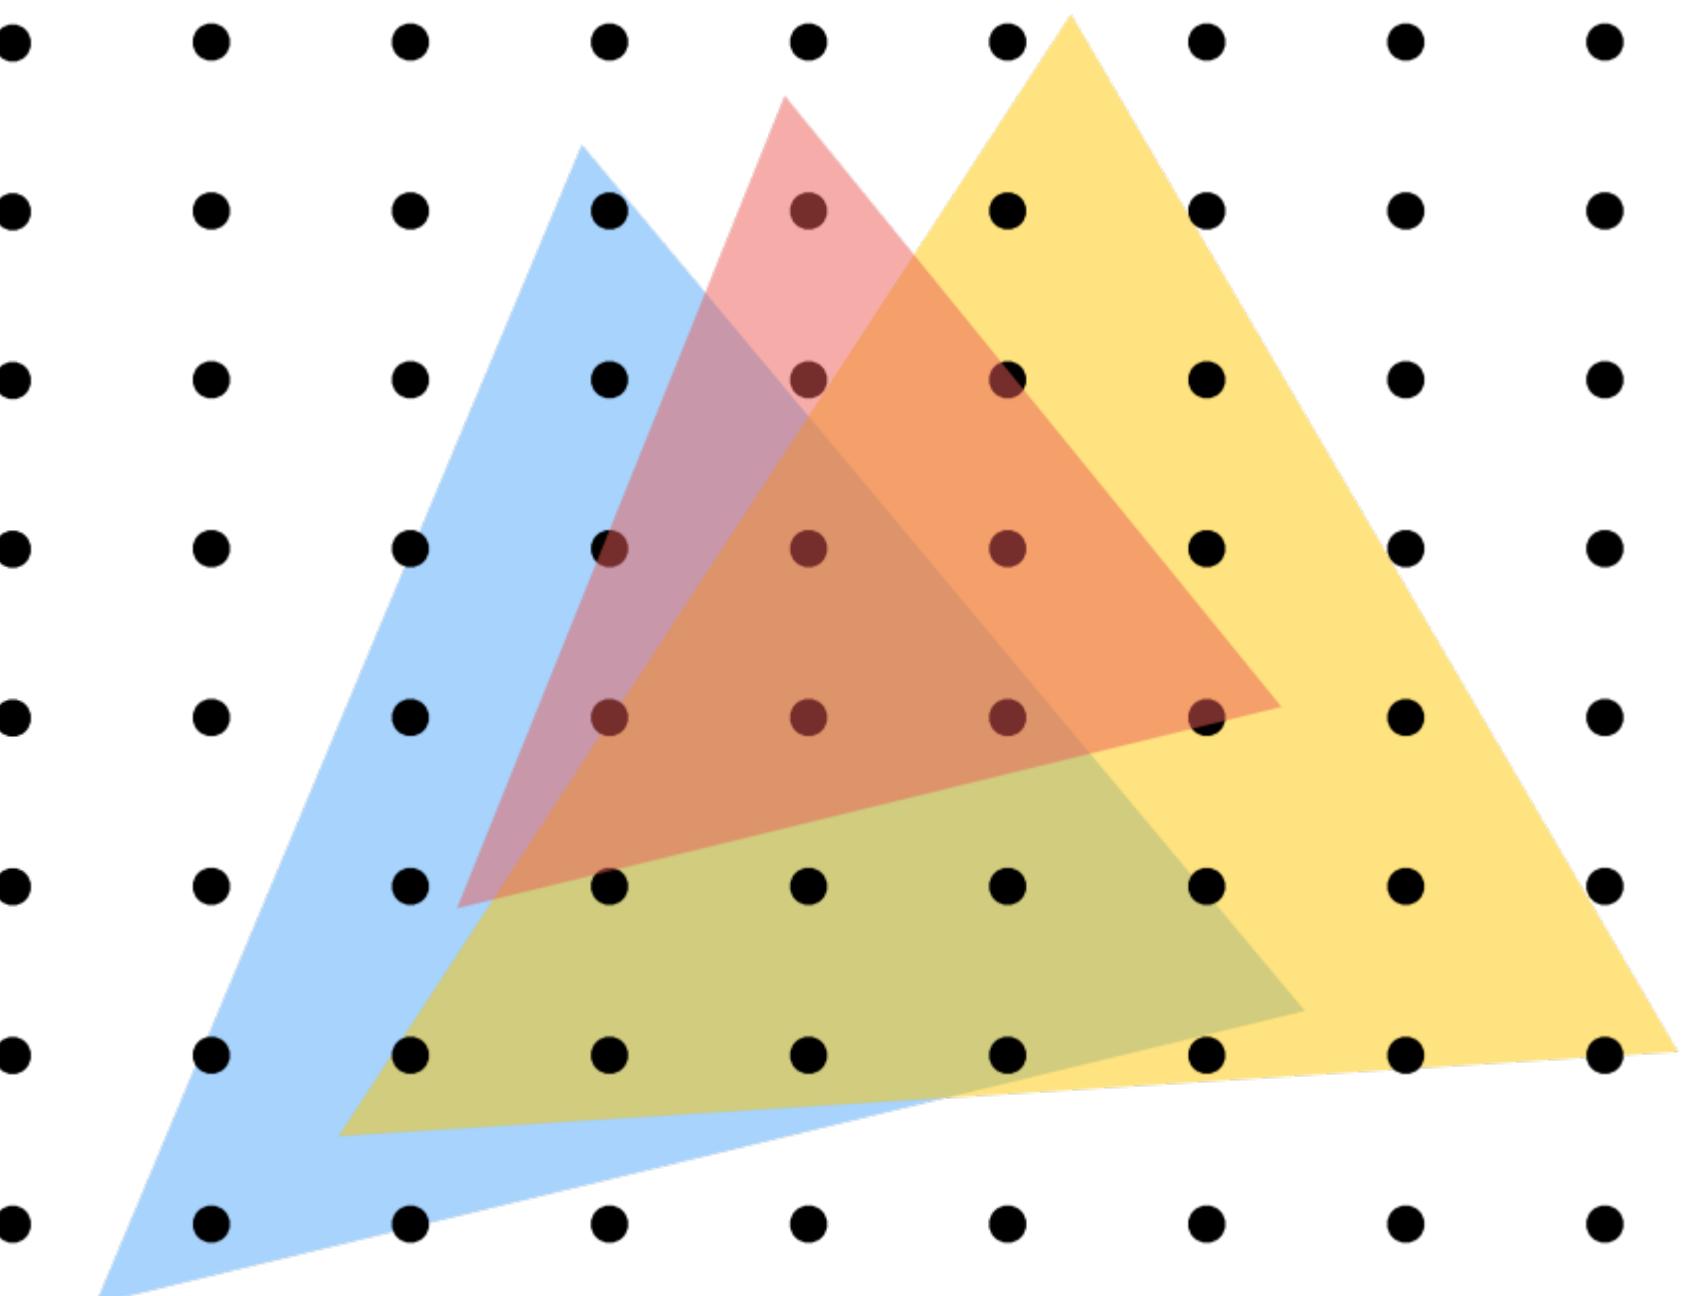
**interpolate vertex attributes
(barycentric coordinates)**

Occlusion

Occlusion: which triangle is visible at each covered sample point?



Opaque Triangles

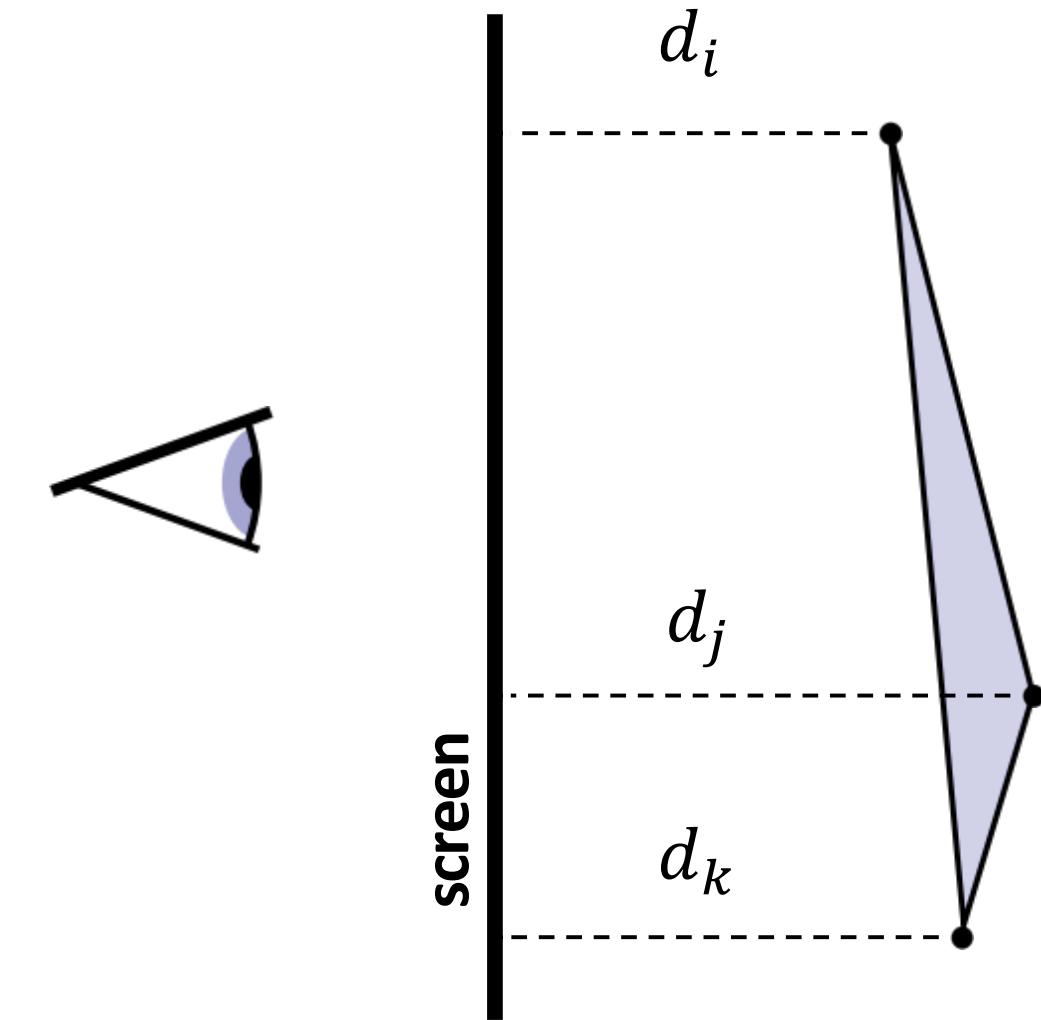
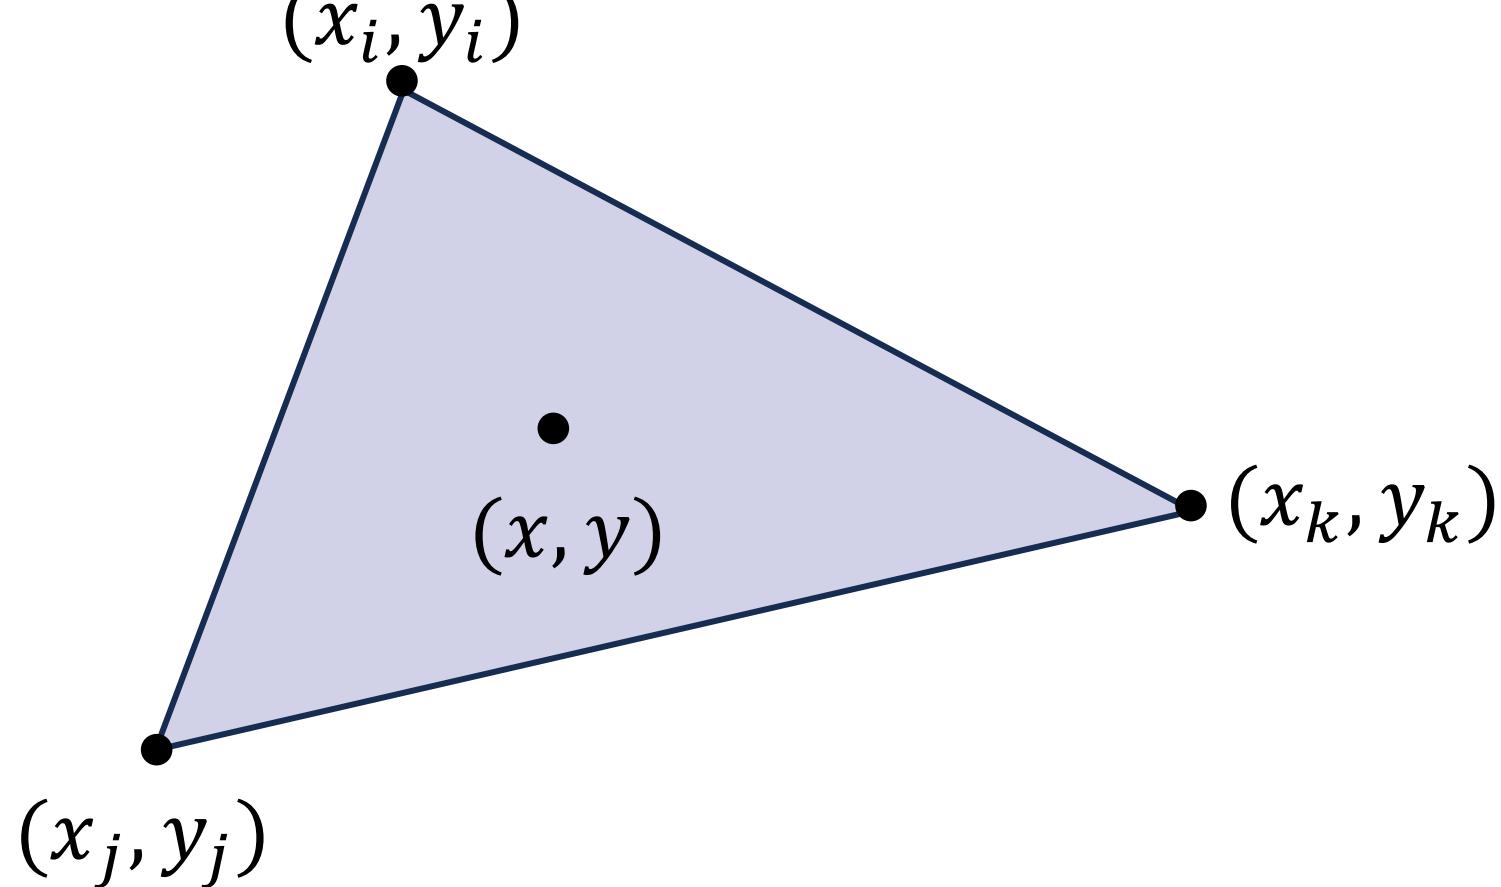


50% transparent triangles

Sampling Depth

Assume we have a triangle given by:

- the projected 2D coordinates (x_i, y_i) of each vertex
- the “depth” d_i of each vertex (i.e., distance from the viewer)

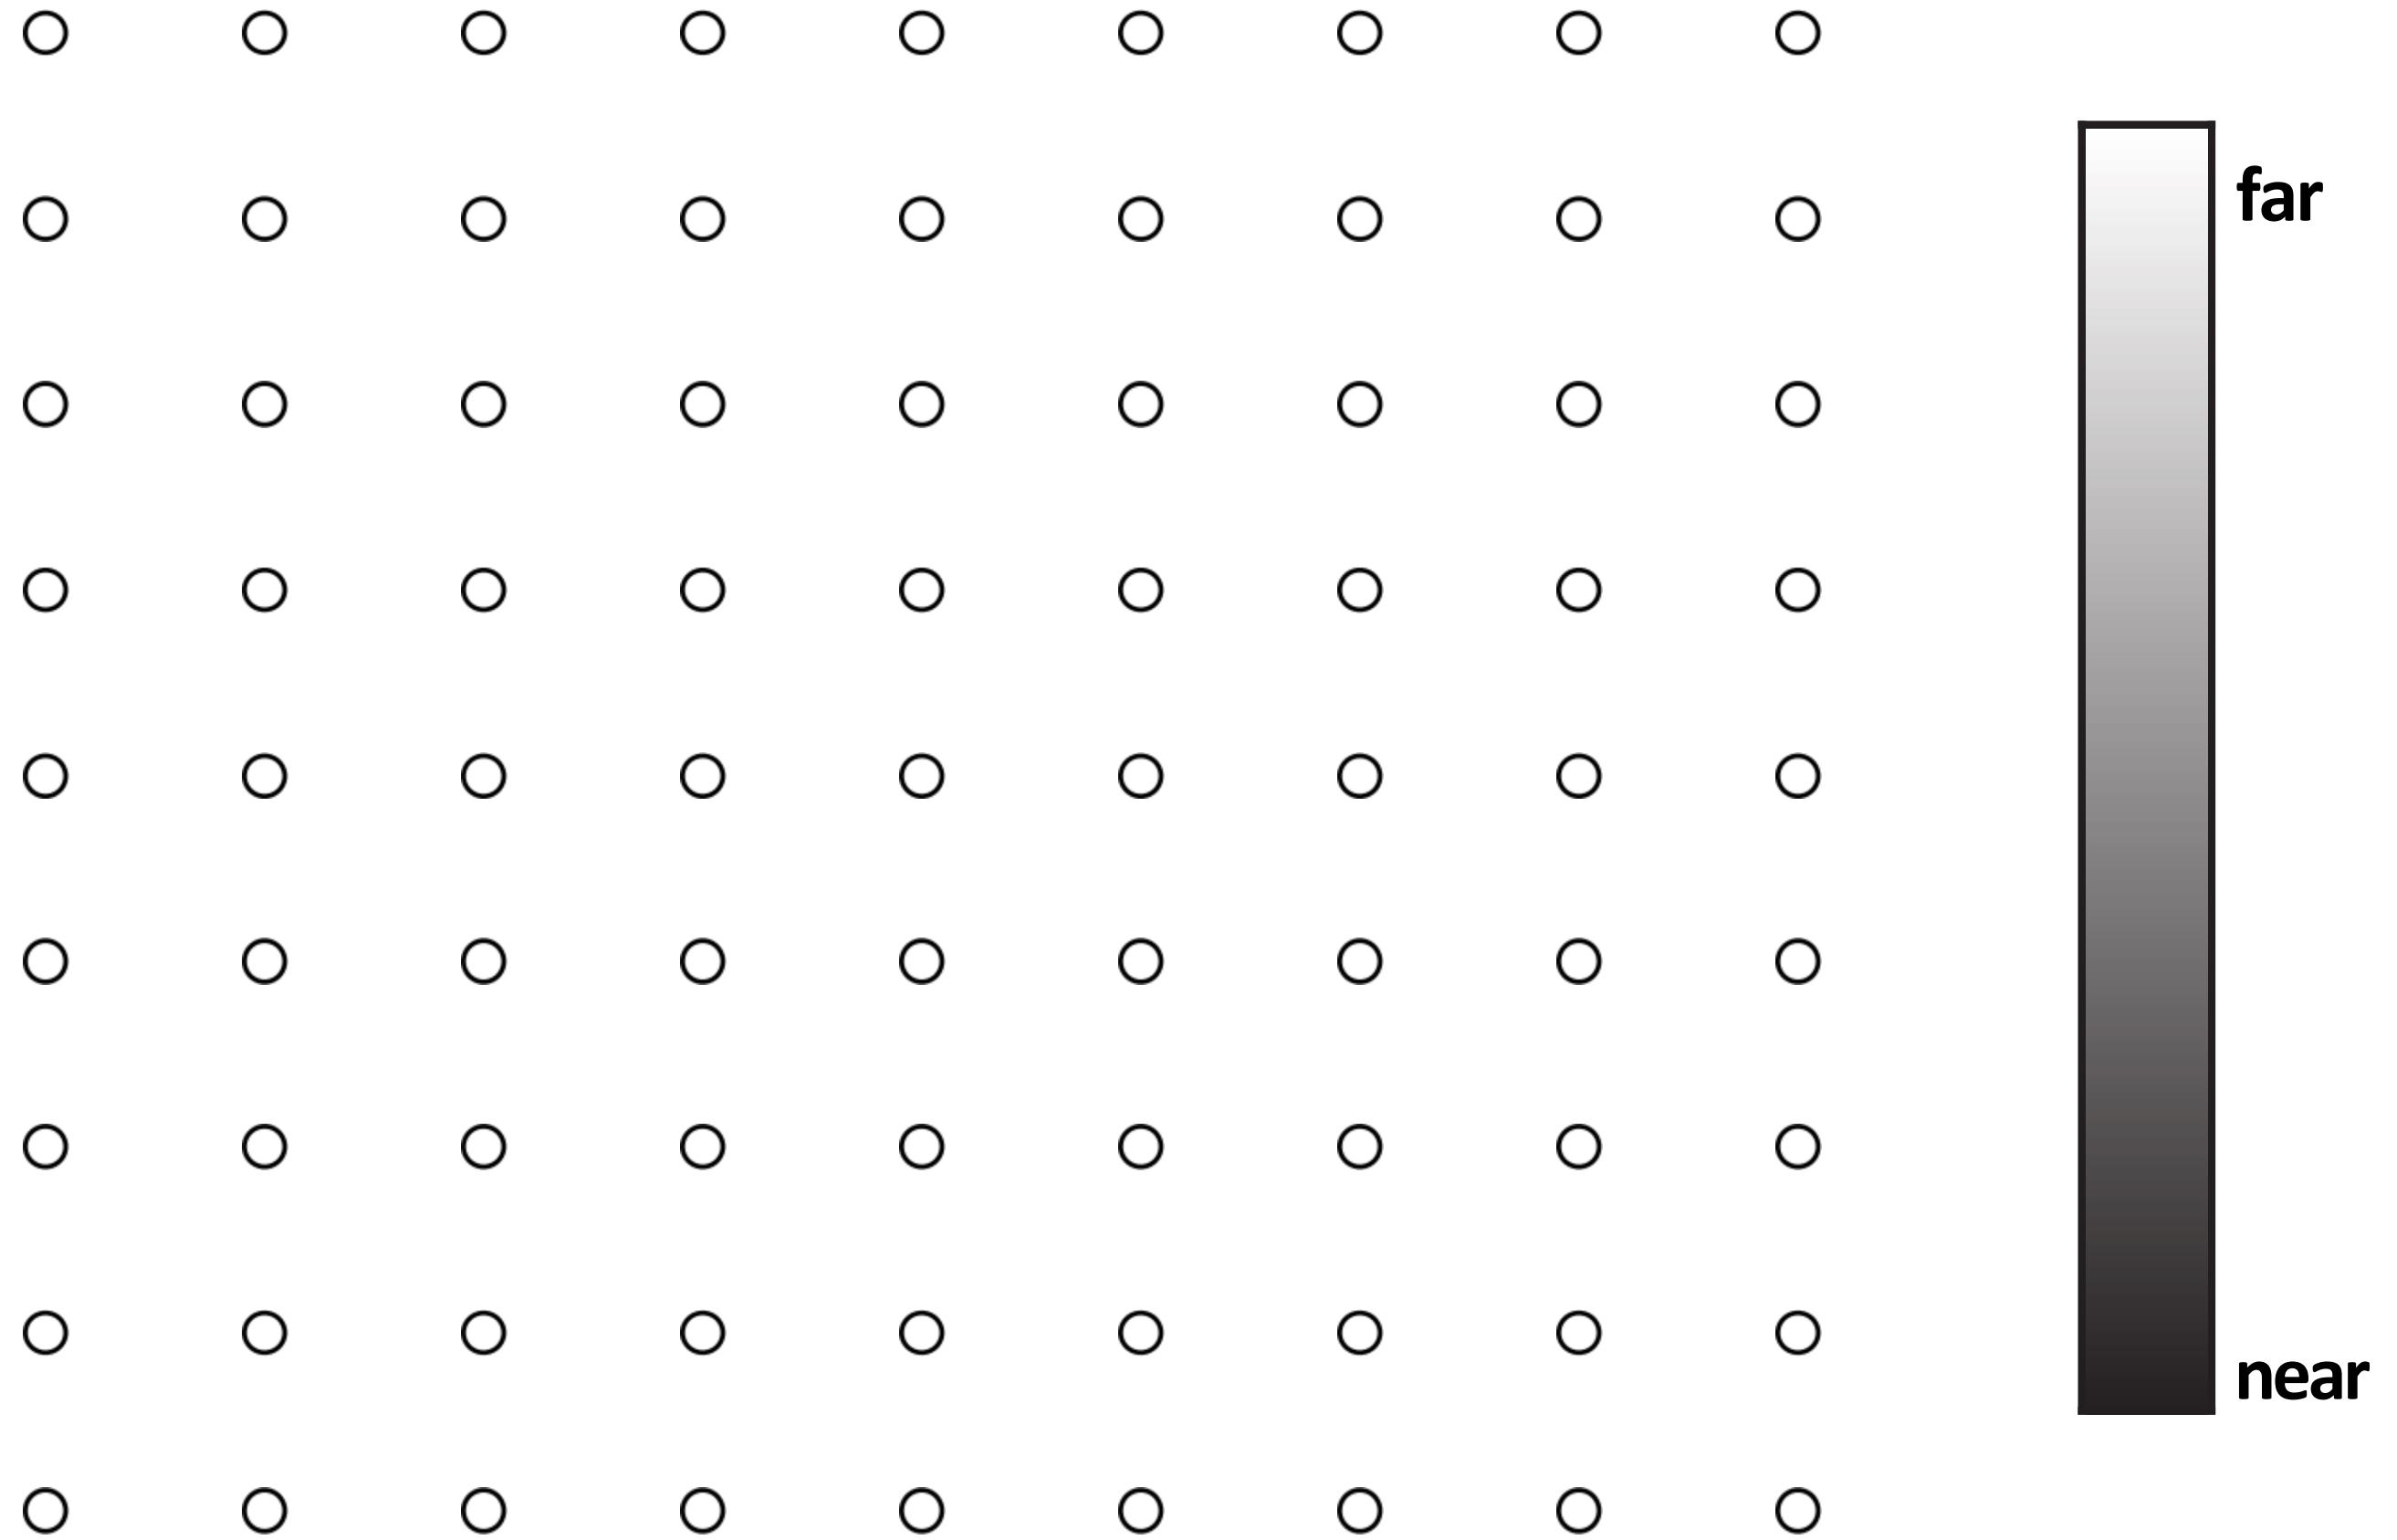


Q: How do we compute the depth d at a given sample point (x, y) ?

A: Interpolate it using barycentric coordinates—just like any other attribute that varies linearly over the triangle

The depth-buffer (Z-buffer)

For each sample, depth-buffer stores the depth of the **closest** triangle seen so far

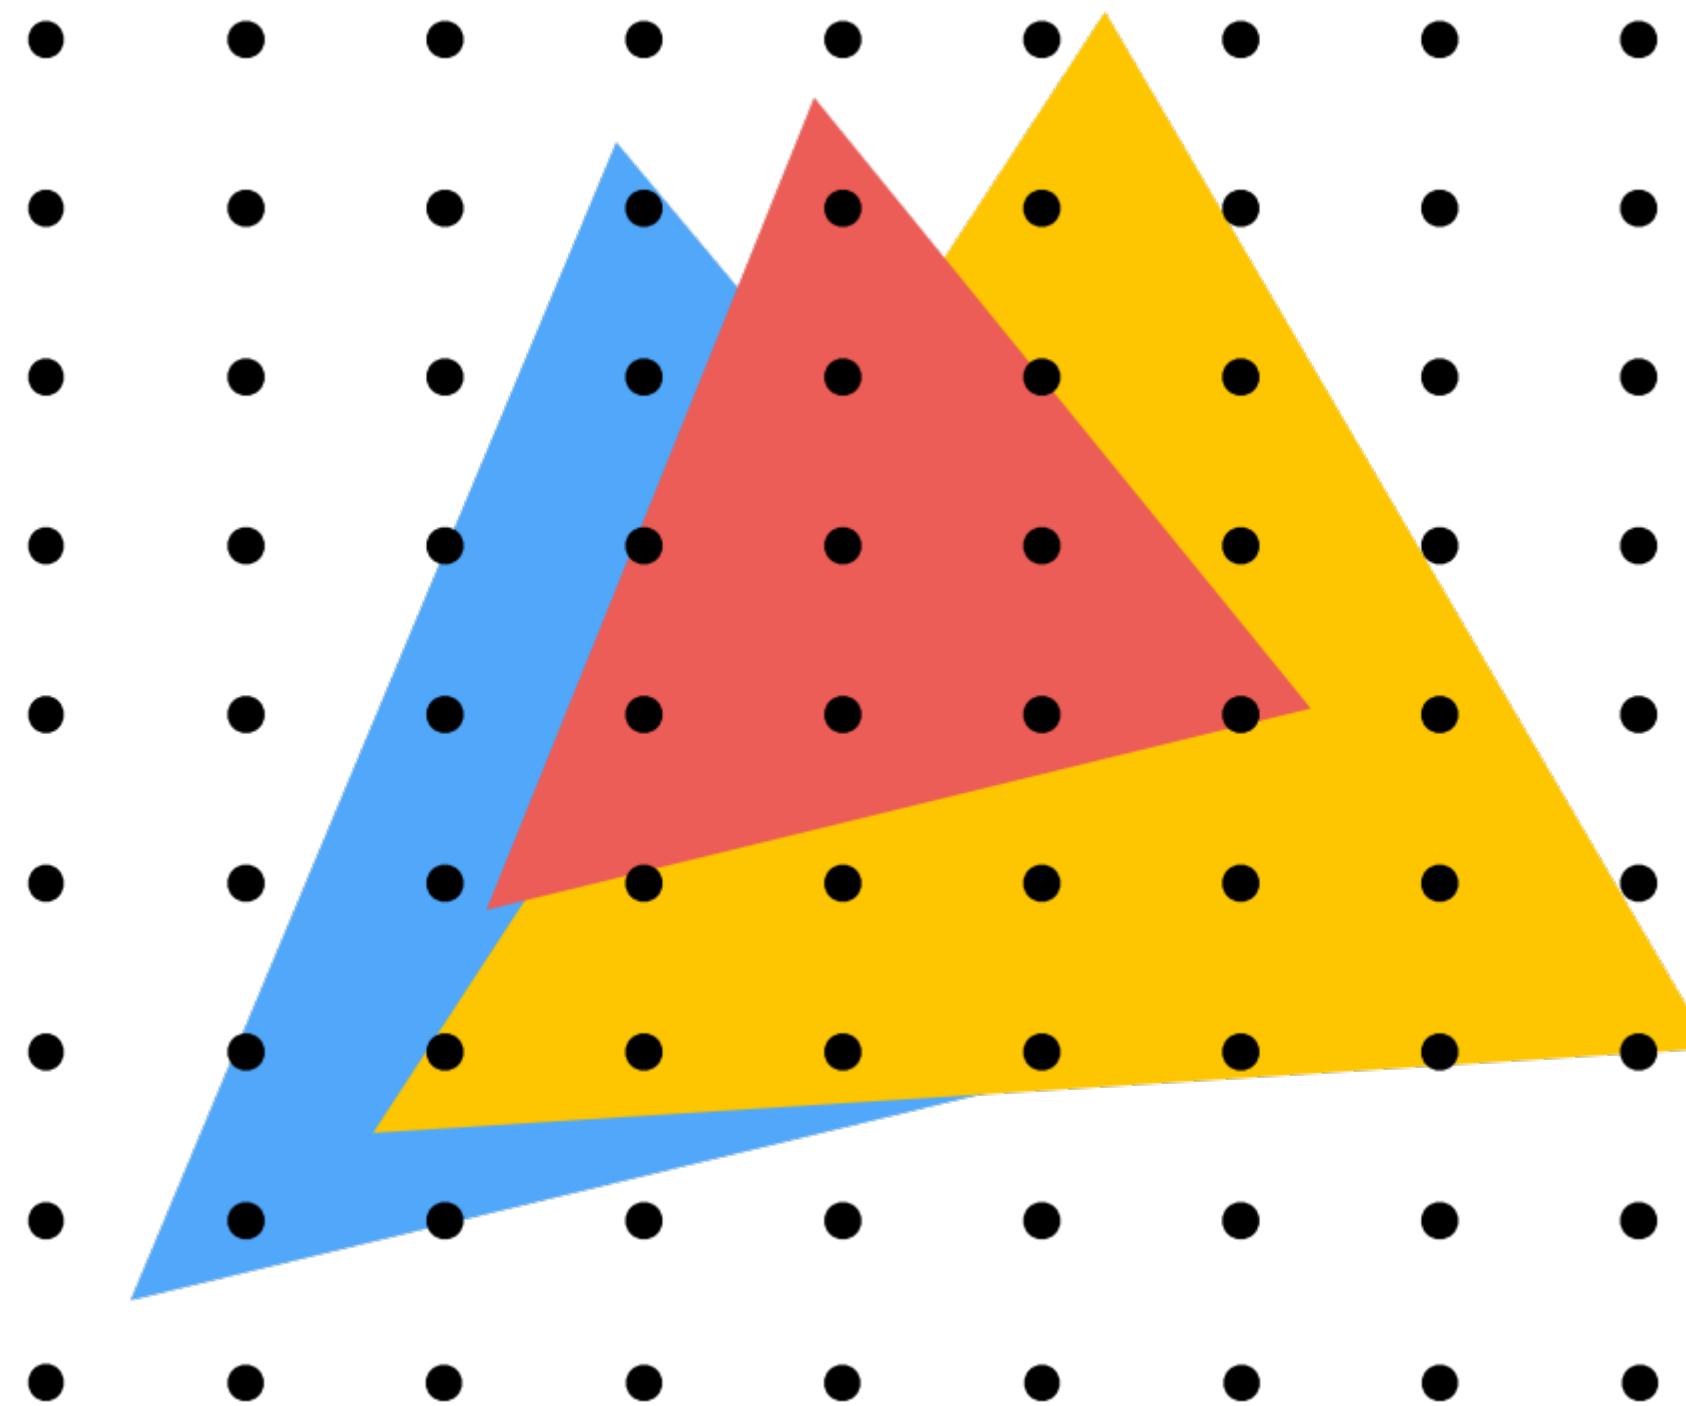


Initialize all depth buffer values to “infinity” (max value)

Depth buffer example

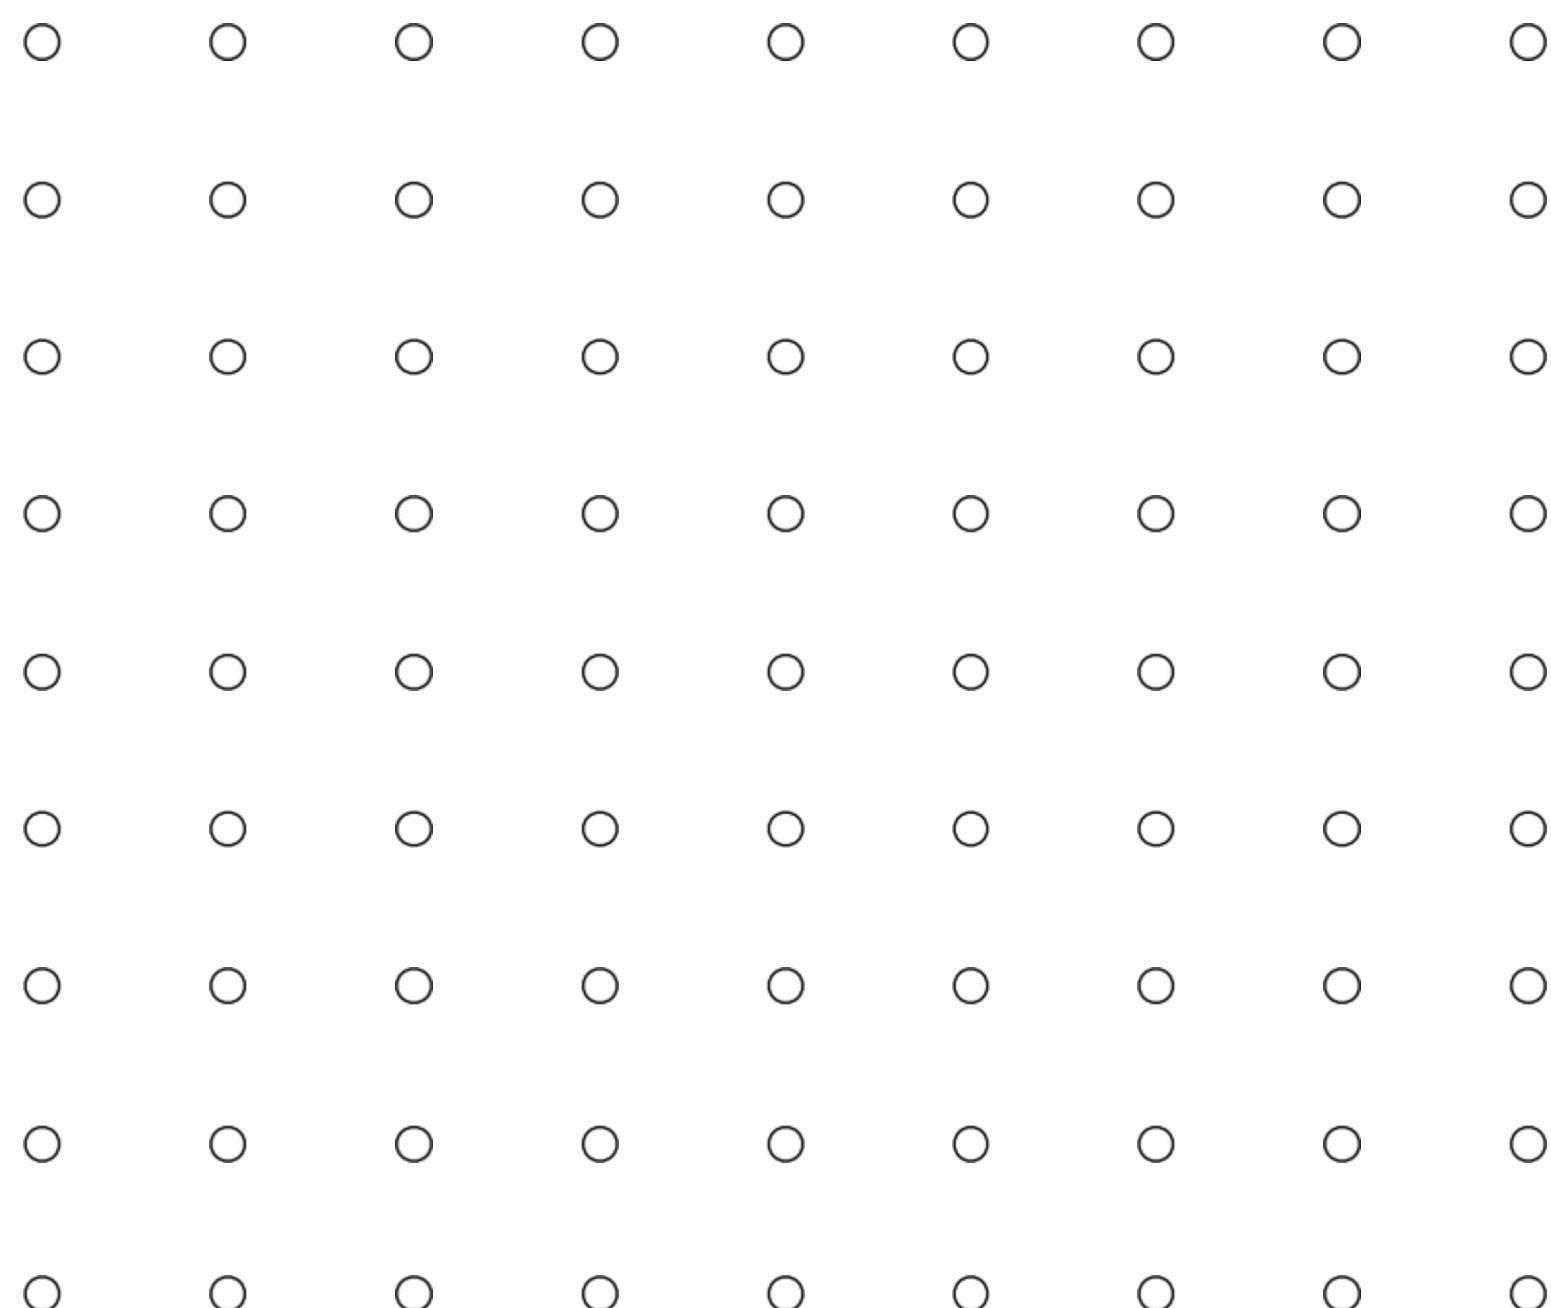


Example: rendering three opaque triangles



Occlusion using the depth-buffer (Z-buffer)

Processing yellow triangle:
depth = 0.5



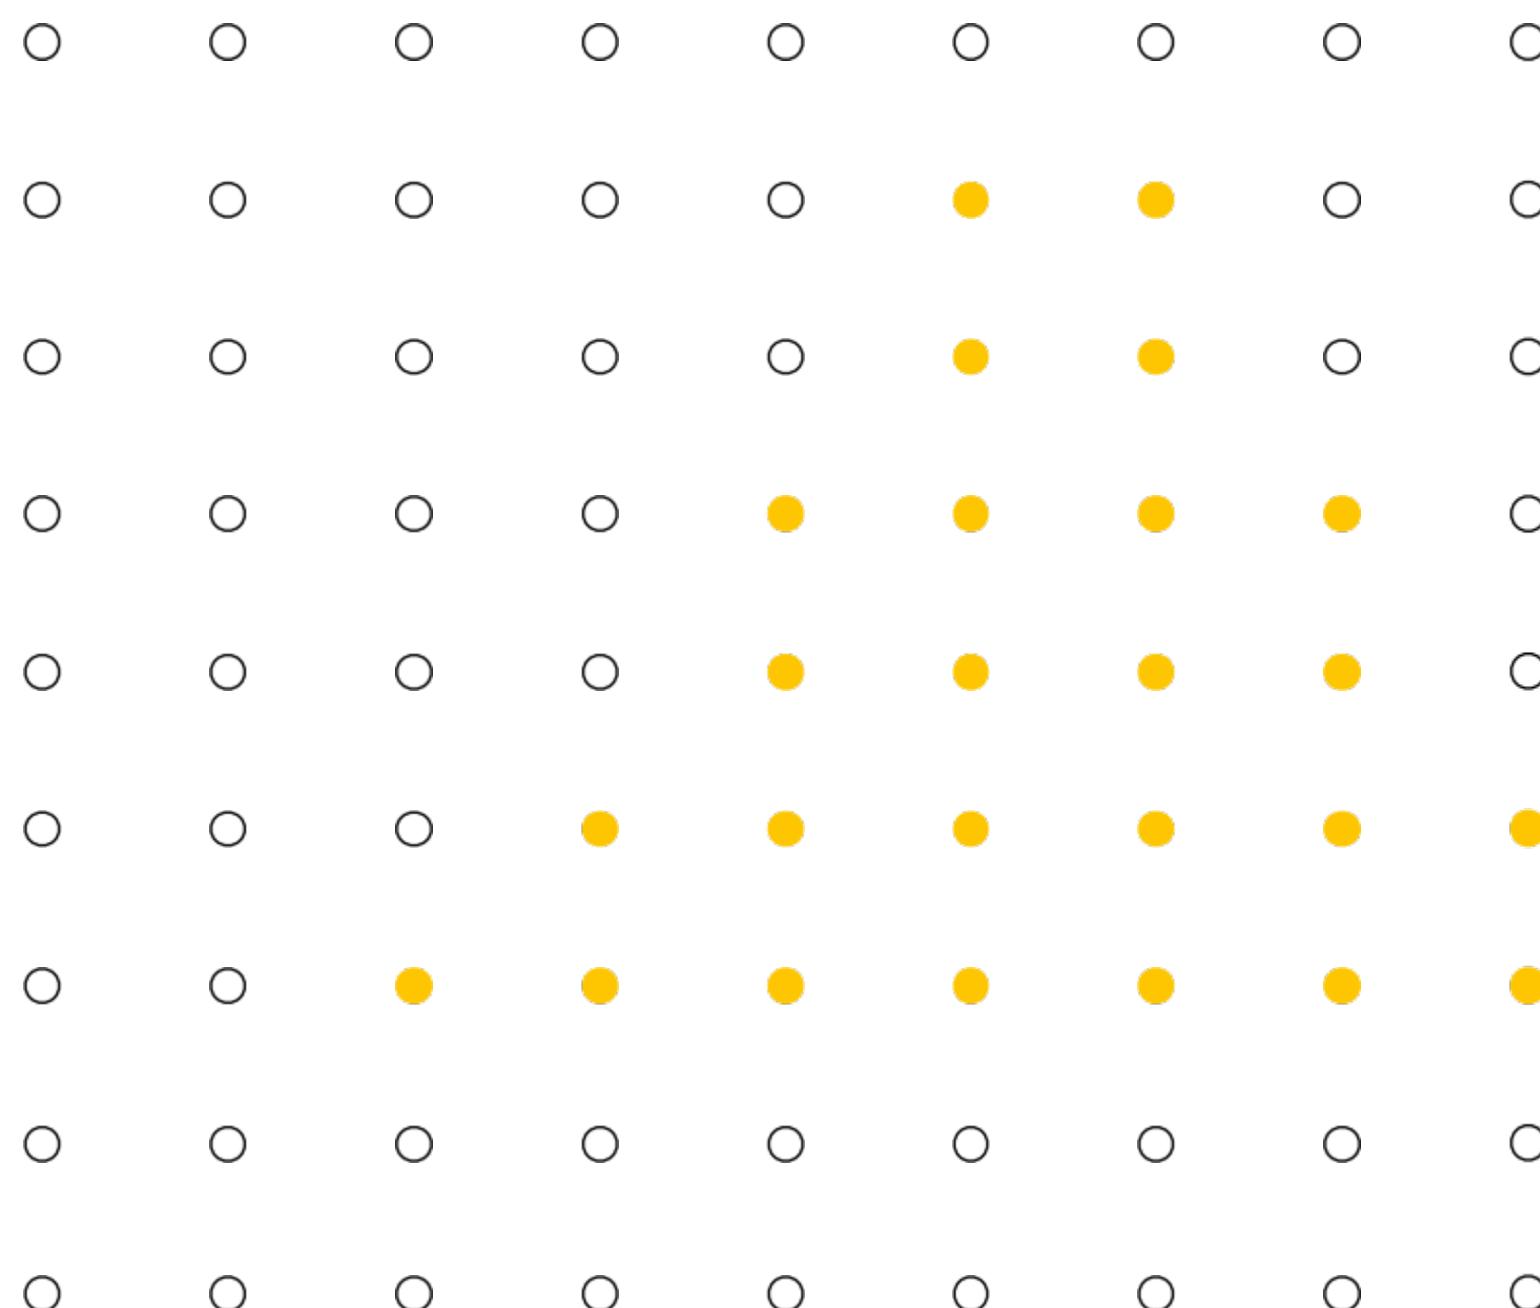
Color buffer contents



Depth buffer contents

Occlusion using the depth-buffer (Z-buffer)

After processing yellow triangle:



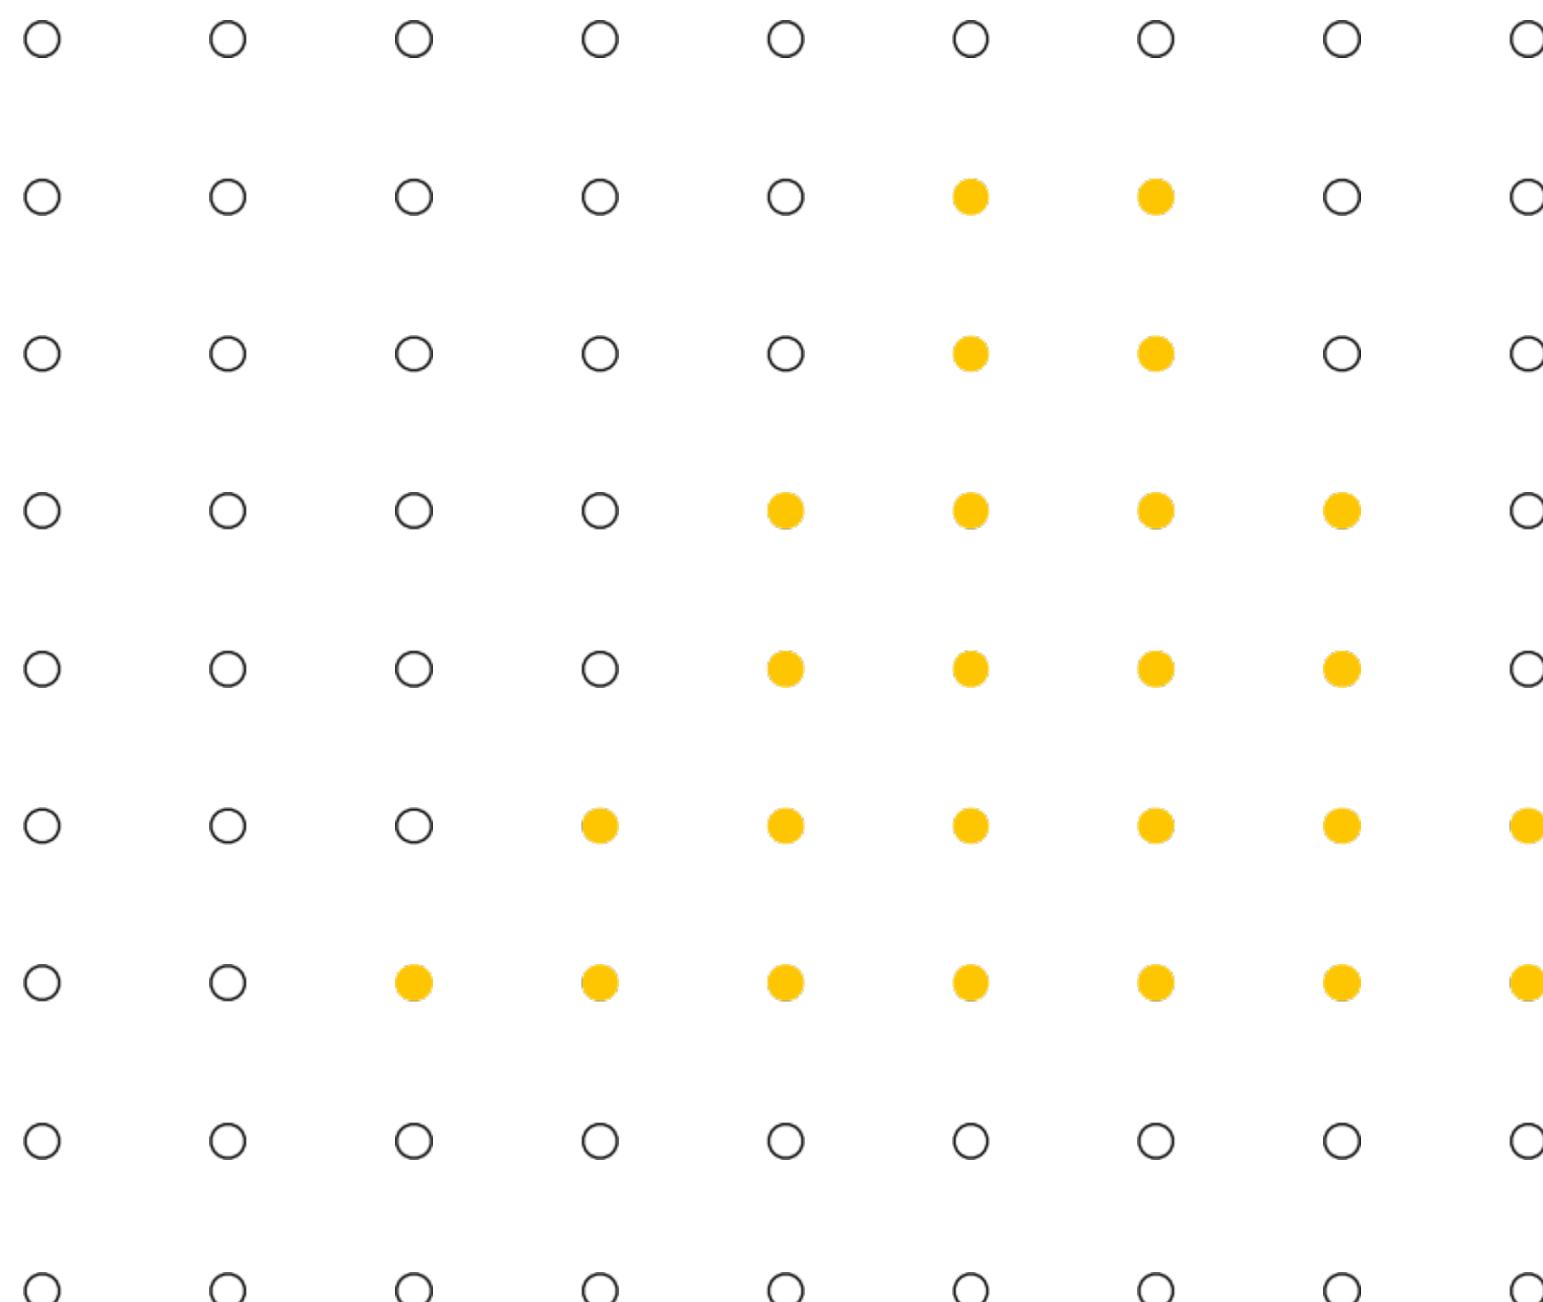
Color buffer contents



Depth buffer contents

Occlusion using the depth-buffer (Z-buffer)

Processing blue triangle:
depth = 0.75



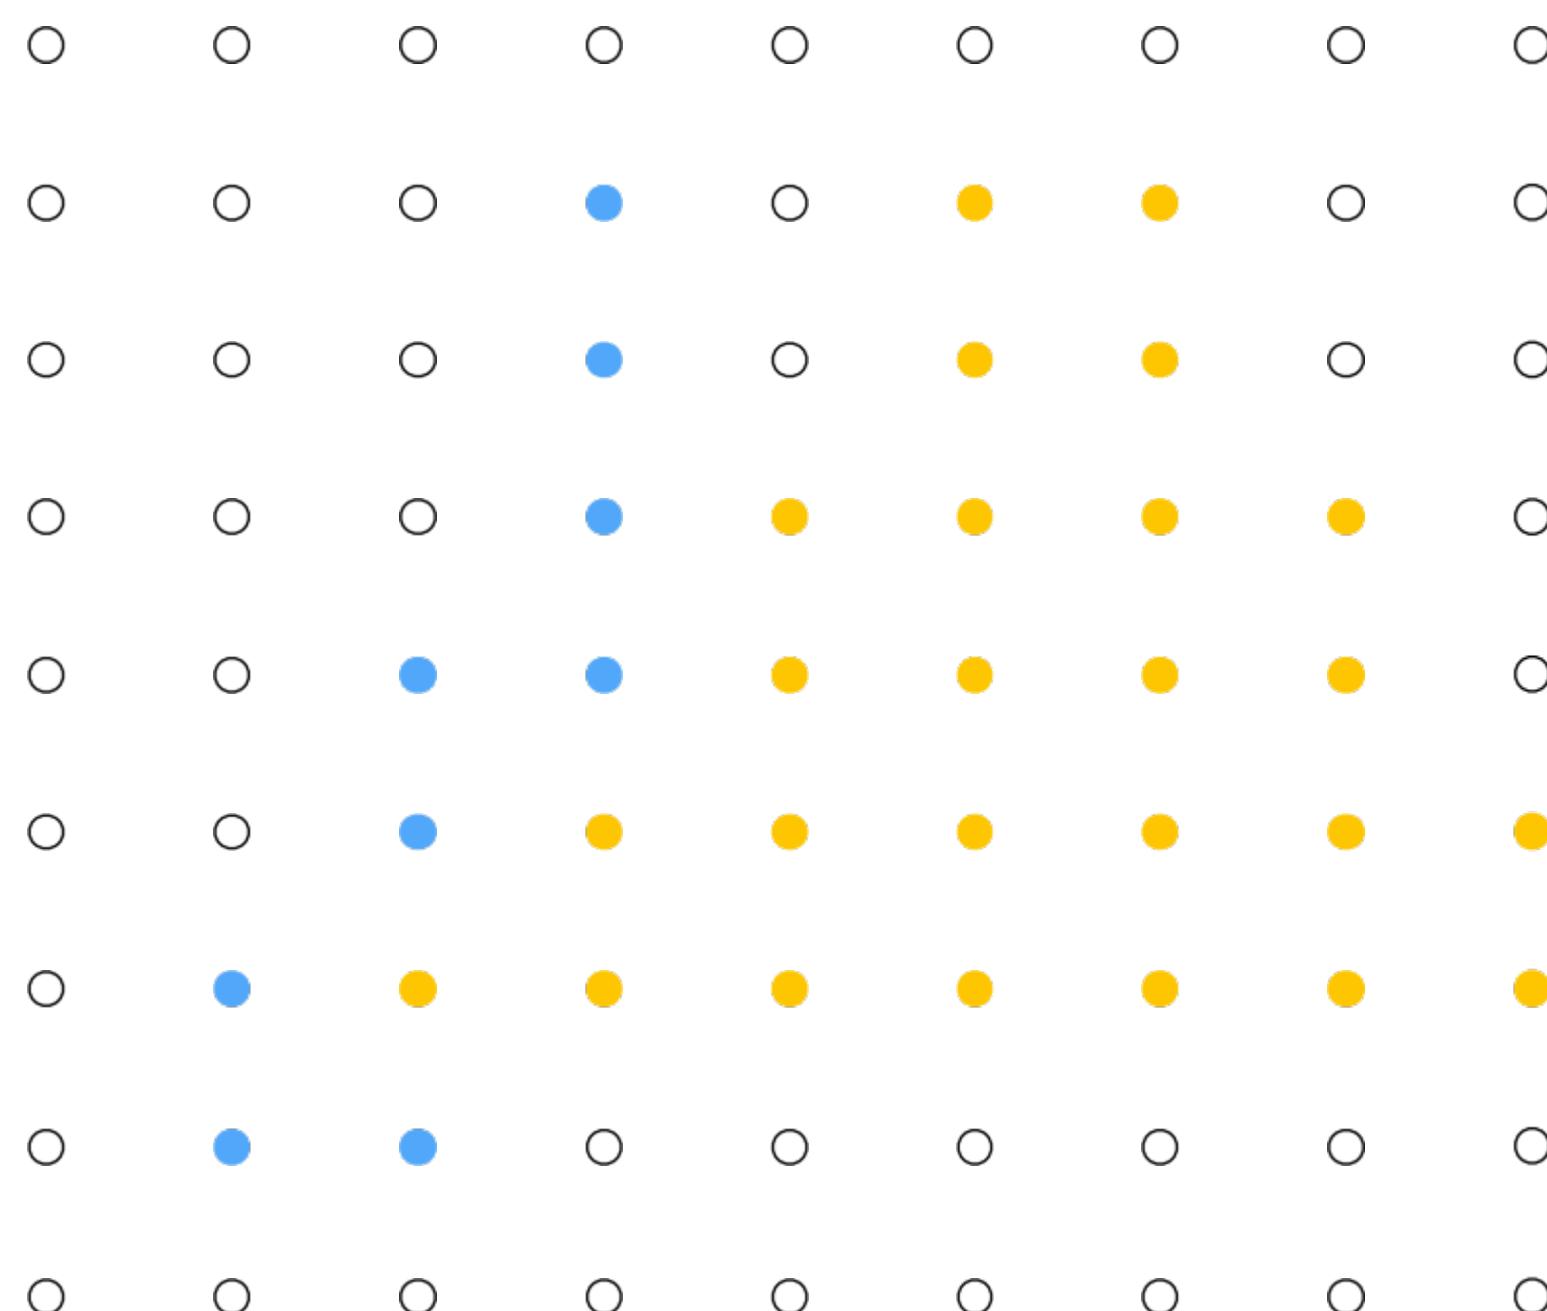
Color buffer contents



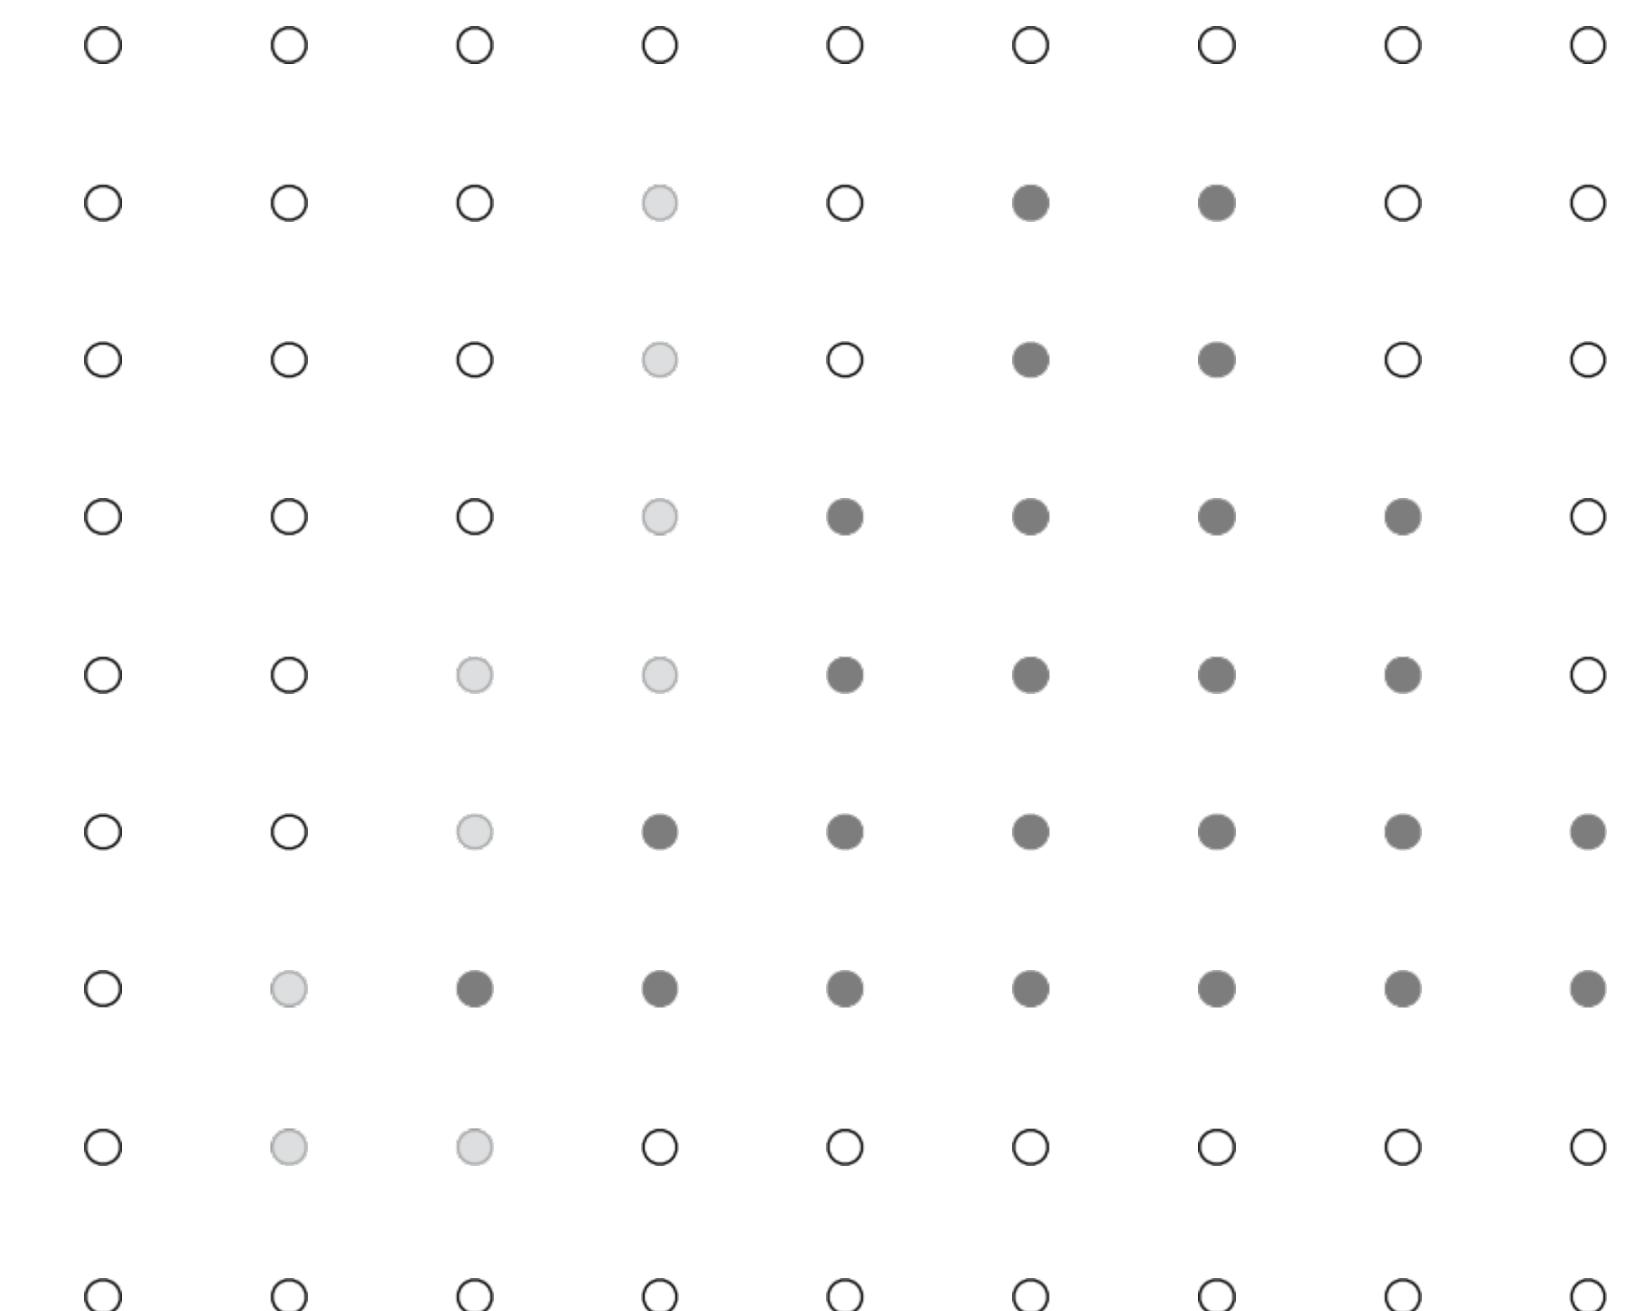
Depth buffer contents

Occlusion using the depth-buffer (Z-buffer)

After processing blue triangle:



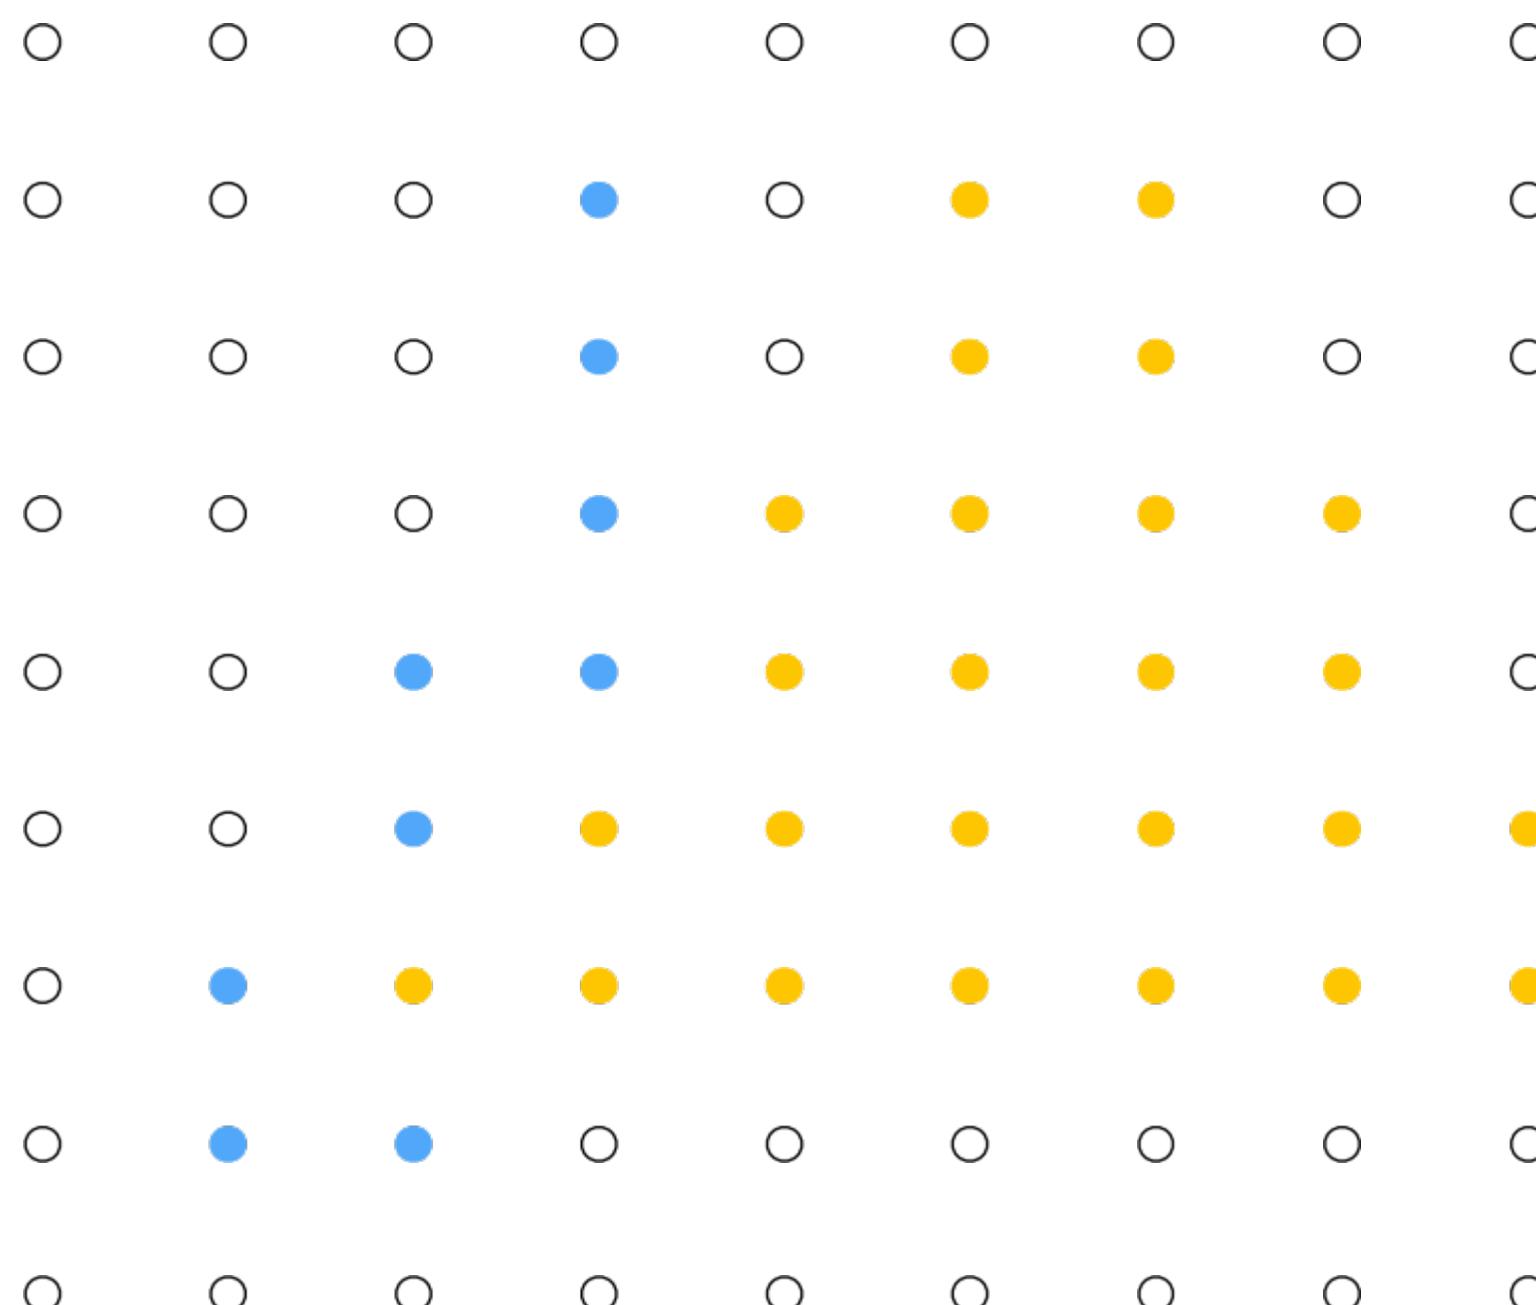
Color buffer contents



Depth buffer contents

Occlusion using the depth-buffer (Z-buffer)

Processing red triangle:
depth = 0.25



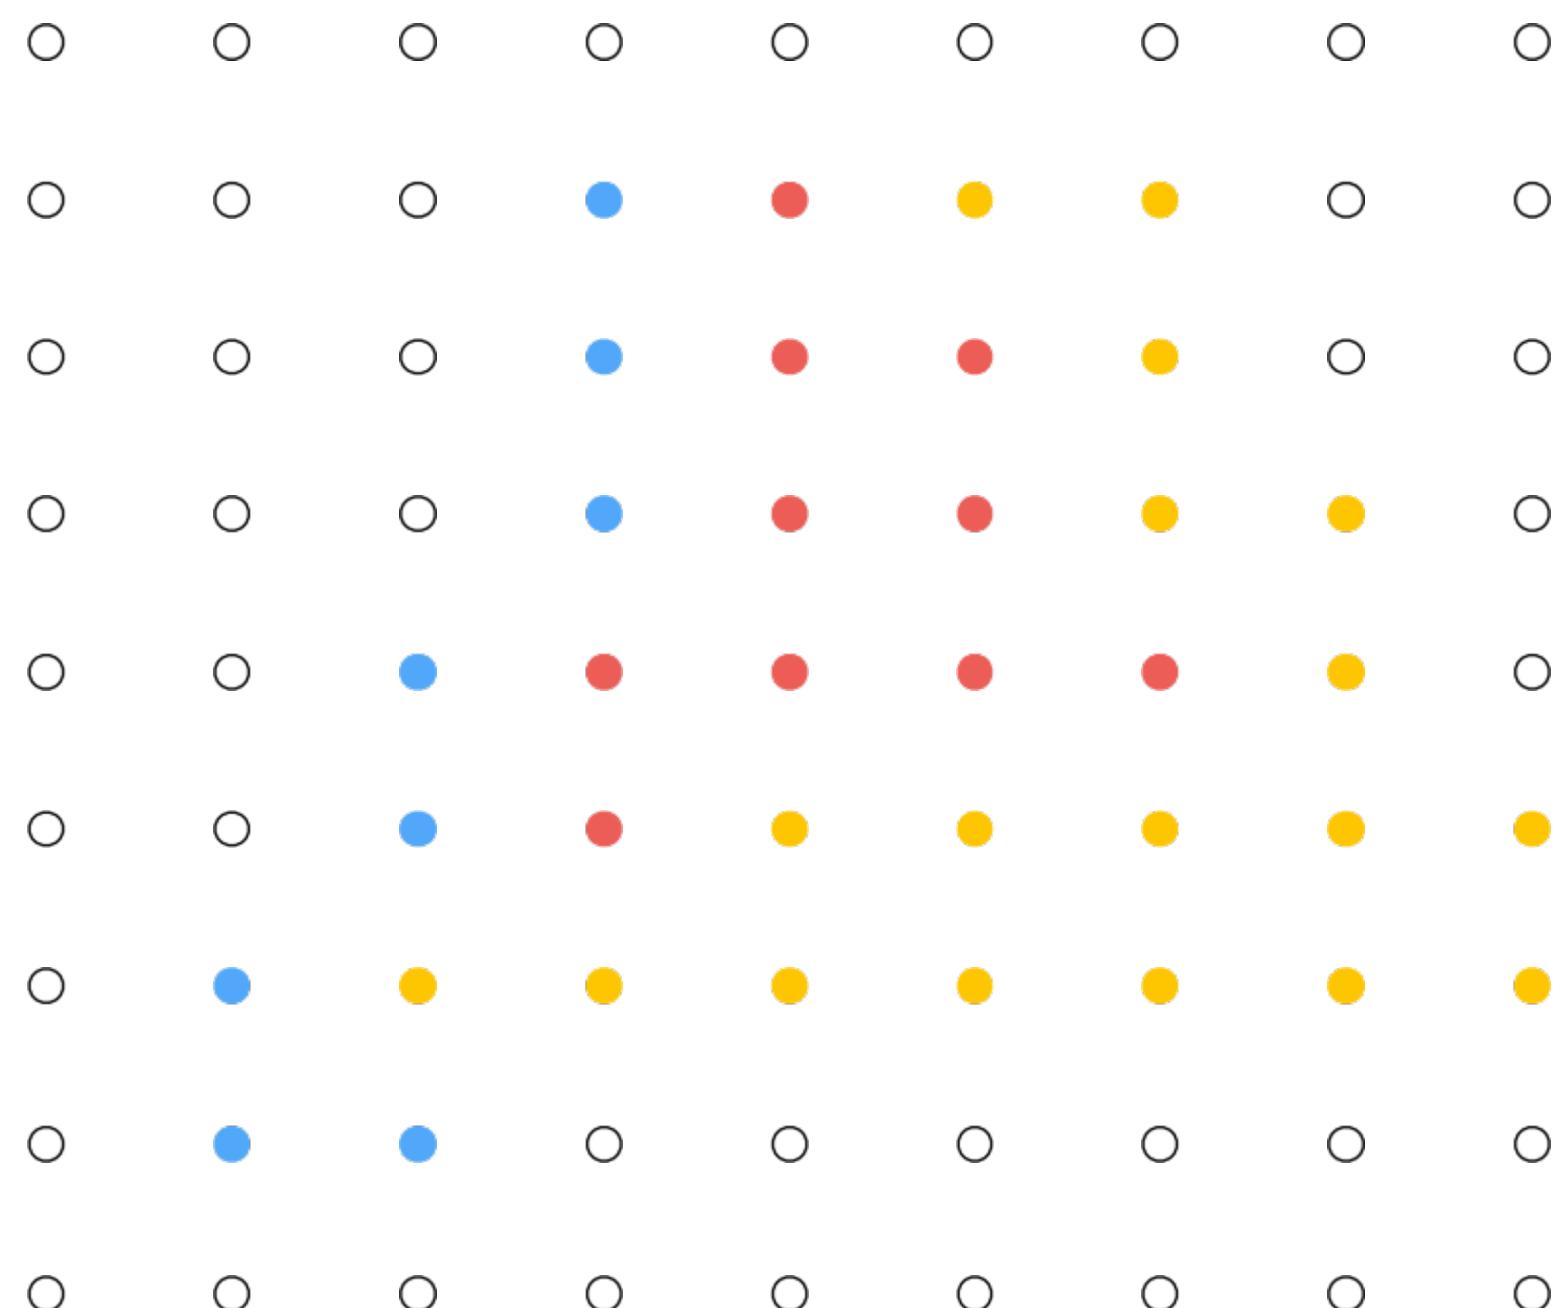
Color buffer contents



Depth buffer contents

Occlusion using the depth-buffer (Z-buffer)

After processing red triangle:



Color buffer contents



Depth buffer contents

Occlusion using the depth buffer

```
bool pass_depth_test(d1, d2)
{
    return d1 < d2;
}
```

```
draw_sample(x, y, d, c) //new depth d & color c at (x,y)
{
    if( pass_depth_test( d, zbuffer[x][y] ))
    {
        // triangle is closest object seen so far at this
        // sample point. Update depth and color buffers.
        zbuffer[x][y] = d;    // update zbuffer
        color[x][y] = c;     // update color buffer
    }
    // otherwise, we've seen something closer already;
    // don't update color or depth
}
```

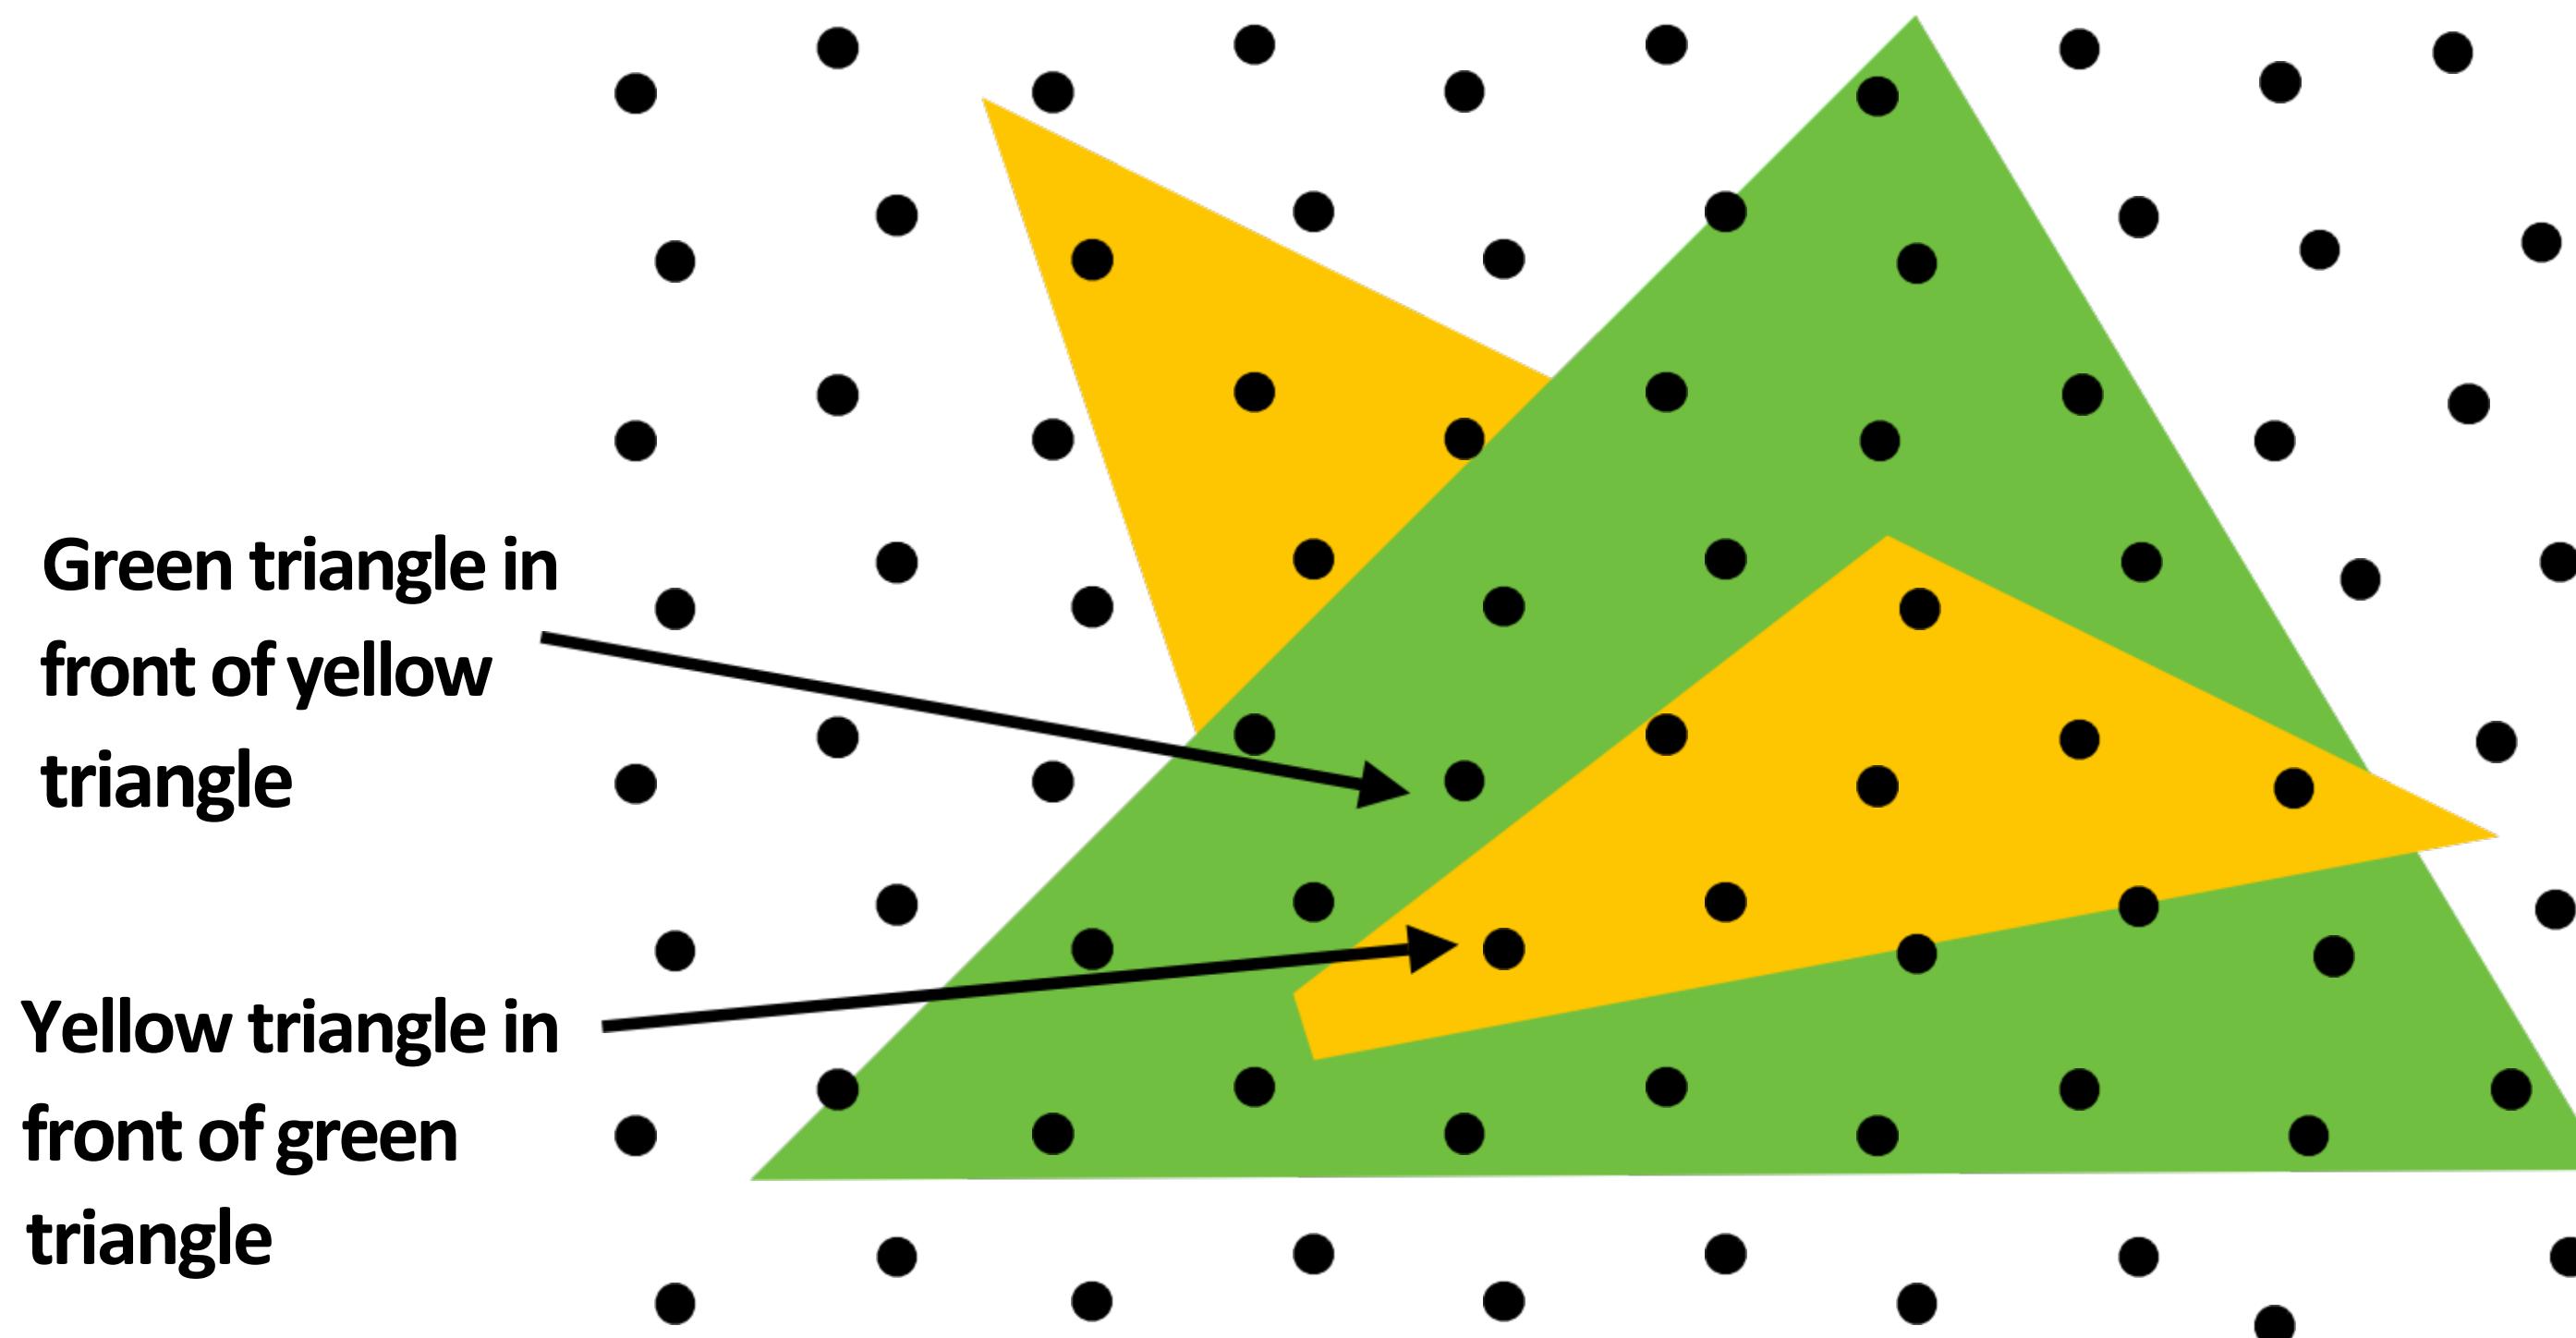
Depth + Intersection

Q: Does depth-buffer algorithm handle interpenetrating surfaces?

A: Of course!

Occlusion test is based on depth of triangles at a given sample point.

Relative depth of triangles maybe different at different sample points.



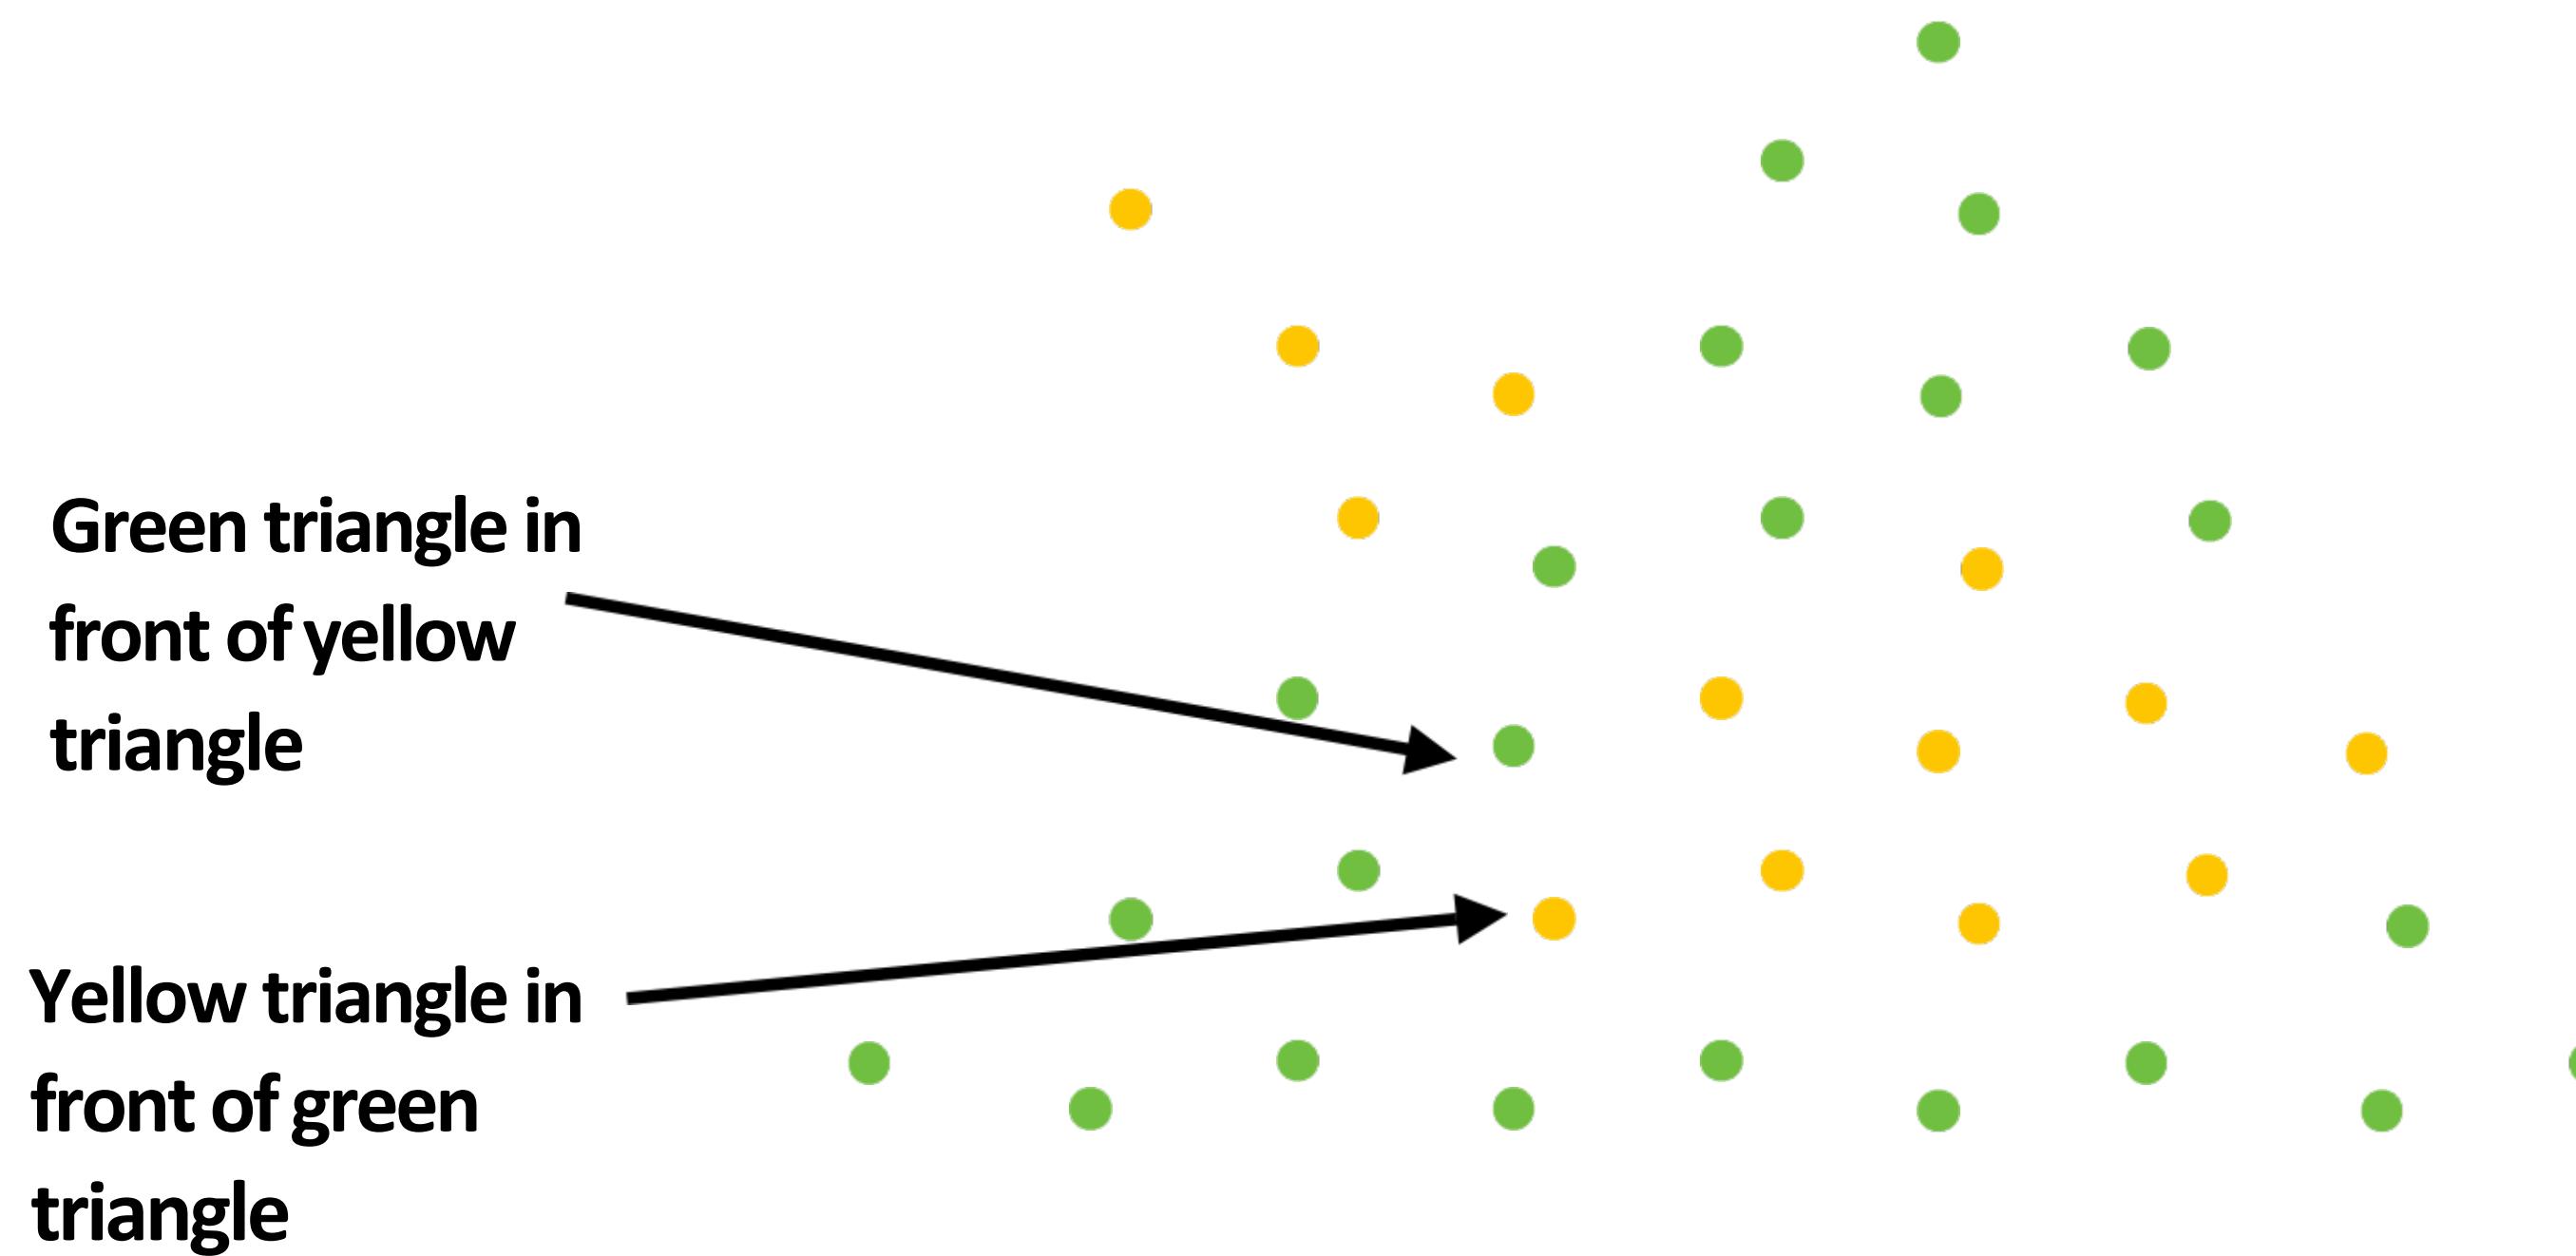
Intersection

Q: Does depth-buffer algorithm handle interpenetrating surfaces?

A: Of course!

Occlusion test is based on depth of triangles at a given sample point.

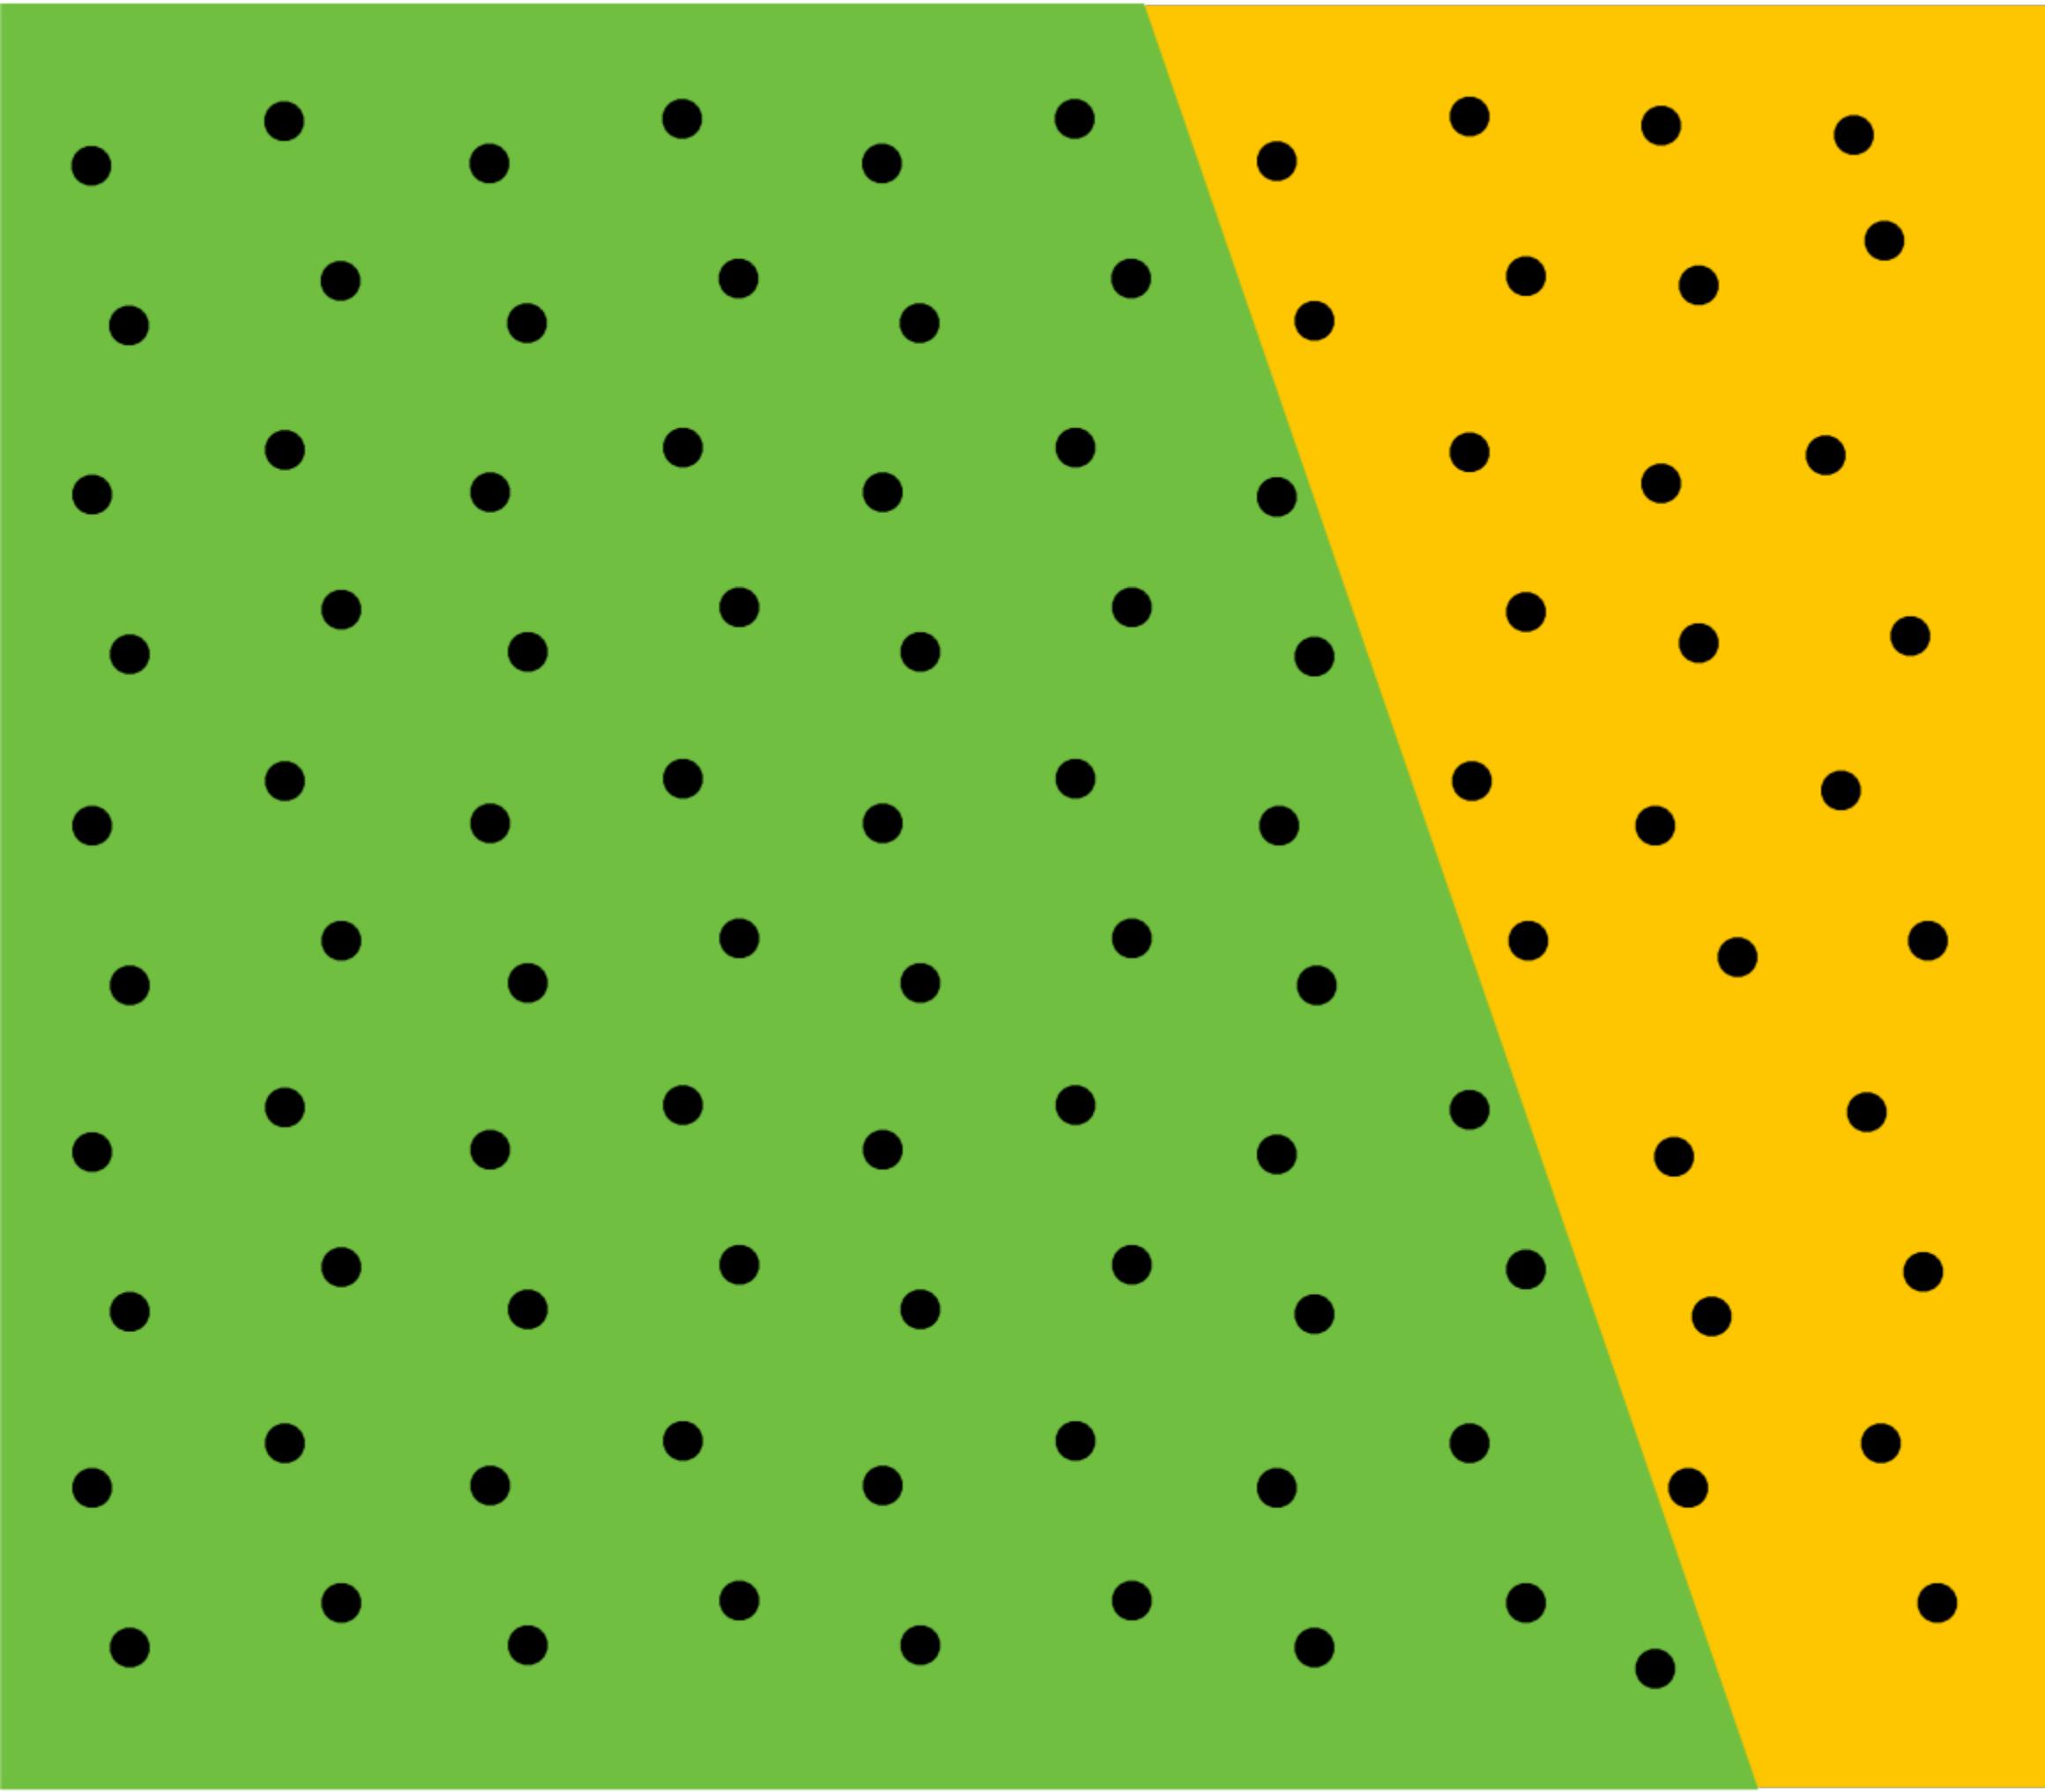
Relative depth of triangles maybe different at different sample points.



Depth + Supersampling

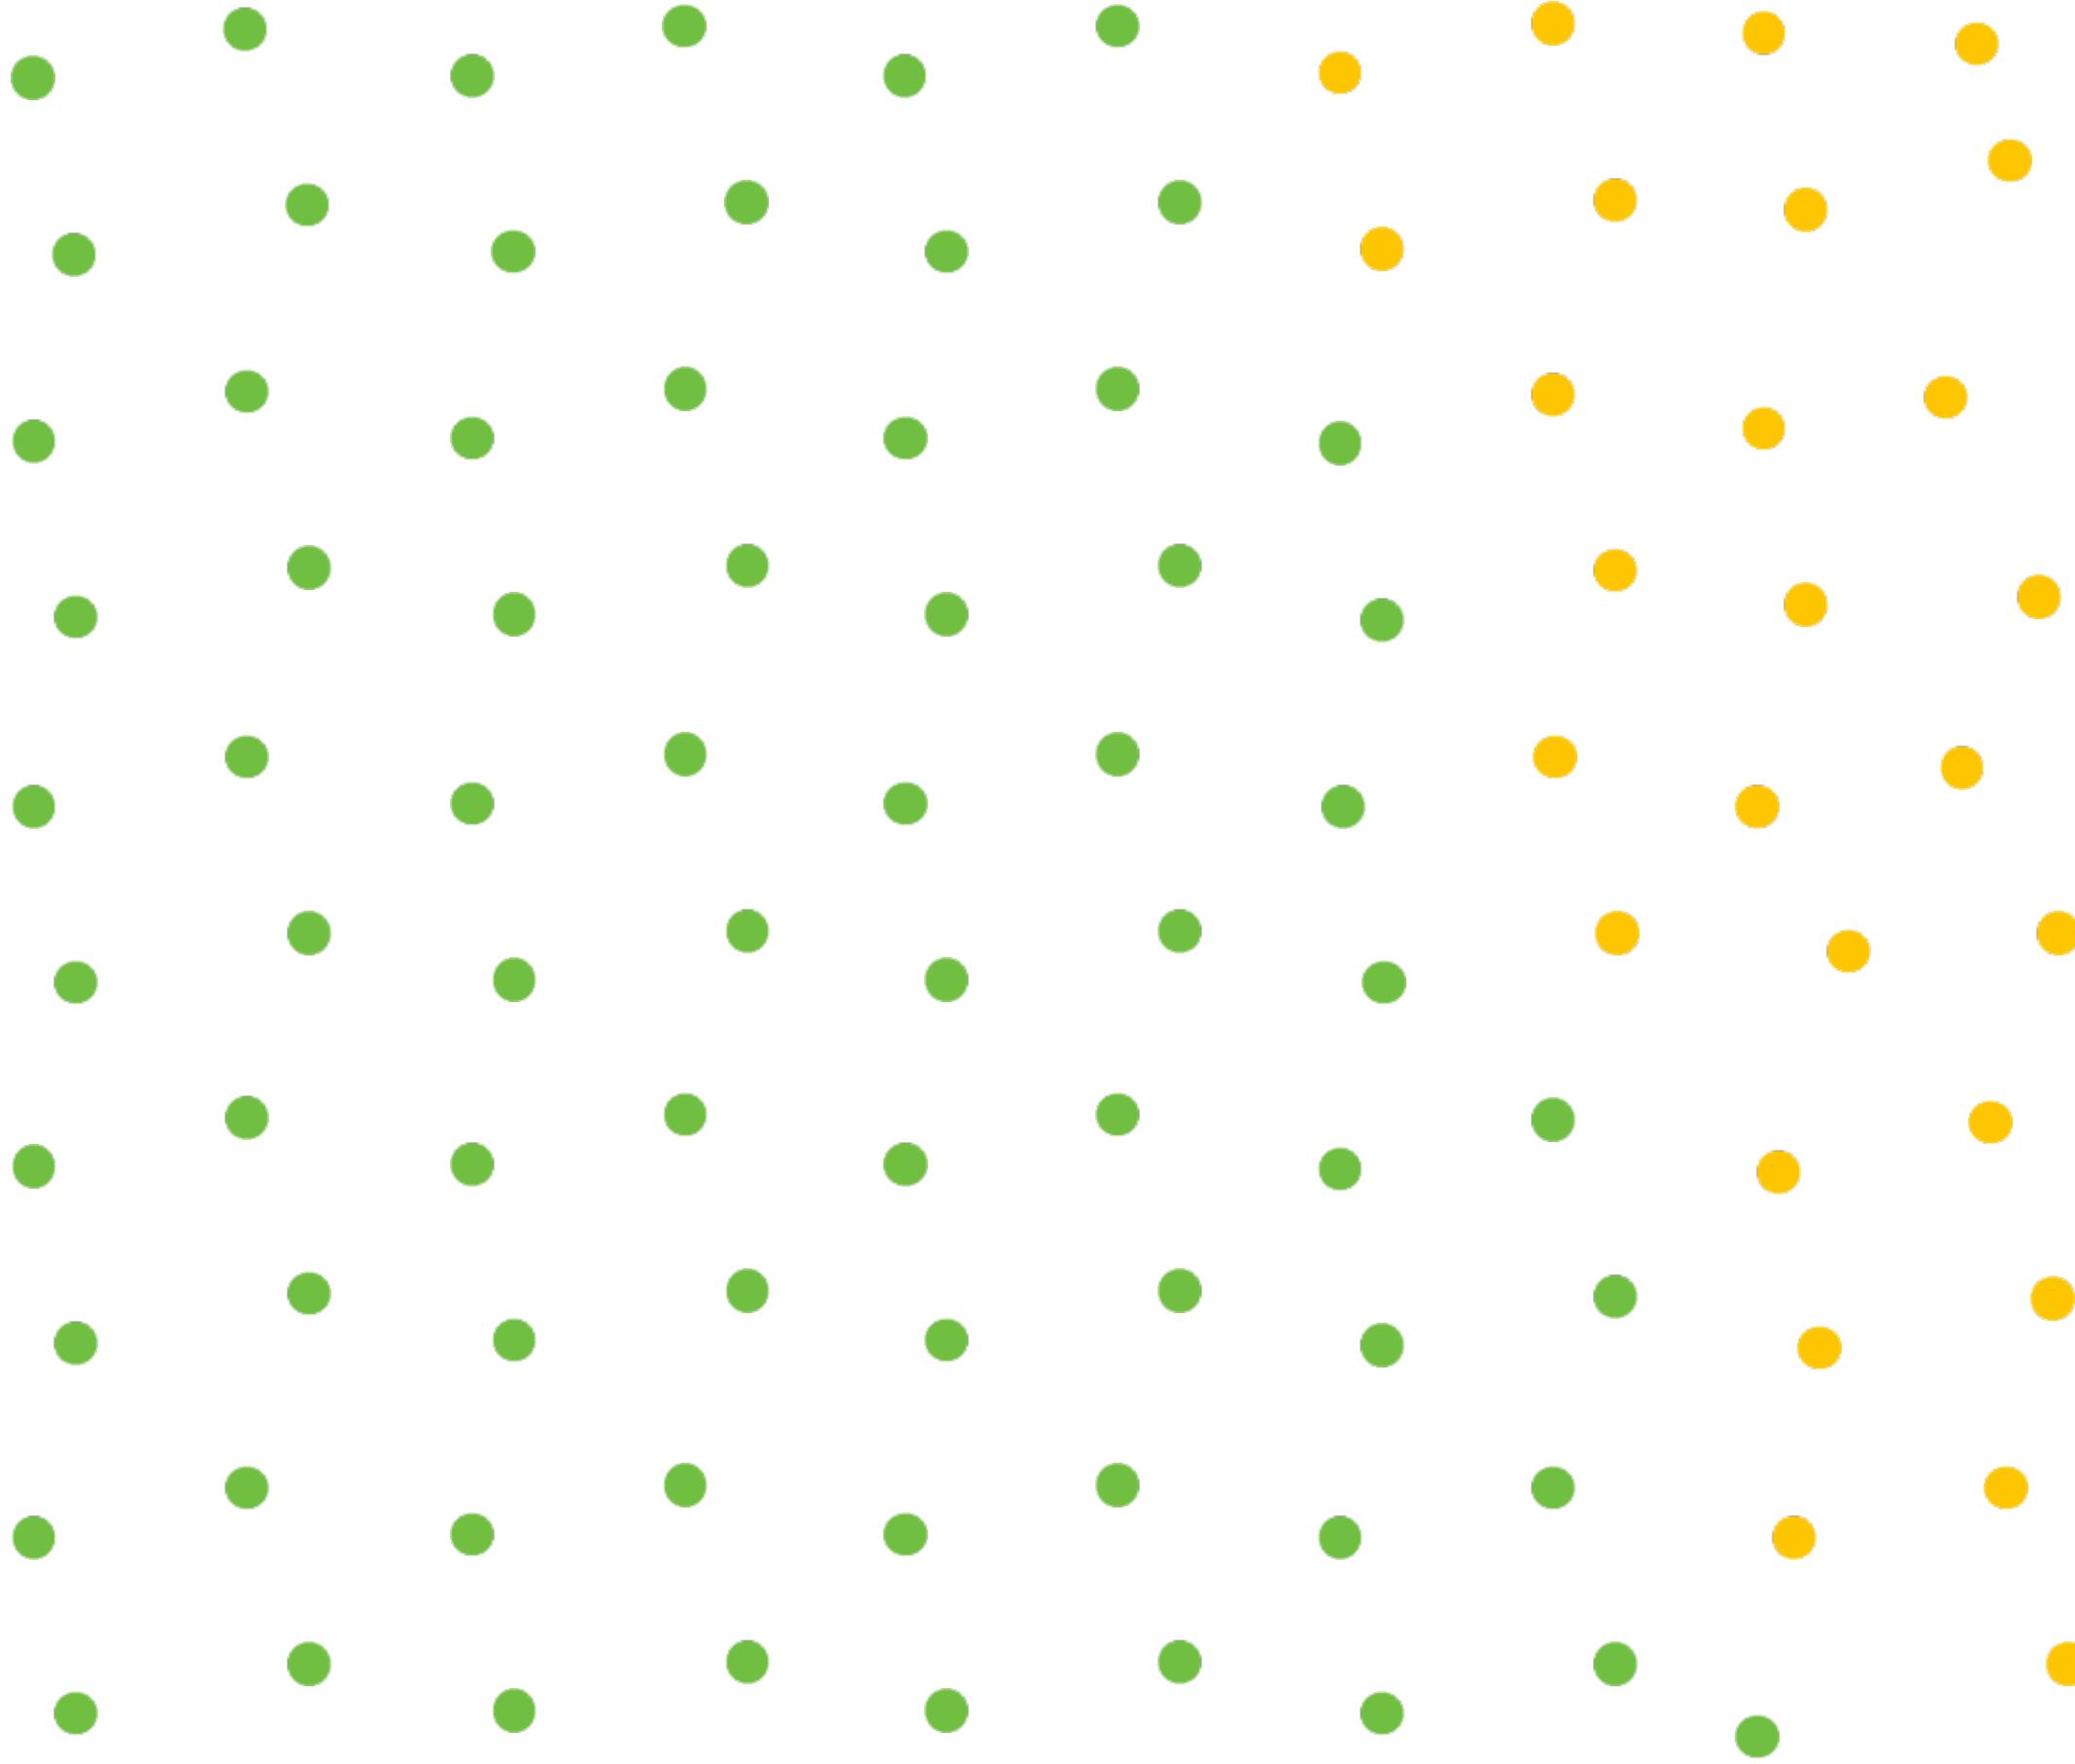
Q: Does depth buffer work with super sampling?

A: Yes! If done per (super) sample.



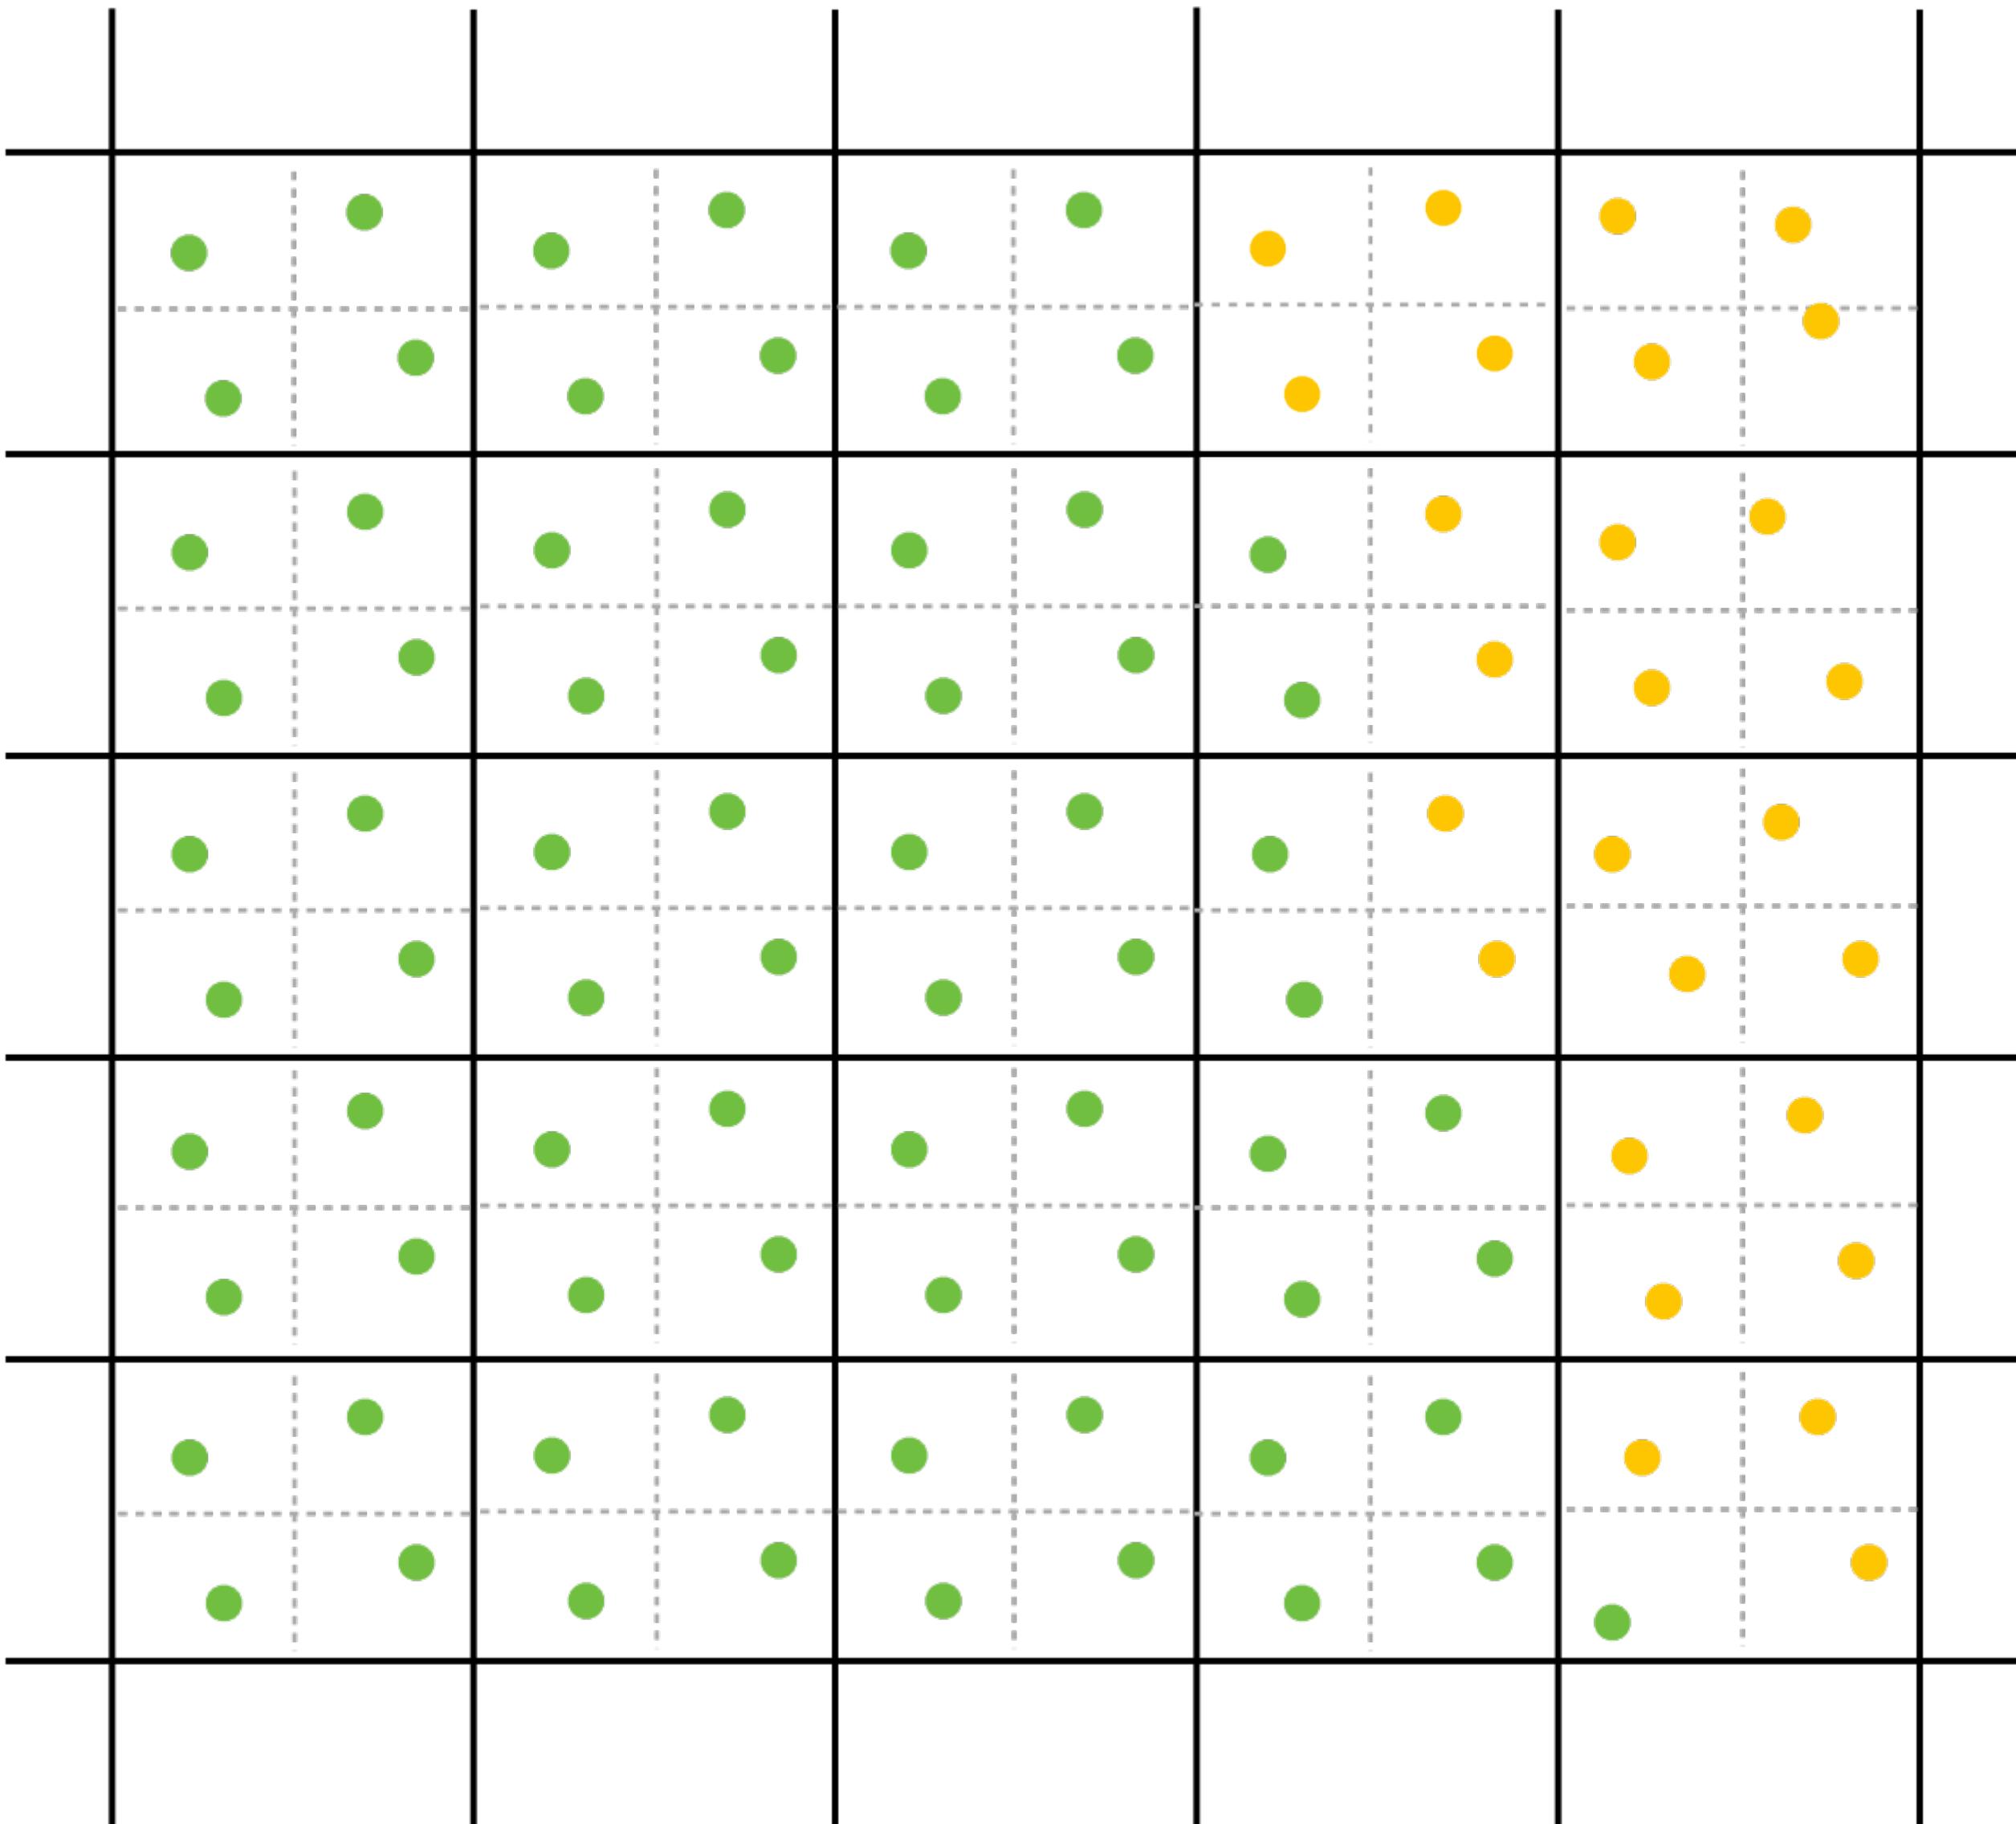
(Here: green triangle occludes yellow triangle)

Depth + Supersampling

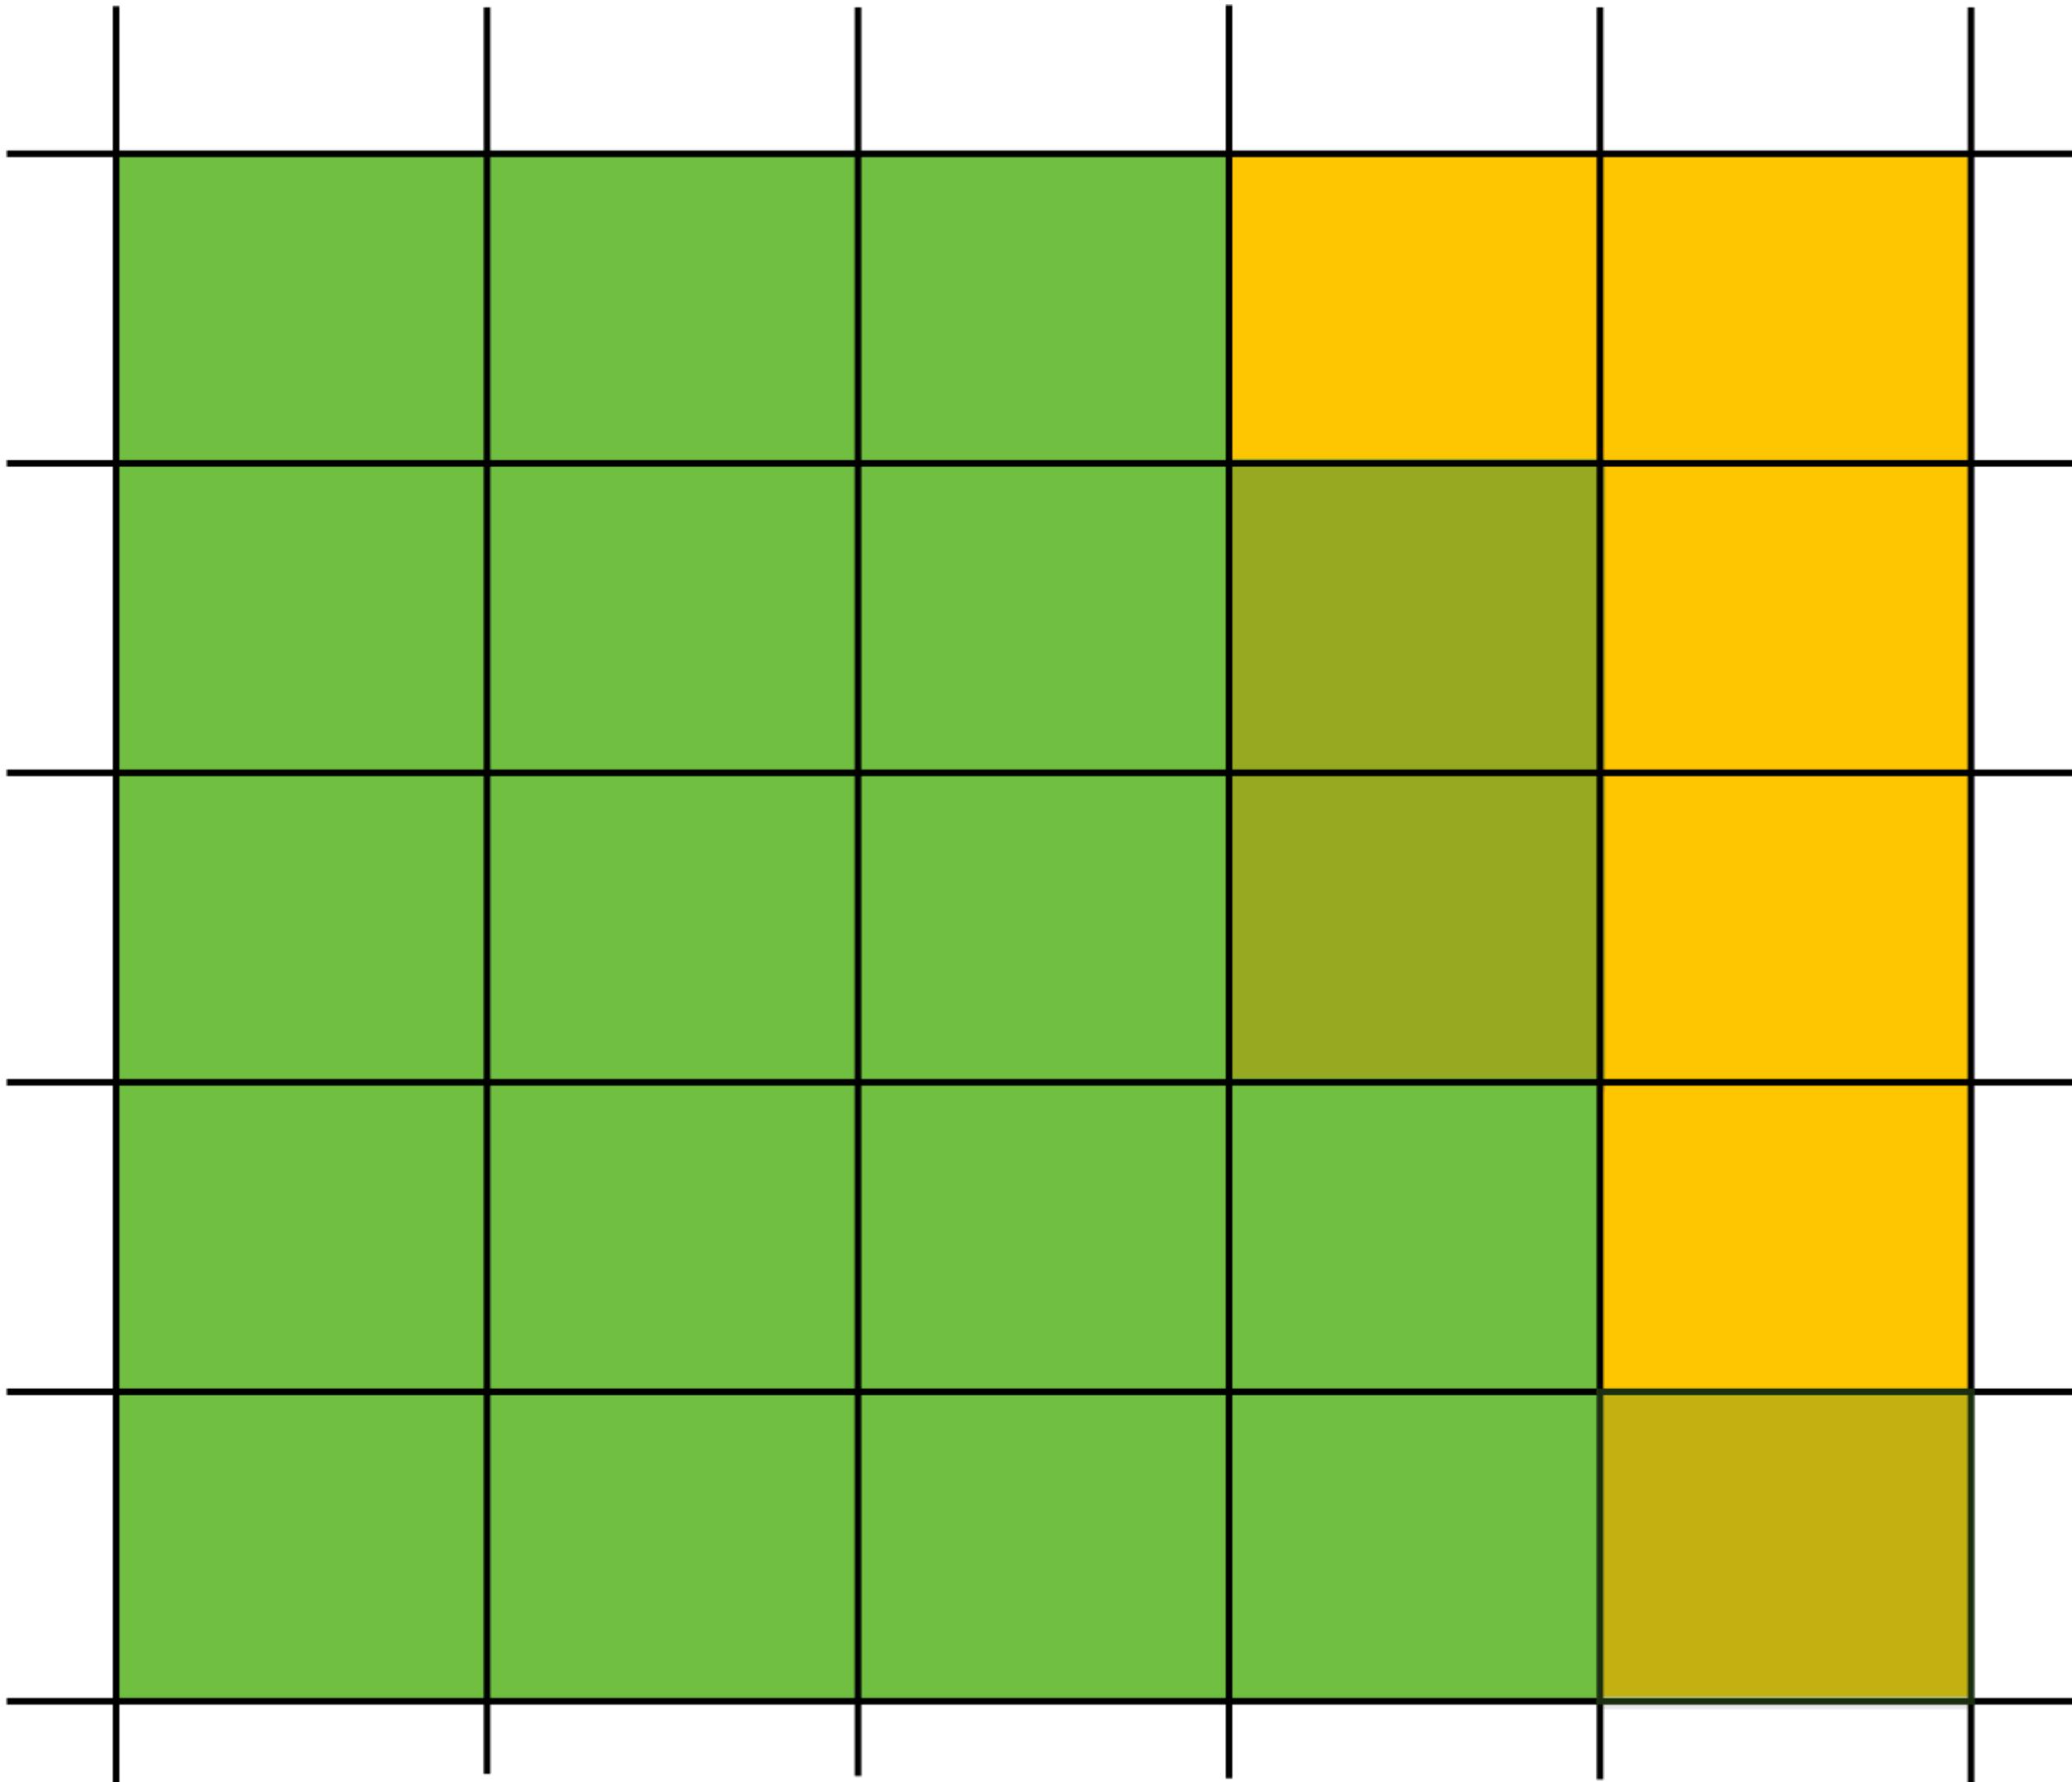


Color of super samples after rasterizing w/ depth buffer

Color buffer contents (4 samples per pixel)



Final resampled result



Note anti-aliasing of edge due to filtering of green and yellow samples

Summary: occlusion using a depth buffer

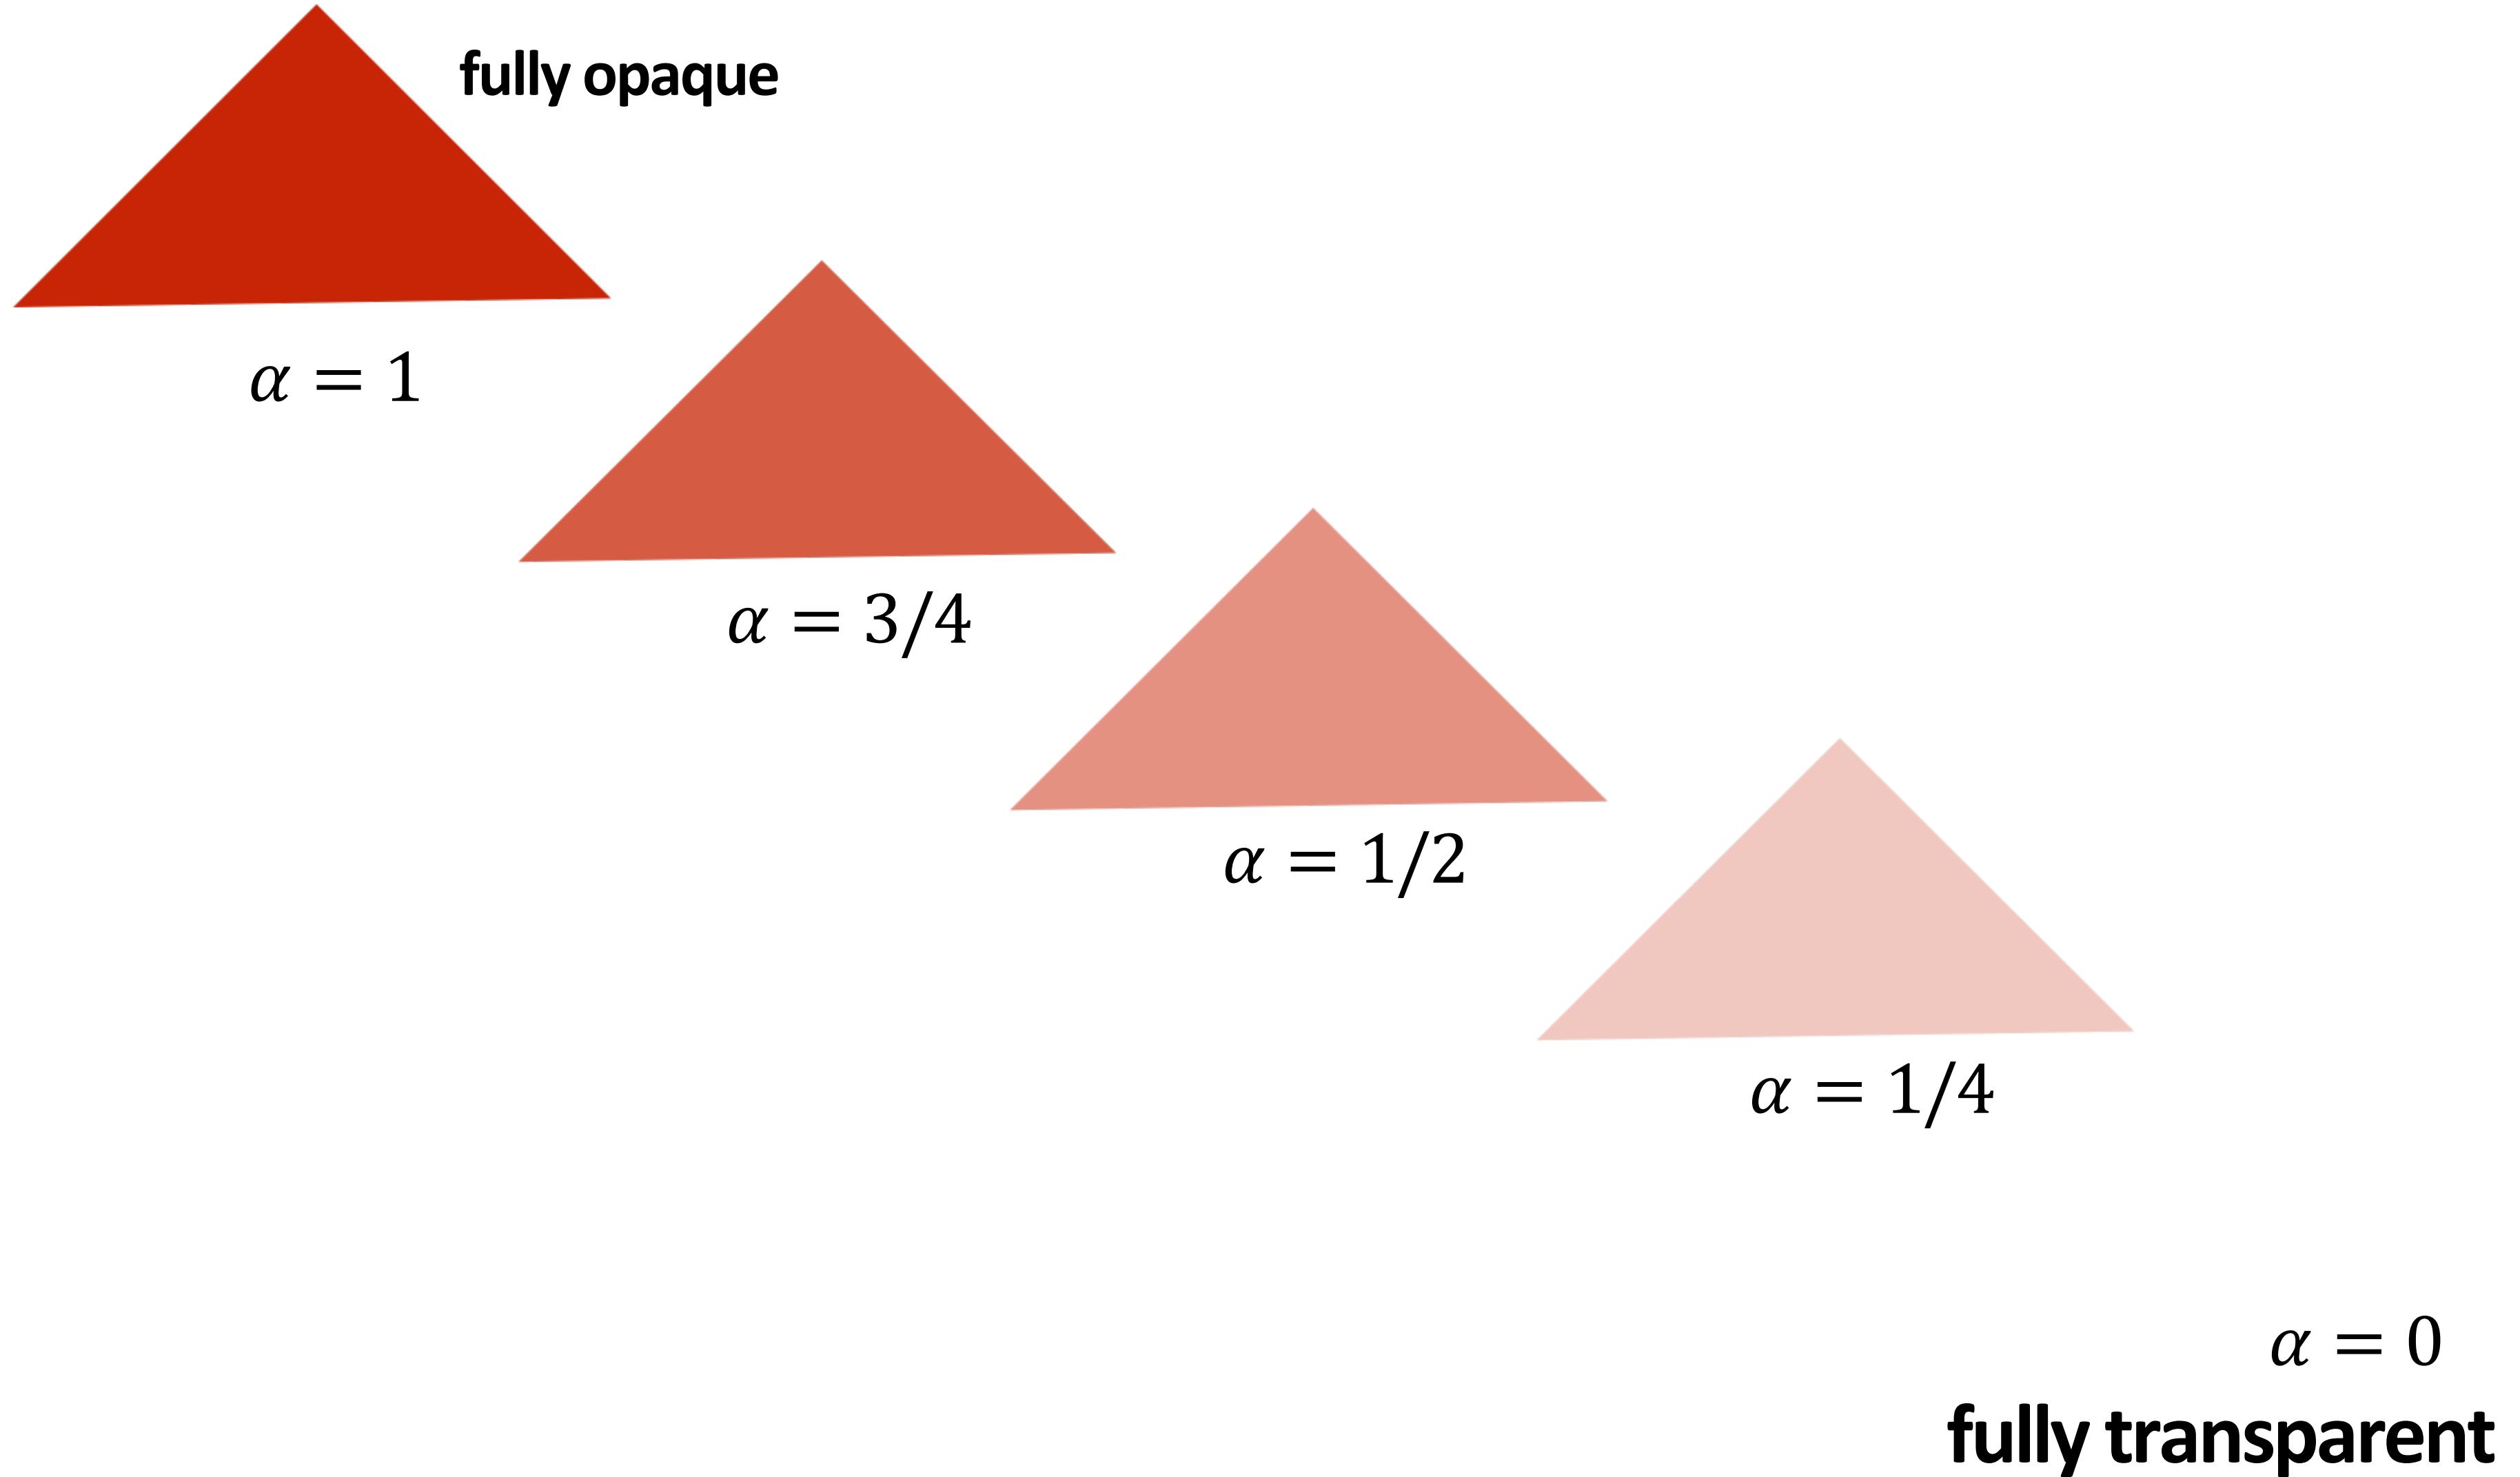
- Store one depth value per (super) sample—not one per pixel!
- Constant additional space per sample
 - Hence, constant space for depth buffer
 - Doesn't depend on number of overlapping primitives!
- Constant time occlusion test per covered sample
 - Read-modify write of depth buffer if “pass” depth test
 - Just a read if “fail”
- Not specific to triangles: only requires that surface depth can be evaluated at a screen sample point

But what about semi-transparent surfaces?

Compositing

Representing opacity as alpha

An “alpha” value $0 \leq \alpha \leq 1$ describes the opacity of an object

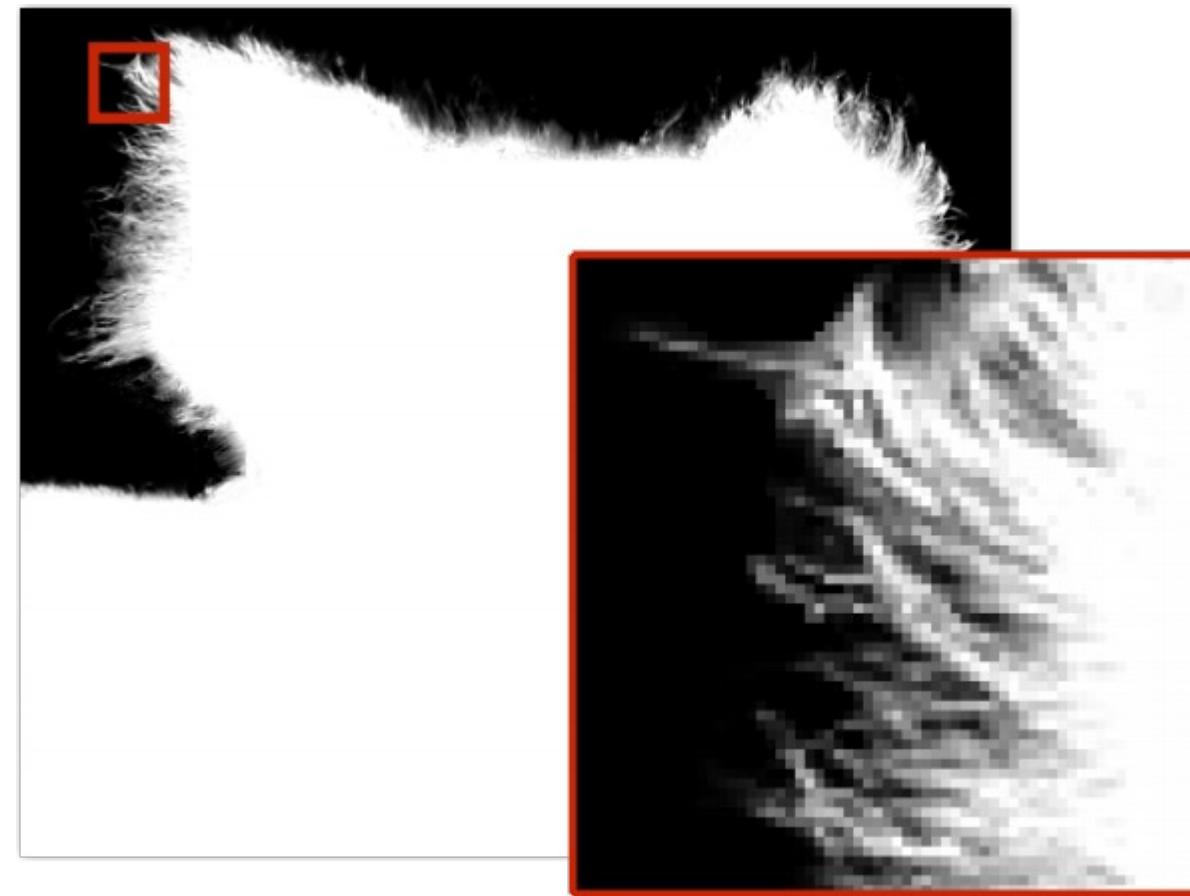


Alpha channel of an image

color channels



α channel



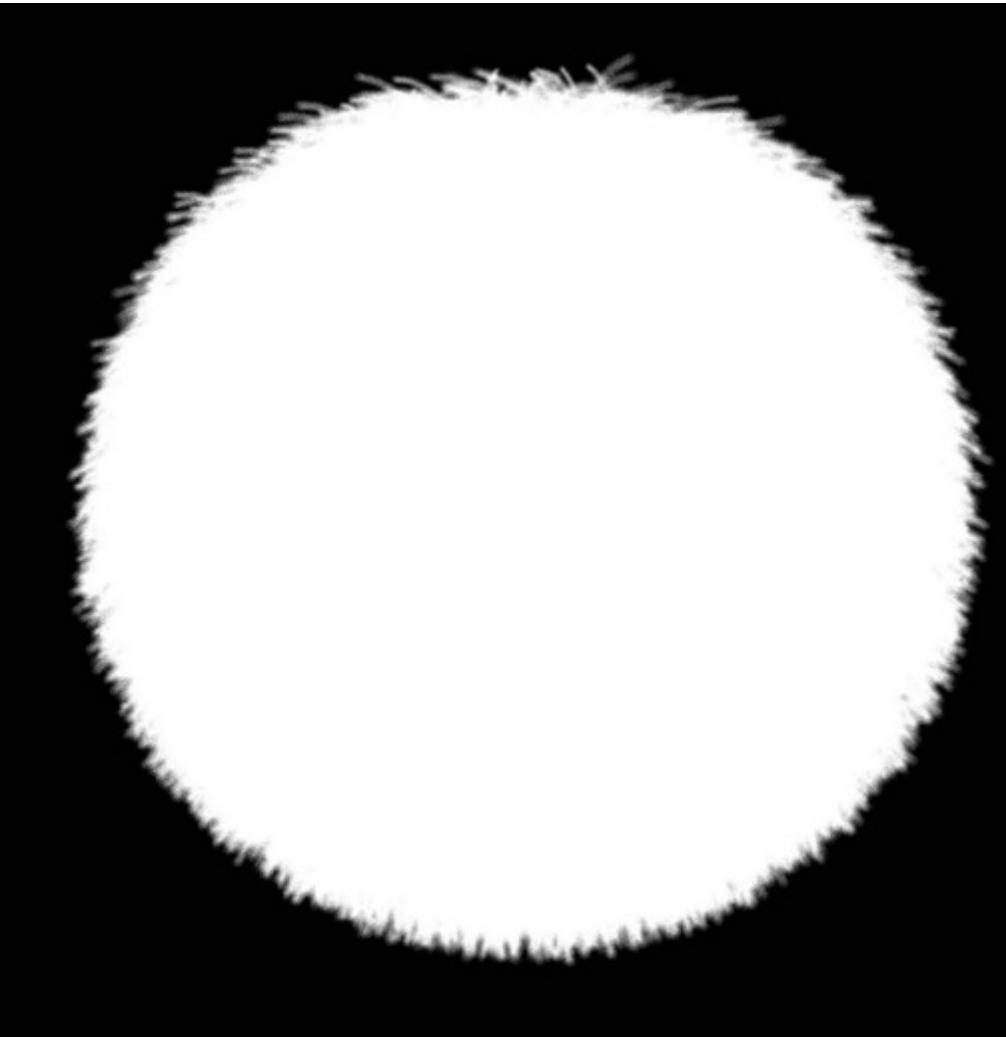
Key idea: can use α channel to composite one image on top of another.

Fringing

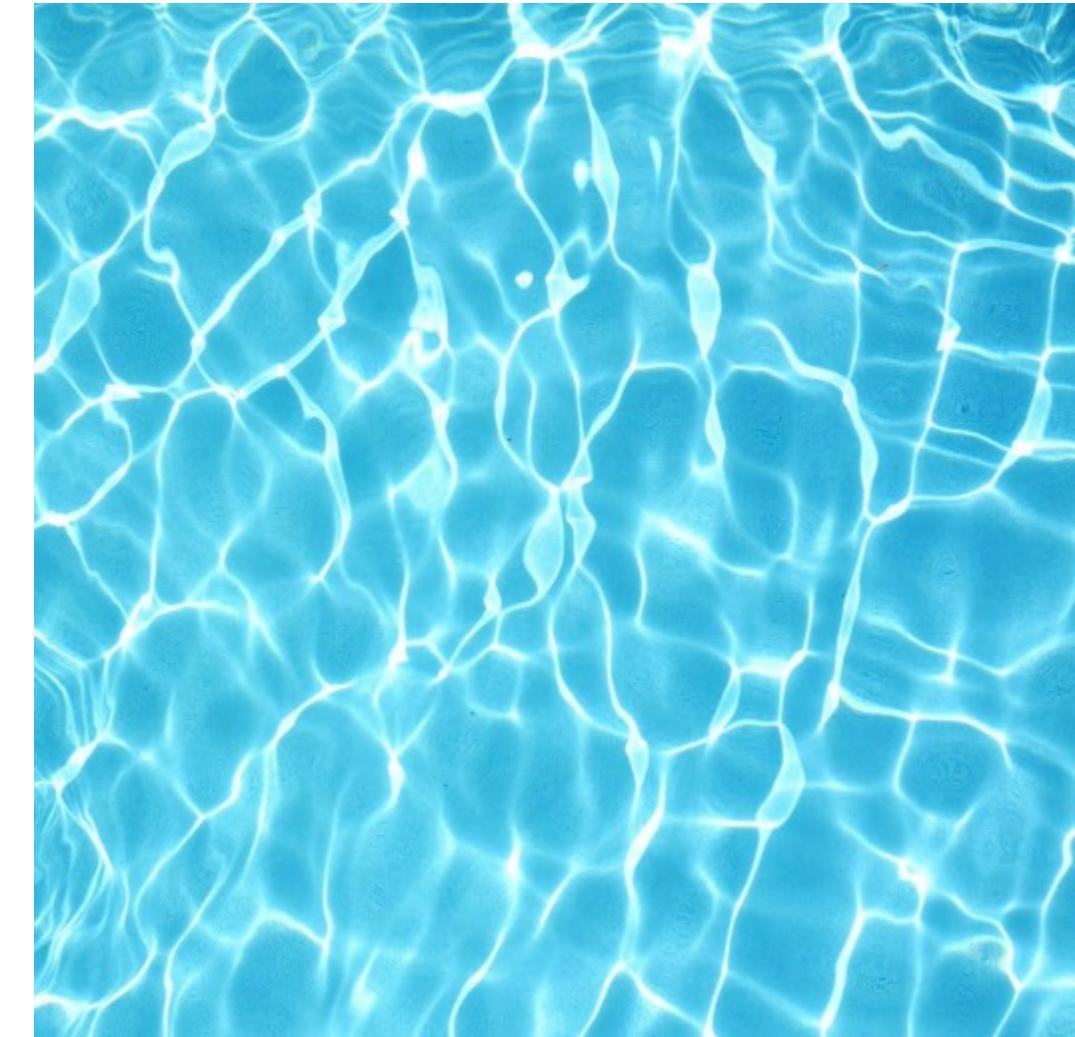
Poor treatment of color/alpha can yield dark “fringing”:



foreground color



foreground alpha



background color

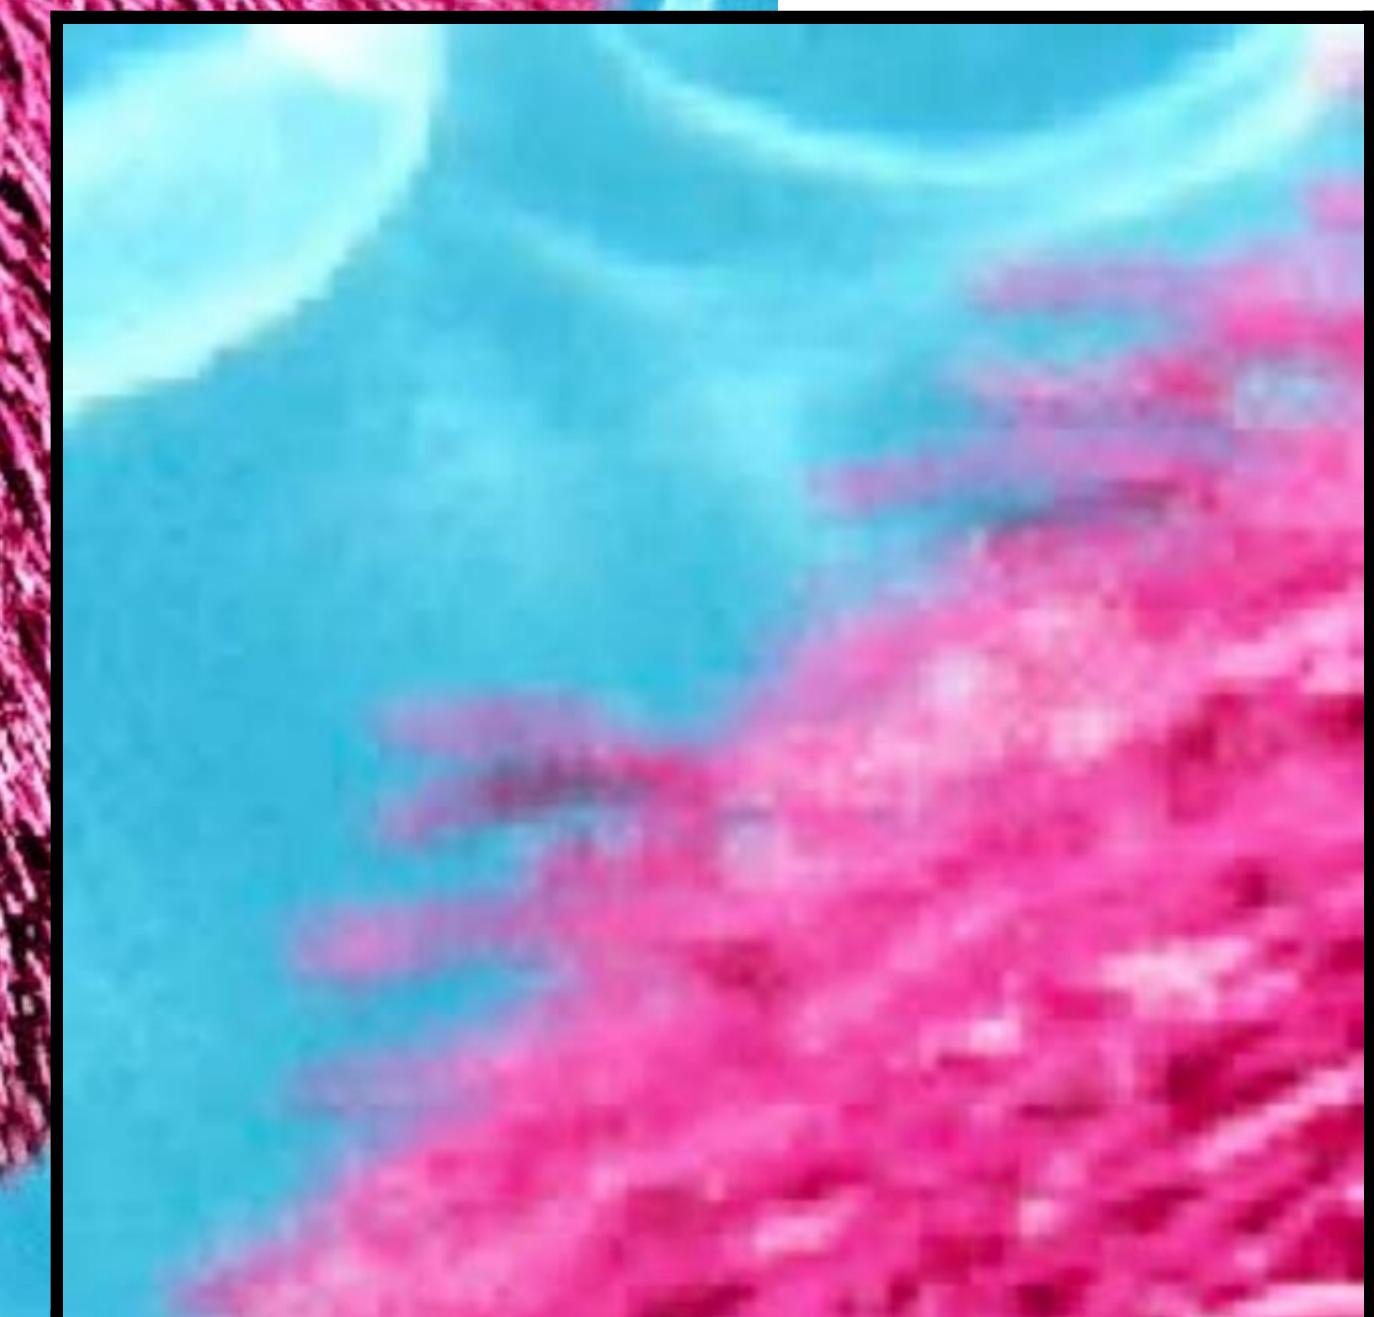


fringing

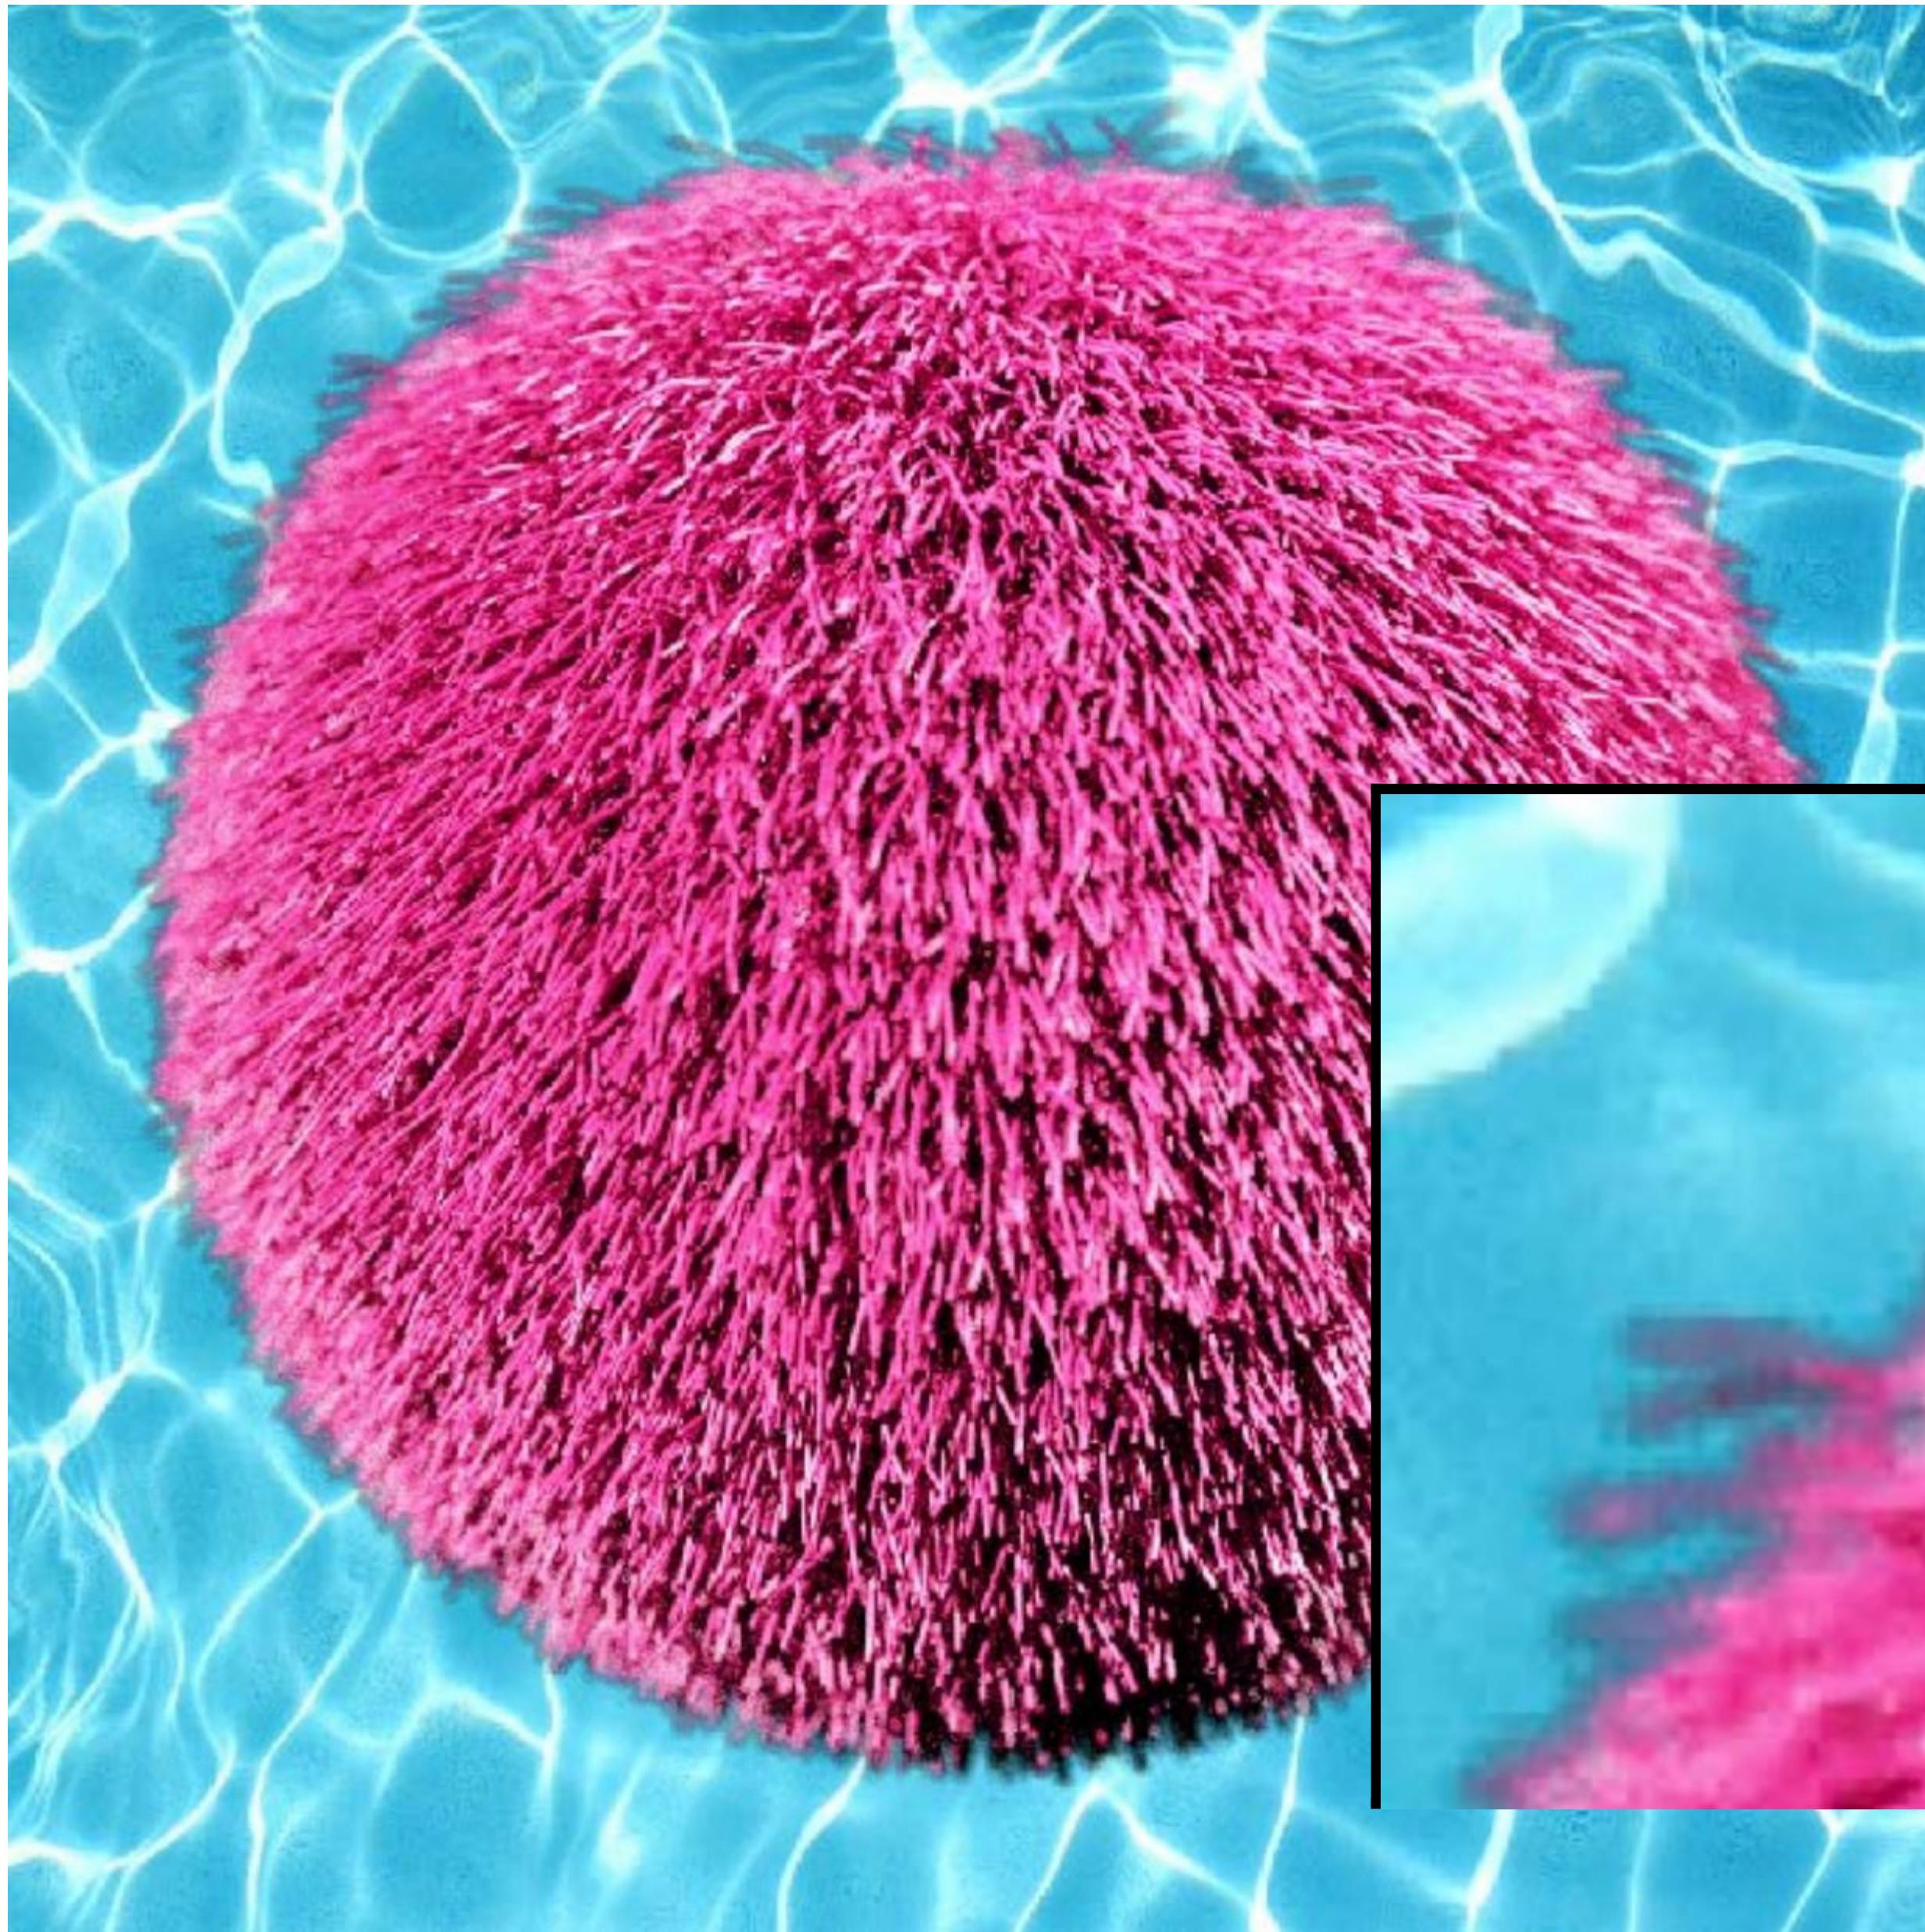


no fringing

No fringing



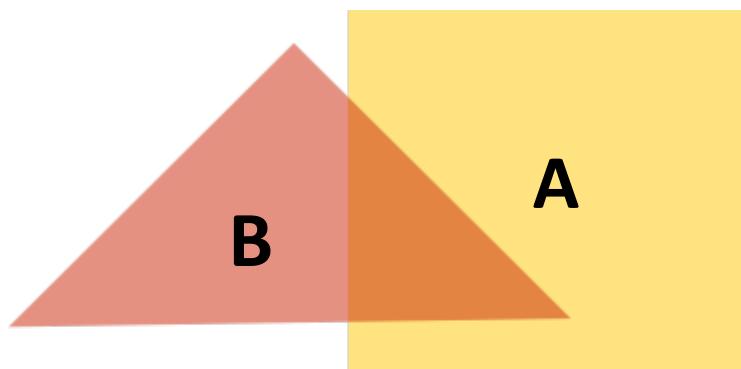
Fringing (...why does this happen?)



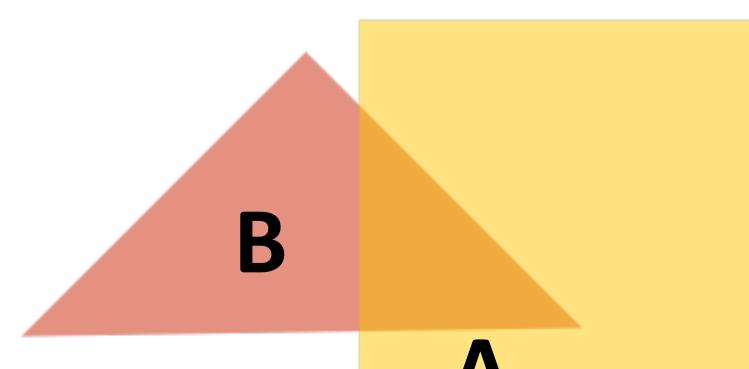
Over operator:

Composites image B with opacity α_B over image A with opacity α_A

Informally, captures behavior of “tinted glass”



B over A



A over B

Notice: “over” is not commutative

$$A \text{ over } B \neq B \text{ over } A$$



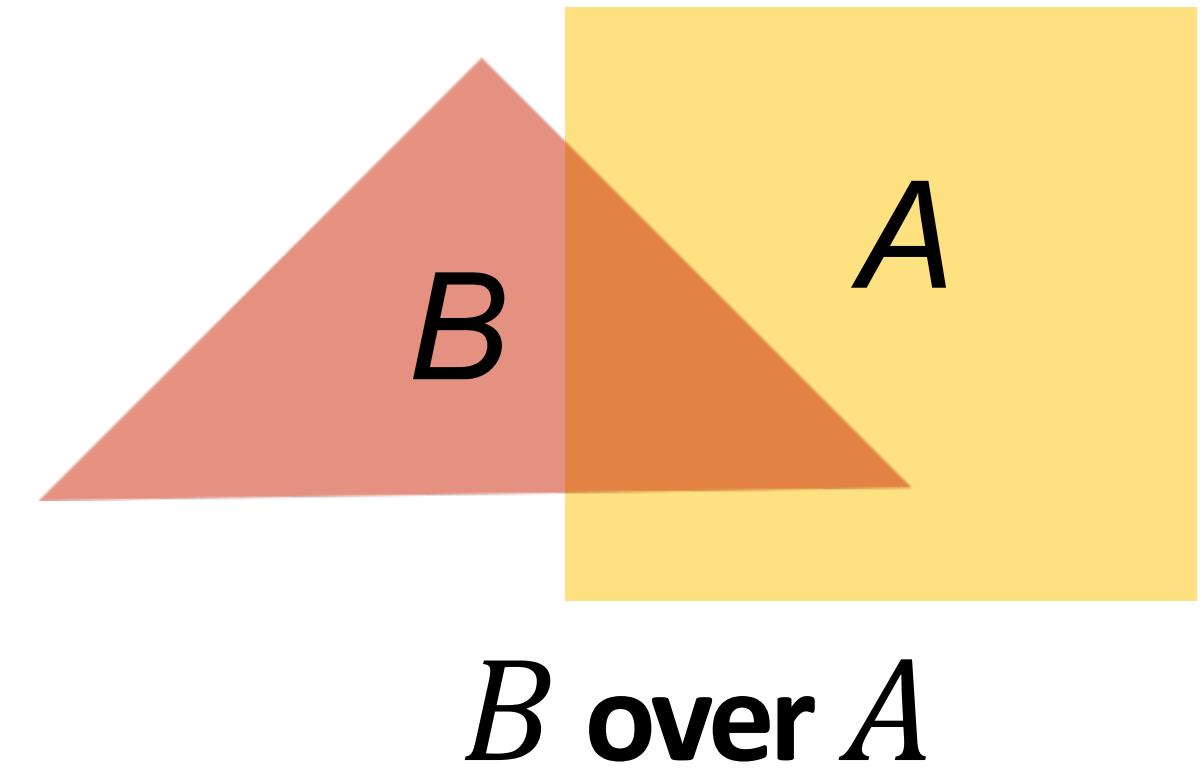
Over operator: non-premultiplied alpha

Composites image B with opacity α_B over image A with opacity α_A

A first attempt:

$$A = (A_r, A_g, A_b)$$

$$B = (B_r, B_g, B_b)$$



Composite color:

$$C = \alpha_B B + (1 - \alpha_B)\alpha_A A$$

↑
appearance of
semi-transparent B

↓
what B lets through
↑
appearance of
semi-transparent A

Composite alpha:

$$\alpha_C = \alpha_B + (1 - \alpha_B)\alpha_A$$

Over operator: premultiplied alpha

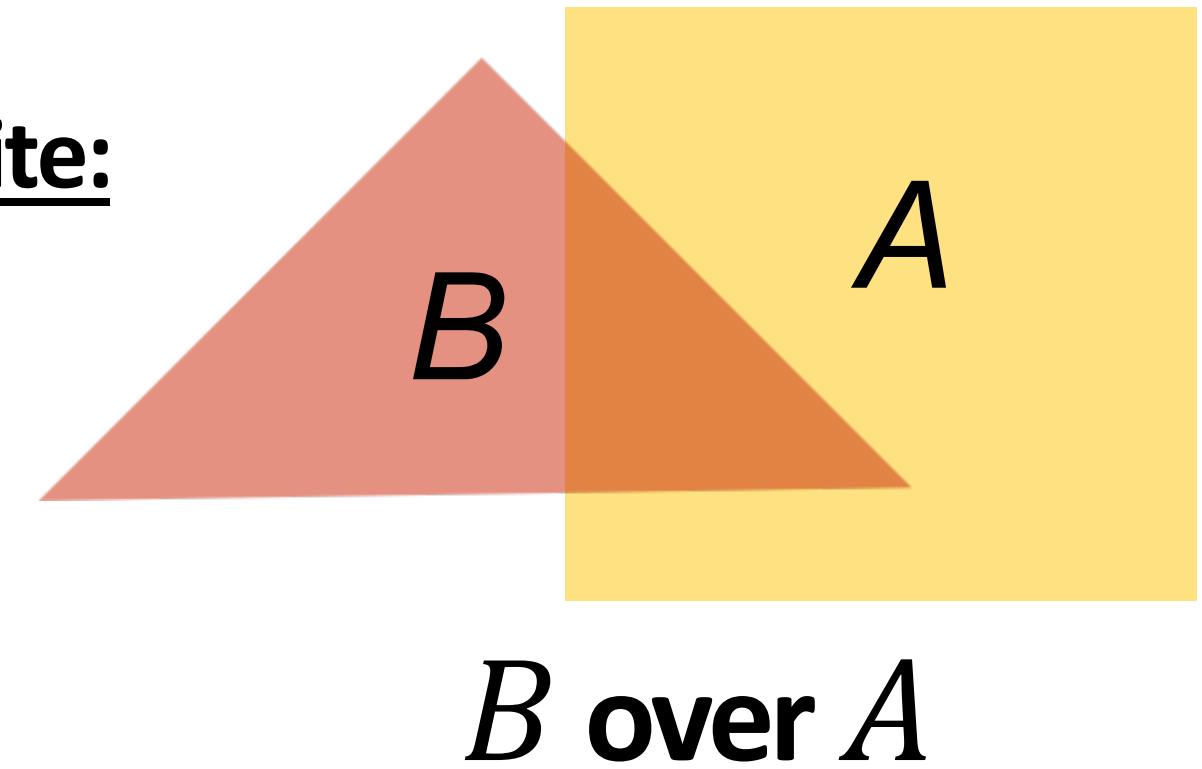
Composites image B with opacity α_B over image A with opacity α_A

Premultiplied alpha-----multiply color by α , then composite:

$$A' = (\alpha_A A_r, \alpha_A A_g, \alpha_A A_b, \alpha_A)$$

$$B' = (\alpha_B B_r, \alpha_B B_g, \alpha_B B_b, \alpha_B)$$

$$C' = B' + (1 - \alpha_B)A'$$

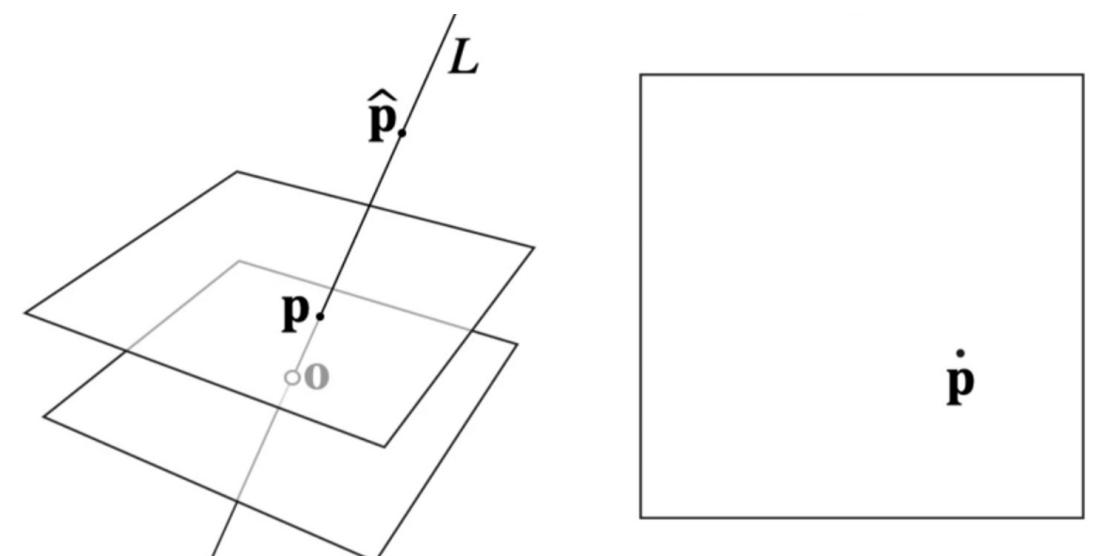


Notice premultiplied alpha composes alpha just like how it composes rgb.
(Non-premultiplied) alpha composes alpha differently than rgb.

“Un-premultiply” to get final color

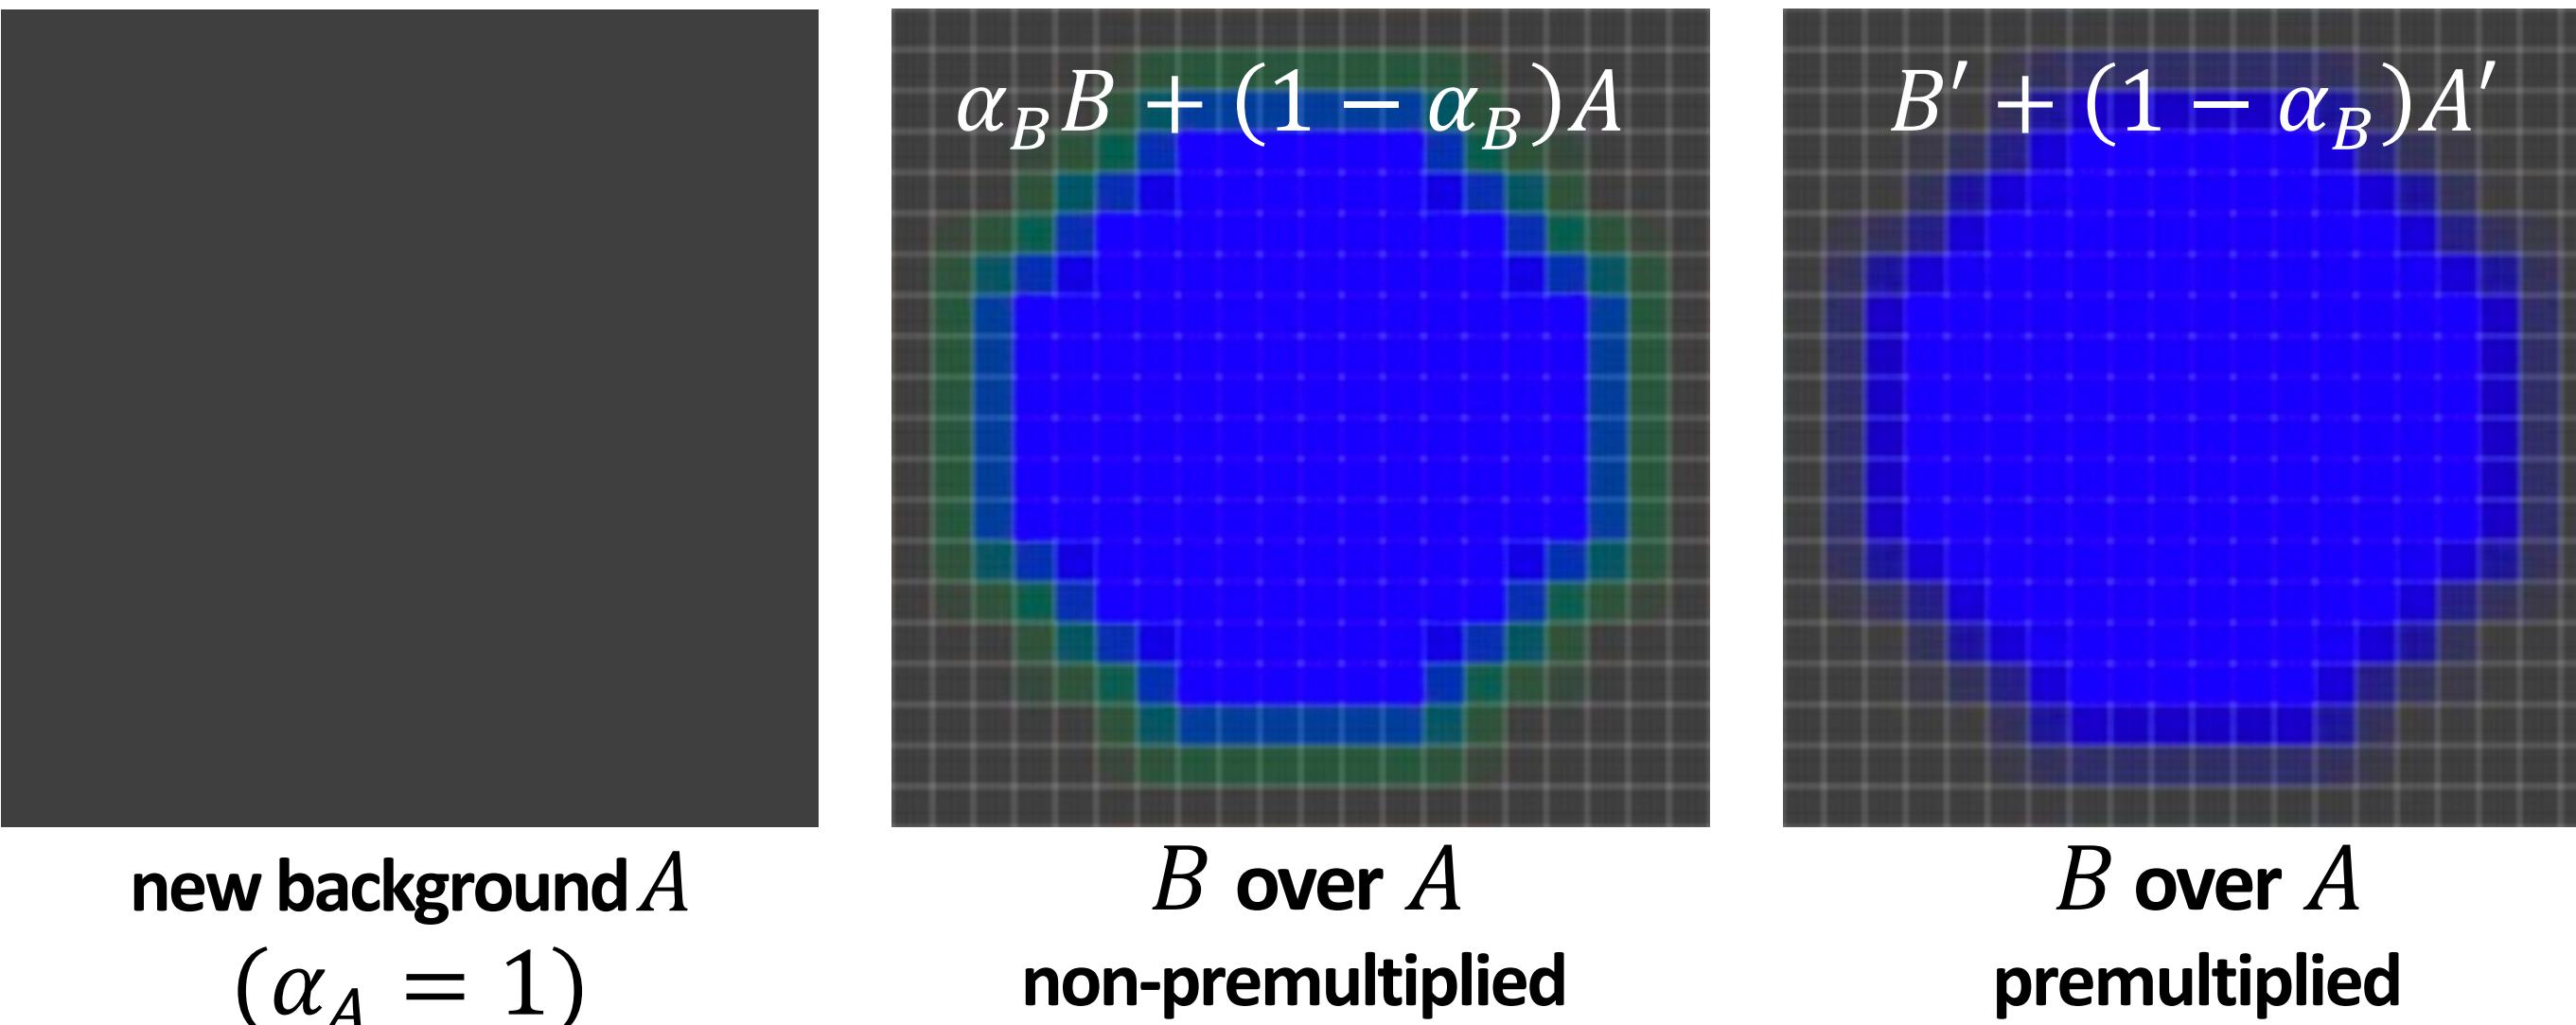
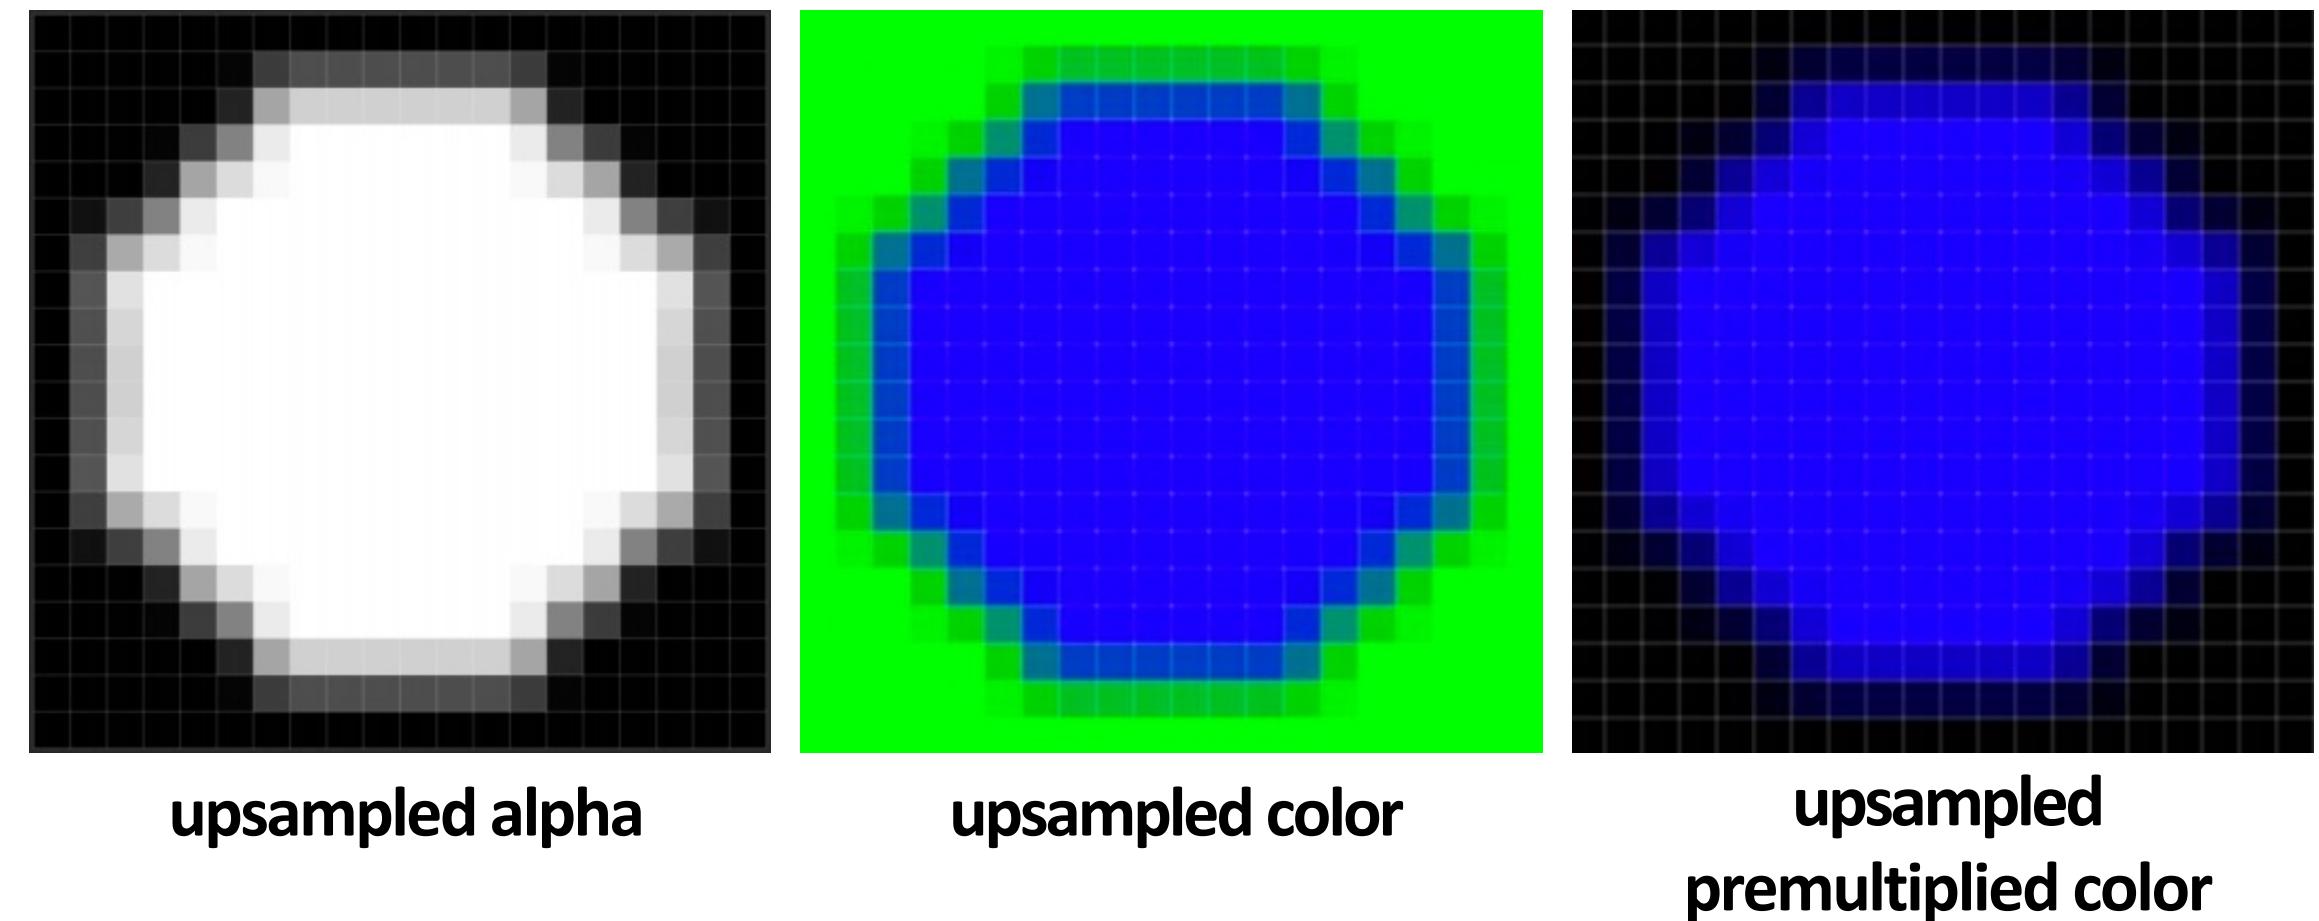
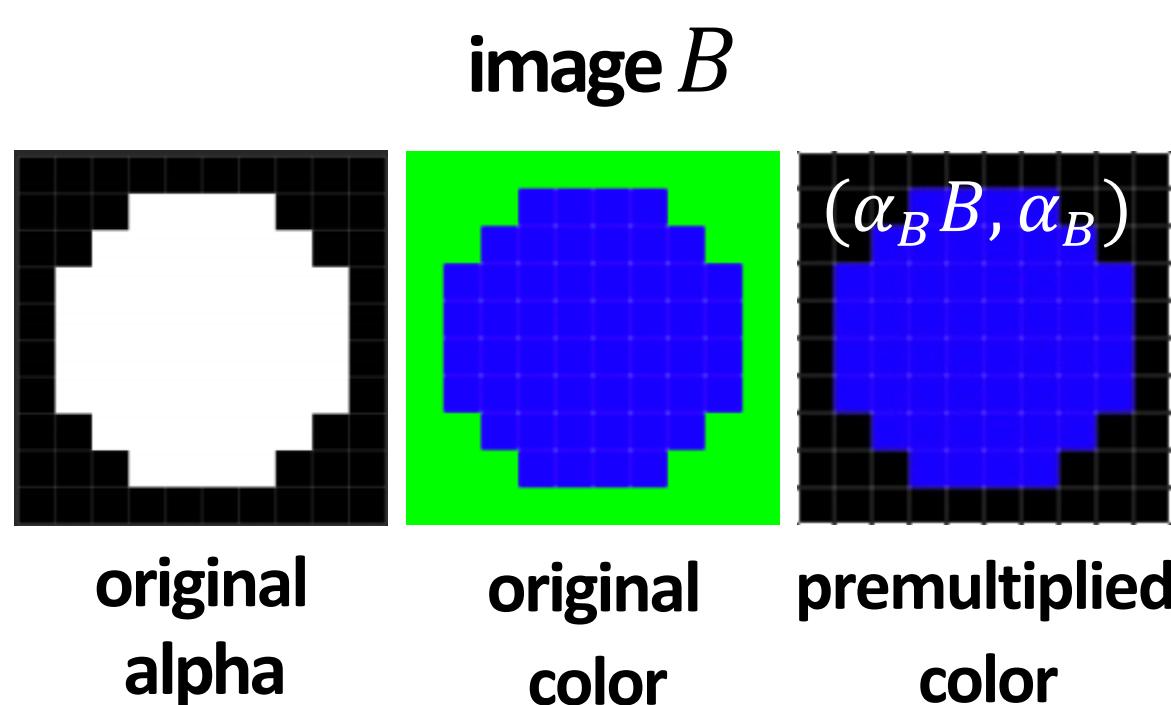
$$(C_r, C_g, C_b, \alpha_C) \rightarrow (C_r / \alpha_C, C_g / \alpha_C, C_b / \alpha_C)$$

Q: Does this division remind you of anything



Compositing with & without premultiplied α

Suppose we upsample an image w/ an α channel, then composite it onto a background:



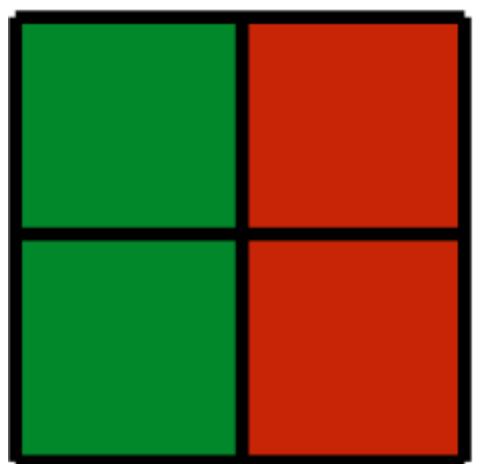
Q: Why do we get the “green fringe” when we don’t premultiply?

Similar problem with non-premultiplied α

Consider pre-filtering (downsampling) a texture with an alpha matte



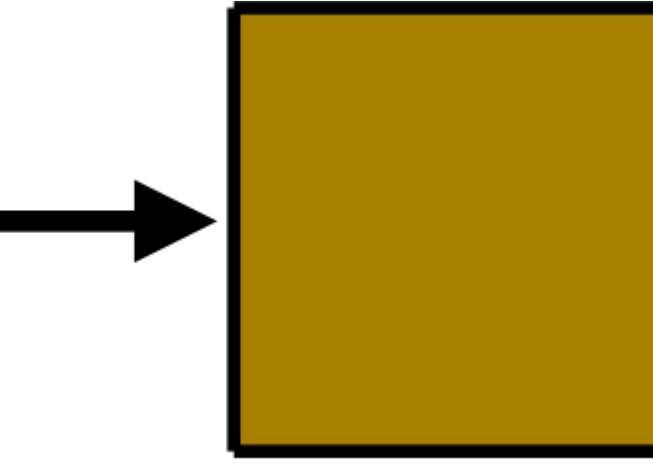
desired downsampled result



input color



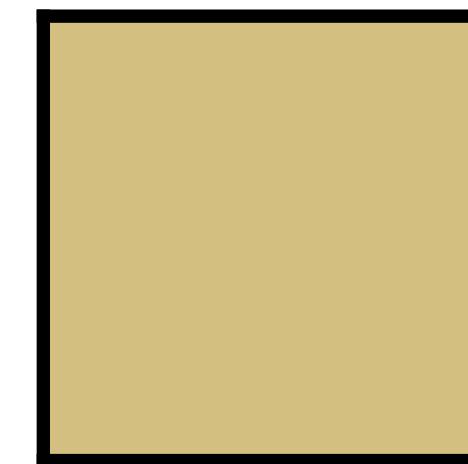
input α



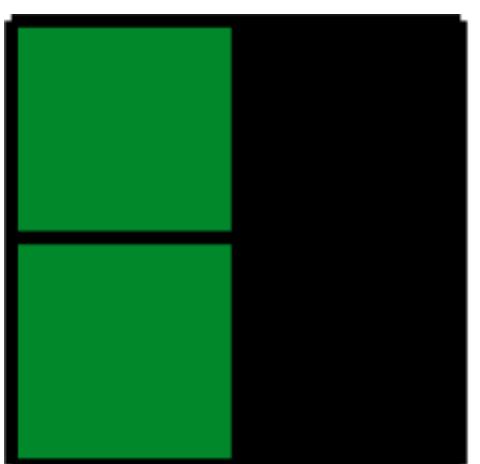
filtered color



filtered α



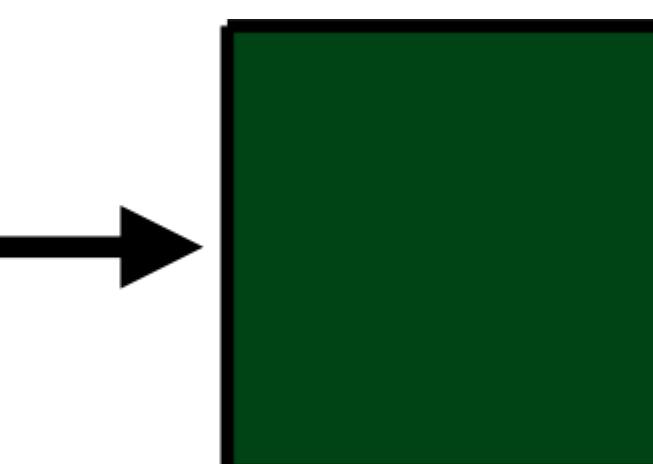
composited over white



premultiplied
color



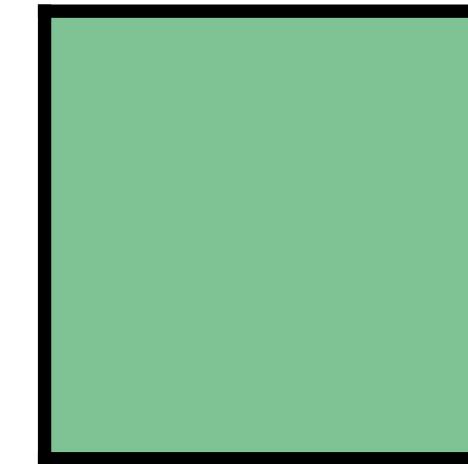
premultiplied
 α



filtered color



filtered α



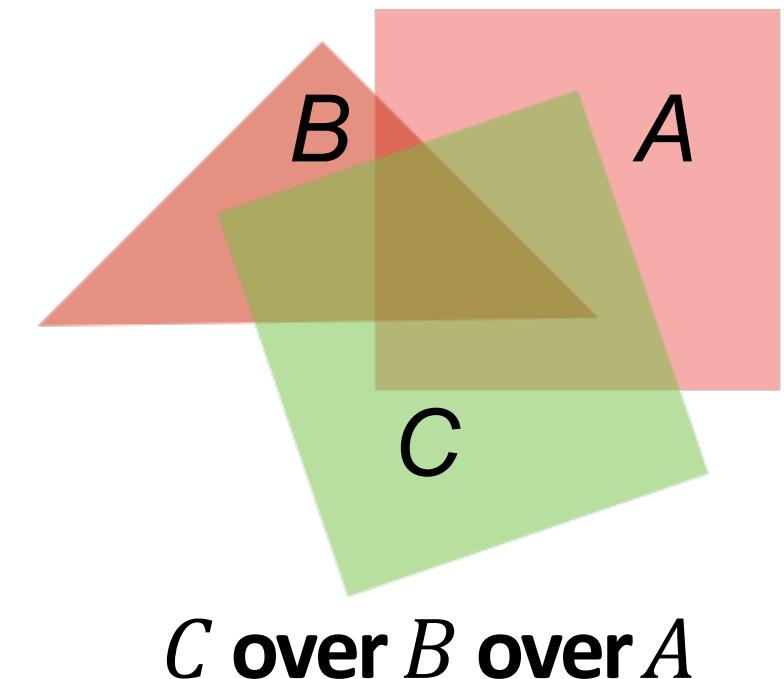
composited over white

More problems: applying “over” repeatedly

Composite image C with opacity α_C over B with opacity α_B over image A with opacity α_A

Premultiplied alpha is closed under composition;
non-premultiplied alpha is not!

Example: composite 50% bright red over 50% bright red
(where “bright red” = $(1,0,0)$, and $\alpha = 0.5$)



non-premultiplied

color

$$0.5(1,0,0) + (1 - 0.5)0.5(1,0,0)$$

↓

too dark!

$$(0.75,0,0)$$

alpha

$$0.5 + (1 - 0.5)0.5 = 0.75$$

premultiplied

color

$$(0.5,0,0,0.5) + (1 - 0.5)(0.5,0,0,0.5)$$

↓

$$(0.75,0,0,0.75)$$

↓

bright red! $(1,0,0)$

alpha

$$\alpha = 0.75$$

Summary: advantages of premultiplied alpha

- Compositing operation treats all channels the same (color and α)
- Fewer arithmetic operations for “over” operation than with non-premultiplied representation
- Closed under composition (repeated “over” operations)
- Better representation for filtering (upsampling/downsampling) images with alpha channel
- Fits naturally into rasterization pipeline (homogeneous coordinates)

Strategy for drawing semi-transparent primitives

Assuming all primitives are semi-transparent, and color values are encoded with premultiplied alpha, here's a strategy for rasterizing an image:

```
over(c1, c2)
{
    return c1.rgb + (1-c1.a) * c2.rgb;
```

```
update_color_buffer( x, y, sample_color, sample_depth)
{
    if (pass_depth_test(sample_depth, zbuffer[x][y])
    {
        // (how) should we update depth buffer here??
        color[x][y] = over(sample_color, color[x][y]);
    }
}
```

Q: What is the assumption made by this implementation?

Triangles must be rendered in back to front order!

Putting it all together

What if we have a mixture of opaque and transparent triangles?

Step 1: render opaque primitives (in any order) using depth-buffered occlusion

If pass depth test, triangle overwrites value in color buffer at sample

Step 2: disable depth buffer update, render semi-transparent surfaces in back-to-front order. If pass depth test, triangle is composited Over contents of color buffer at sample

